



# 第8章 异常控制流

信号和非本地跳转 Signals and Nonlocal Jumps

100076202: 计算机系统导论



任课教师:

宿红毅 张艳 黎有琦 李秀星

原作者:

Randal E. Bryant and David R. O'Hallaron

Carnegie  
Mellon  
University



# 上次课复习 Review from last lecture

## ■ 异常 Exceptions

- 需要非标准控制流的事件 Events that require nonstandard control flow
- 由外部（中断）或内部（陷阱和故障）生成 Generated externally (interrupts) or internally (traps and faults)

## ■ 进程 Processes

- 在任何给定时间，系统都有多个活动进程 At any given time, system has multiple active processes
- 一次只能在任何单个内核上执行一个进程 Only one can execute at a time on any single core
- 每个进程似乎都可以完全控制处理器+专用内存空间 Each process appears to have total control of processor + private memory space



# 复习 (续) Review (cont.)

- **创建进程 Spawning processes**
  - 调用fork Call `fork`
  - 一次调用, 两次返回 One call, two returns
- **进程完成 Process completion**
  - 调用exit Call `exit`
  - 调用一次, 不返回 One call, no return
- **回收和等待进程 Reaping and waiting for processes**
  - 调用wait或waitpid Call `wait` or `waitpid`
- **加载和运行程序 Loading and running programs**
  - 调用execve (或变种) Call `execve` (or variant)
  - 调用一次, (正常) 不返回 One call, (normally) no return

# execve: 加载并运行程序

## execve: Loading and Running Programs



- `int execve(char *filename, char *argv[], char *envp[])`
- 加载并在当前进程运行: **Loads and runs in the current process:**
  - 可执行文件 `filename` Executable file `filename`
    - 可以是目标文件或以“#!解释器”开始的脚本文件 Can be object file or script file beginning with `#!interpreter` (e.g., `#!/bin/bash`)
  - 参数列表 `argv` ...with argument list `argv`
    - 按照规则 `argv[0]` 为文件名 By convention `argv[0]==filename`
  - 和环境变量列表 `envp` ...and environment variable list `envp`
    - “名字=值”串 “name=value” strings (e.g., `USER=droh`)
    - `getenv`, `putenv`, `printenv`
- 覆盖代码、数据和栈 **Overwrites code, data, and stack**
  - 维持PID、打开文件和信号上下文 Retains PID, open files and signal context
- 调用一次，从不返回 **Called once and never returns**
  - 除非如果发生错误 ...except if there is an error



# 异常控制流存在系统每个层次

## ECF Exists at All Levels of a System

### ■ 异常 **Exceptions**

- 硬件和操作系统内核软件
- Hardware and operating system kernel software

### ■ 进程上下文切换 **Process Context Switch**

- 硬件时钟和内核软件
- Hardware timer and kernel software

### ■ 信号 **Signals**

- 内核软件和应用软件
- Kernel software and application software

### ■ 非局部跳转 **Nonlocal jumps**

- 应用代码 Application code

**Previous Lecture**  
以前的课

**This Lecture**  
本次课

教材和补充幻灯片  
**Textbook and  
supplemental slides**

# (部分) 分类

## (partial) Taxonomy

内核处理 Handled in kernel

用户进程处理 Handled in user process

异常控制流ECF

异步Asynchronous

同步 Synchronous

中断Interrupts

信号Signals

陷阱 Traps

故障 Faults

终止Aborts

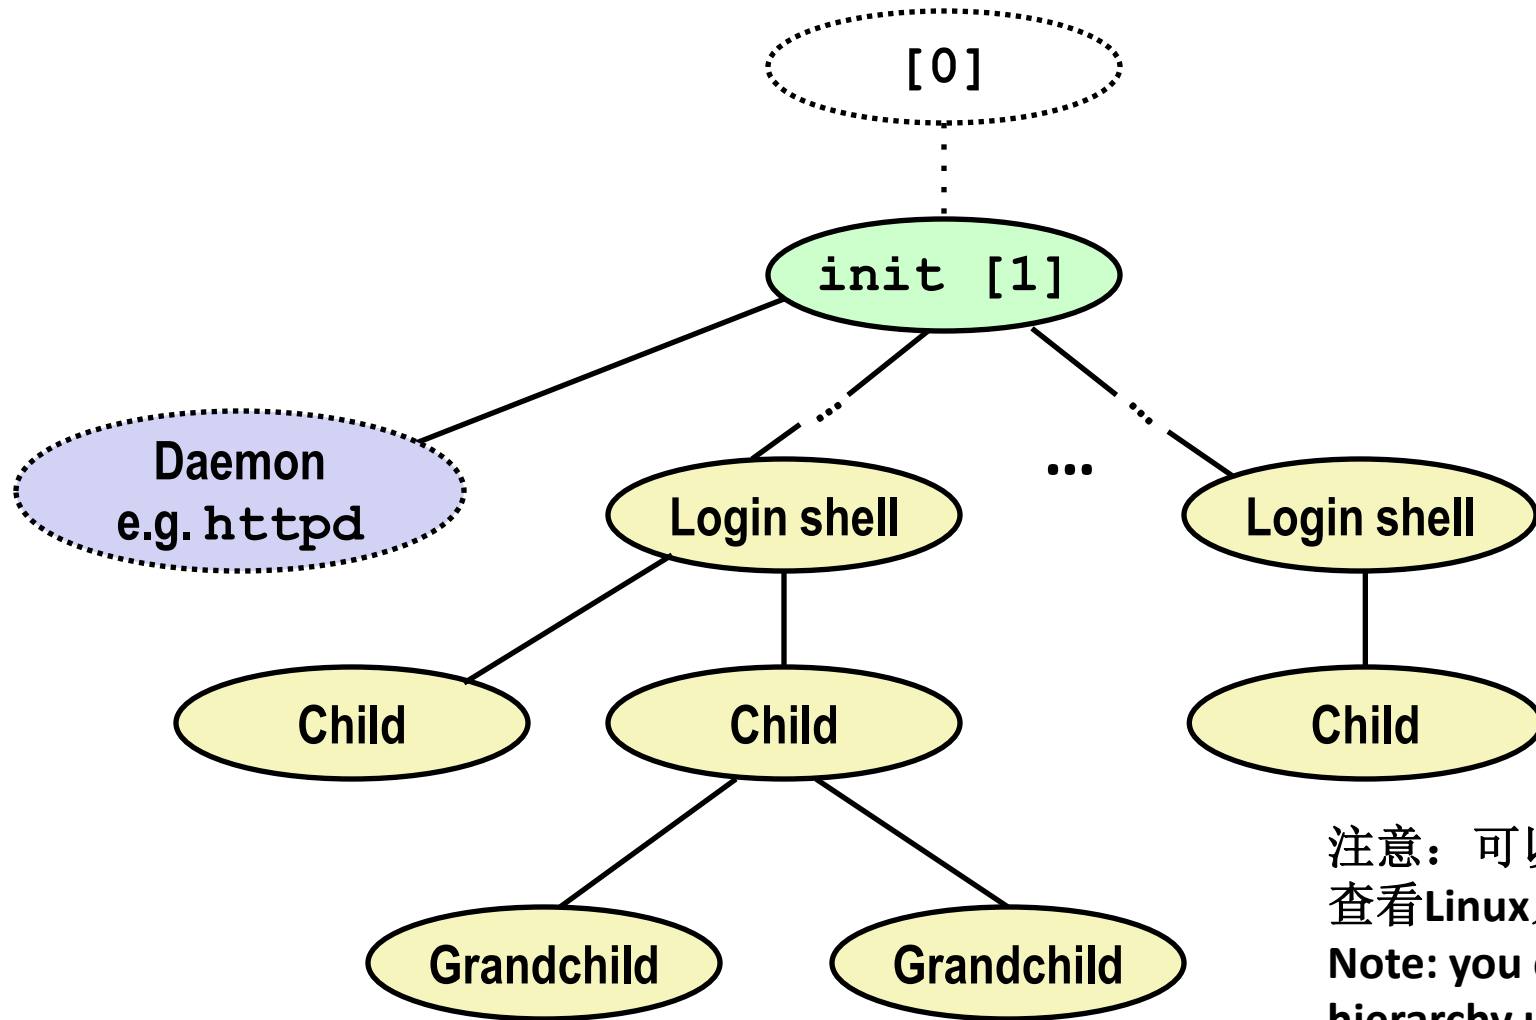


# 议题

- 外壳 **Shells**
- 信号 **Signals**
- 非局部跳转 **Nonlocal jumps**



# Linux进程树 Linux Process Hierarchy



注意：可以用`ps tree`命令  
查看Linux系统的进程树  
Note: you can view the  
hierarchy using the Linux  
`ps tree` command





# Shell程序 Shell Programs

- Shell是按照用户要求运行程序的应用程序 **A *shell* is an application program that runs programs on behalf of the user**
  - **sh** 最早的 Original Unix shell (Stephen Bourne, AT&T Bell Labs, 1977)
  - **csch/tcsch** BSD Unix C shell
  - **bash** 默认的 “Bourne-Again” Shell (default Linux shell)
- **简单shell Simple shell**
  - 教材p753页处描述 Described in the textbook, starting at p. 753
  - 一个非常基础的shell实现 Implementation of a very elementary shell
  - 目的 Purpose
    - 理解当输入了命令后究竟发生了什么事情 Understand what happens when you type commands
    - 理解进程控制操作的使用和操作 Understand use and operation of process control operations



# 简单shell示例 Simple Shell Example

```
linux> ./shellex
```

```
> /bin/ls -l csapp.c 必须给出程序的全路径名 Must give full pathnames for programs
```

```
-rw-r--r-- 1 bryant users 23053 Jun 15 2015 csapp.c
```

```
> /bin/ps
```

PID	TTY	TIME	CMD
31542	pts/2	00:00:01	tcsh
32017	pts/2	00:00:00	shellex
32019	pts/2	00:00:00	ps

```
> /bin/sleep 10 & 后台运行程序 Run program in background
```

```
32031 /bin/sleep 10 &
```

```
> /bin/ps
```

PID	TTY	TIME	CMD
31542	pts/2	00:00:01	tcsh
32024	pts/2	00:00:00	emacs
32030	pts/2	00:00:00	shellex
32031	pts/2	00:00:00	sleep
32033	pts/2	00:00:00	ps

**Sleep正在后台运行**

**Sleep is running**

```
> quit
```

**in background**



# 简单shell实现

## Simple Shell Implementation

### ■ 基本循环 Basic loop

- 从命令行读一行 Read line from command line
- 执行请求的操作 Execute the requested operation
  - 内置命令（仅实现一个命令是**quit**） Built-in command (only one implemented is **quit**)
  - 从文件加载和执行程序 Load and execute program from file

```
int main(int argc, char** argv)
{
    char cmdline[MAXLINE]; /* command line */

    while (1) {
        /* read */
        printf("> ");
        fgets(cmdline, MAXLINE, stdin);
        if (feof(stdin))
            exit(0);

        /* evaluate */
        eval(cmdline);
    }
    ...
}
```

*shellex.c*

执行的过程就是一系列读/求值的步骤  
Execution is a sequence of read/evaluate steps



# 简单的Shell eval函数 Simple Shell eval Function

```
void eval(char *cmdline)
{
    char *argv[MAXARGS]; /* Argument list execve() */
    char buf[MAXLINE]; /* Holds modified command line */
    int bg; /* Should the job run in bg or fg? */
    pid_t pid; /* Process id */

    strcpy(buf, cmdline);
    bg = parseline(buf, argv);
    if (argv[0] == NULL)
        return; /* Ignore empty lines */

    if (!builtin_command(argv)) {
        if ((pid = Fork()) == 0) { /* Child runs user job */
            if (execve(argv[0], argv, environ) < 0) {
                printf("%s: Command not found.\n", argv[0]);
                exit(0);
            }
        }

        /* Parent waits for foreground job to terminate */
        if (!bg) {
            int status;
            if (waitpid(pid, &status, 0) < 0)
                unix_error("waitfg: waitpid error");
        }
        else
            printf("%d %s", pid, cmdline);
    }
    return;
}
```

# 简单的Shell eval函数 Simple Shell eval Function



```
void eval(char *cmdline)
{
    char *argv[MAXARGS]; /* Argument list execve() */
    char buf[MAXLINE];    /* Holds modified command line */
    int bg;               /* Should the job run in bg or fg? */
    pid_t pid;            /* Process id */

    strcpy(buf, cmdline);
    bg = parseline(buf, argv);
}
```

**Parseline**函数将buf解析成argv并返回是否输入行以&结尾  
**parseline** will parse 'buf' into 'argv' and return whether or not input line ended in '&'

# 简单的Shell eval函数 Simple Shell eval Function



```
void eval(char *cmdline)
{
    char *argv[MAXARGS]; /* Argument list execve() */
    char buf[MAXLINE];    /* Holds modified command line */
    int bg;               /* Should the job run in bg or fg? */
    pid_t pid;            /* Process id */

    strcpy(buf, cmdline);
    bg = parseline(buf, argv);
    if (argv[0] == NULL)
        return; /* Ignore empty lines */
```

忽略空行

Ignore empty lines.

# 简单的Shell eval函数 Simple Shell eval Function



```
void eval(char *cmdline)
{
    char *argv[MAXARGS]; /* Argument list execve() */
    char buf[MAXLINE];    /* Holds modified command line */
    int bg;               /* Should the job run in bg or fg? */
    pid_t pid;            /* Process id */

    strcpy(buf, cmdline);
    bg = parseline(buf, argv);
    if (argv[0] == NULL)
        return; /* Ignore empty lines */

    if (!builtin_command(argv)) {
```

如果是“内置”命令，那么在这个程序此处处理它。否则创建进程(fork) / 执行(exec) 在argv[0]中指定的程序

If it is a 'built in' command, then handle it here in this program.  
Otherwise fork/exec the program specified in argv[0]

# 简单的Shell eval函数 Simple Shell eval Function



```
void eval(char *cmdline)
{
    char *argv[MAXARGS]; /* Argument list execve() */
    char buf[MAXLINE];    /* Holds modified command line */
    int bg;               /* Should the job run in bg or fg? */
    pid_t pid;            /* Process id */

    strcpy(buf, cmdline);
    bg = parseline(buf, argv);
    if (argv[0] == NULL)
        return; /* Ignore empty lines */

    if (!builtin_command(argv)) {
        if ((pid = fork()) == 0) { /* Child runs user job */
```

创建子进程/Create child



# 简单的Shell eval函数 Simple Shell eval Function



```
void eval(char *cmdline)
{
    char *argv[MAXARGS]; /* Argument list execve() */
    char buf[MAXLINE];    /* Holds modified command line */
    int bg;               /* Should the job run in bg or fg? */
    pid_t pid;            /* Process id */

    strcpy(buf, cmdline);
    bg = parseline(buf, argv);
    if (argv[0] == NULL)
        return; /* Ignore empty lines */

    if (!builtin_command(argv)) {
        if ((pid = fork()) == 0) { /* Child runs user job */
            if (execve(argv[0], argv, environ) < 0) {
                printf("%s: Command not found.\n", argv[0]);
                exit(0);
            }
        }
    }
}
```

启动argv[0].

记住execve仅在出错时返回

Start **argv[0]**.

Remember **execve** only returns on error.

# 简单的Shell eval函数 Simple Shell eval Function



```
void eval(char *cmdline)
{
    char *argv[MAXARGS]; /* Argument list execve() */
    char buf[MAXLINE];    /* Holds modified command line */
    int bg;               /* Should the job run in bg or fg? */
    pid_t pid;            /* Process id */

    strcpy(buf, cmdline);
    bg = parseline(buf, argv);
    if (argv[0] == NULL)
        return; /* Ignore empty lines */

    if (!builtin_command(argv)) {
        if ((pid = fork()) == 0) { /* Child runs user job */
            if (execve(argv[0], argv, environ) < 0) {
                printf("%s: Command not found.\n", argv[0]);
                exit(0);
            }
        }

        /* Parent waits for foreground job to terminate */
        if (!bg) {
            int status;
            if (waitpid(pid, &status, 0) < 0)
                unix_error("waitfg: waitpid error");
        }
    }
}
```

如果子进程在前台运行，等待直到子进程完成  
If running child in foreground, wait until it is done.

# 简单的Shell eval函数 Simple Shell eval Function



```
void eval(char *cmdline)
{
    char *argv[MAXARGS]; /* Argument list execve() */
    char buf[MAXLINE];    /* Holds modified command line */
    int bg;               /* Should the job run in bg or fg? */
    pid_t pid;            /* Process id */

    strcpy(buf, cmdline);
    bg = parseline(buf, argv);
    if (argv[0] == NULL)
        return; /* Ignore empty lines */

    if (!builtin_command(argv)) {
        if ((pid = fork()) == 0) { /* Child runs user job */
            if (execve(argv[0], argv, environ) < 0) {
                printf("%s: Command not found.\n", argv[0]);
                exit(0);
            }
        }

        /* Parent waits for foreground job to terminate */
        if (!bg) {
            int status;
            if (waitpid(pid, &status, 0) < 0)
                unix_error("waitfg: waitpid error");
        }
        else
            printf("%d %s", pid, cmdline);
    }
    return;
}
```

如果子进程在后台运行，打印pid并继续做其它事情

If running child in background, print pid and continue doing other stuff.

# 简单的Shell eval函数 Simple Shell eval Function



```
void eval(char *cmdline)
{
    char *argv[MAXARGS]; /* Argument list execve() */
    char buf[MAXLINE];    /* Holds modified command line */
    int bg;               /* Should the job run in bg or fg? */
    pid_t pid;            /* Process id */

    strcpy(buf, cmdline);
    bg = parseline(buf, argv);
    if (argv[0] == NULL)
        return; /* Ignore empty lines */

    if (!builtin_command(argv)) {
        if ((pid = fork()) == 0) { /* Child runs user job */
            if (execve(argv[0], argv, environ) < 0) {
                printf("%s: Command not found.\n", argv[0]);
                exit(0);
            }
        }

        /* Parent waits for foreground job to terminate */
        if (!bg) {
            int status;
            if (waitpid(pid, &status, 0) < 0)
                unix_error("waitfg: waitpid error");
        }
        else
            printf("%d %s", pid, cmdline);
    }
    return;
}
```

哎呀。此代码有问题。  
Oops. There is a problem with this code.

# 简单Shell程序存在的问题

## Problem with Simple Shell Example



- **Shell设计成无限循环运行 Shell designed to run indefinitely**
  - 不应该积累不需要的资源 Should not accumulate unneeded resources
    - 内存 Memory
    - 子进程 Child processes
    - 文件描述符 File descriptors
- **例子shell只能等待并回收前台作业 Our example shell correctly waits for and reaps foreground jobs**
- **后台作业怎么办？ But what about background jobs?**
  - 终止后变成僵尸 Will become zombies when they terminate
  - 由于shell不会终止，所以永远不会被回收 Will never be reaped because shell (typically) will not terminate
  - 会造成系统内存泄露并耗尽内核内存 Will create a memory leak that could run the kernel out of memory



# 可以利用ECF解决 ECF to the Rescue!

- 解决方案：异常控制流 **Solution: Exceptional control flow**
  - 在后台进程处理完成后，内核打断正常处理流程并提醒我们 The kernel will interrupt regular processing to alert us when a background process completes
  - Unix系统中这种提醒的机制是信号 In Unix, the alert mechanism is called a *signal*



# 议题

- 外壳 **Shells**
- 信号 **Signals**
- 非局部跳转 **Nonlocal jumps**

# 信号 Signals



- 信号是一条小消息，用来通知一个进程某种类型的事件在系统中发生了 **A *signal* is a small message that notifies a process that an event of some type has occurred in the system**
  - 类似于异常和中断 Akin to exceptions and interrupts
  - 由内核发送给一个进程（有时是根据另一个进程的请求） Sent from the kernel (sometimes at the request of another process) to a process
  - 信号的类型是用1-30的小整型标识 Signal type is identified by small integer ID's (1-30)
  - 信号的唯一信息就是这个ID以及信号达到的事实 Only information in a signal is its ID and the fact that it arrived

<i>ID</i>	<i>Name</i>	<i>Default Action</i>	<i>Corresponding Event</i>
2	SIGINT	Terminate	用户输入ctrl-c User typed ctrl-c
9	SIGKILL	Terminate	杀死程序（不能覆盖或被忽略） Kill program (cannot override or ignore)
11	SIGSEGV	Terminate	段错误 Segmentation violation
14	SIGALRM	Terminate	时钟信号 Timer signal
17	SIGCHLD	Ignore	子进程停止或者终止 Child stopped or terminated





# 信号概念：发送一个信号

## Signal Concepts: Sending a Signal

- 内核通过更新目标进程上下文的某些状态来**发送**（传递）一个信号给**目标进程** Kernel **sends** (delivers) a signal to a **destination process** by updating some state in the context of the destination process
- 内核发送信号是由于以下原因之一 Kernel sends a signal for one of the following reasons:
  - 内核侦测到除零错误（SIGFPE）或者子进程终止（SIGCHLD）等系统事件 Kernel has detected a system event such as divide-by-zero (SIGFPE) or the termination of a child process (SIGCHLD)
  - 另外一个进程调用了kill系统调用显式请求内核发送一个信号给目标进程 Another process has invoked the **kill** system call to explicitly request the kernel to send a signal to the destination process

# 信号概念：发送一个信号

## Signal Concepts: Sending a Signal



用户级  
User level

进程B  
Process B

进程A  
Process A

进程C  
Process C

内核  
kernel

		挂起 Pending for A
		挂起 Pending for B
		挂起 Pending for C

		阻塞 Blocked for A
		阻塞 Blocked for B
		阻塞 Blocked for C

# 信号概念：发送一个信号

## Signal Concepts: Sending a Signal



用户级  
User level

进程B  
Process B

进程A  
Process A

进程C  
Process C

Sends to C

			Pending for A
			Pending for B
			Pending for C

			Blocked for A
			Blocked for B
			Blocked for C

内核  
kernel

# 信号概念：发送一个信号

## Signal Concepts: Sending a Signal



用户级  
User level

进程B  
Process B

进程A  
Process A

进程C  
Process C

内核  
kernel

			Pending for A
			Pending for B
	1		Pending for C

			Blocked for A
			Blocked for B
			Blocked for C

# 信号概念：发送一个信号Signal

## Concepts: Sending a Signal



用户级  
User level

进程B  
Process B

进程A  
Process A

进程C  
Process C

内核  
kernel

			Pending for A
			Pending for B
1			Pending for C

			Blocked for A
			Blocked for B
			Blocked for C

Received by C

# 信号概念：发送一个信号Signal

## Concepts: Sending a Signal



用户级  
User level

进程B  
Process B

进程A  
Process A

进程C  
Process C

内核  
kernel

			Pending for A
			Pending for B
		0	Pending for C

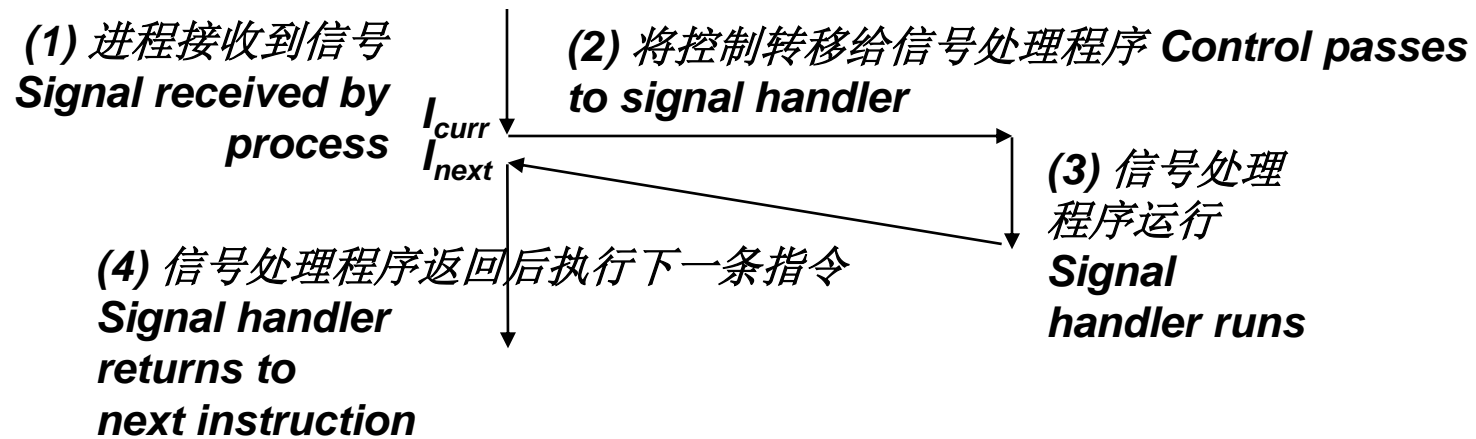
			Blocked for A
			Blocked for B
			Blocked for C



# 信号概念：接收一个信号

## Signal Concepts: Receiving a Signal

- 目标进程**接收**信号是由于系统内核强制其对某个信号的发送做出响应  
A destination process **receives** a signal when it is forced by the kernel to react in some way to the delivery of the signal
- 可能的响应方式 **Some possible ways to react:**
  - **忽略**信号（什么也不做） **Ignore** the signal (do nothing)
  - **终止进程**（可以选择对信息转储） **Terminate** the process (with optional core dump)
  - **调用**用户级**信号处理函数**对信号进行处理 **Catch** the signal by executing a user-level function called **signal handler**
    - 类似于硬件异常处理函数对异步中断的响应 Akin to a hardware exception handler being called in response to an asynchronous interrupt:



# 信号概念：挂起或者阻塞的信号

## Signal Concepts: Pending and Blocked Signals



- 已经发送但是没有被接收的信号处于 **挂起** 状态 A signal is **pending** if sent but not yet received
  - 任何特定类型的信号最多有一个挂起的 There can be at most one pending signal of any particular type
  - 重要：信号不排队 Important: Signals are not queued
    - 如有某个进程有一个类型为k的信号挂起，则后续发给该进程的k类信号被直接抛弃 If a process has a pending signal of type k, then subsequent signals of type k that are sent to that process are discarded
- 一个进程会 **阻塞** 某种特定类型信号的接收 A process can **block** the receipt of certain signals
  - 阻塞的信号可以发送，但是在解除阻塞前不会被接收 Blocked signals can be delivered, but will not be received until the signal is unblocked
  - 有些信号不能被阻塞（SIGKILL, SIGSTOP）或者仅当其它进程发送（SIGSEGV、SIGILL等）时被阻塞 Some signals cannot be blocked (SIGKILL, SIGSTOP) or can only be blocked when sent by other processes (SIGSEGV, SIGILL, etc)
- 挂起的信号最多被接收一次 A pending signal is received at most once





# 信号概念：挂起/阻塞位

## Signal Concepts: Pending/Blocked Bits

- 内核在每个进程的上下文维护一个挂起和阻塞的比特向量 **Kernel maintains pending and blocked bit vectors in the context of each process**
  - **挂起：表示挂起的信号集合** **pending**: represents the set of pending signals
    - 当发送了一个k类型的信号时系统设置第k个比特位 **Kernel sets bit k in pending** when a signal of type k is delivered
    - 当类型k的信号被接收后系统会将第k个比特位清零 **Kernel clears bit k in pending** when a signal of type k is received
  - **阻塞：表示阻塞的信号集合** **blocked**: represents the set of blocked signals
    - 可以使用**sigprocmask**函数设置或者清除 **Can be set and cleared by using the sigprocmask function**
    - 也称为信号掩码 **Also referred to as the signal mask.**

# 信号概念：发送信号

## Signal Concepts: Sending a Signal



用户级  
User level

进程B  
Process B

进程A  
Process A

进程C  
Process C

Sends to C

内核  
kernel

			Pending for A
			Pending for B
	1		Pending for C

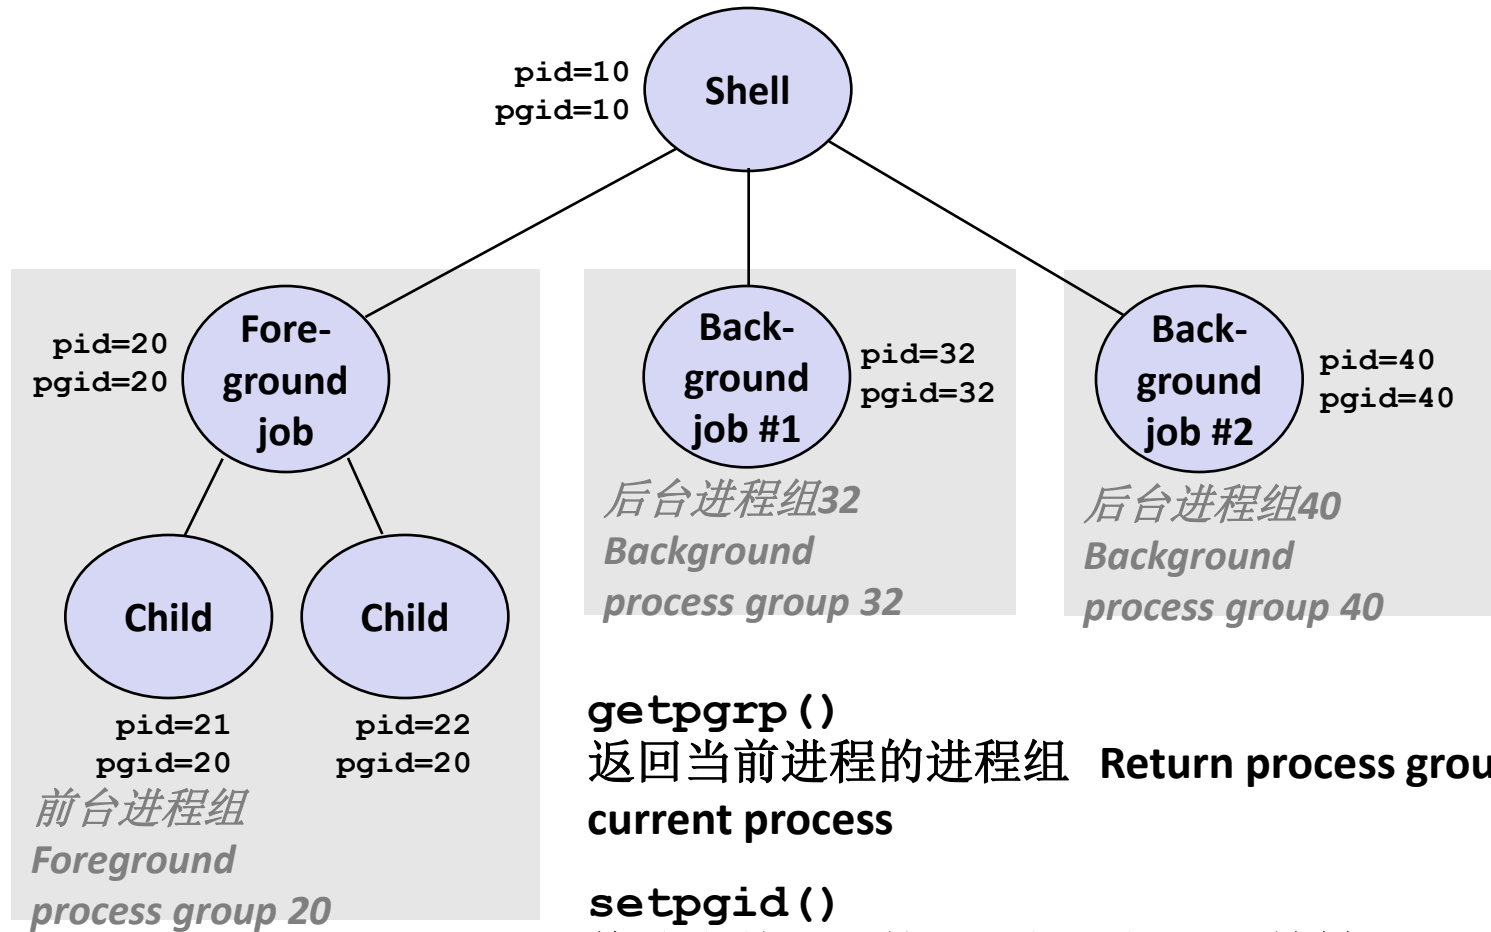
			Blocked for A
			Blocked for B
			Blocked for C



# 发送信号：进程组

## Sending Signals: Process Groups

- 每个进程只属于一个进程组 Every process belongs to exactly one process group



`getpgrp()`

返回当前进程的进程组 Return process group of current process

`setpgid()`

修改当前进程的进程组（细节见教材） Change process group of a process (see text for details)



# 通过/bin/kill程序发送信号

## Sending Signals with /bin/kill Program

- /bin/kill程序可以发送任意信号给一个进程或者进程组 /bin/kill program sends arbitrary signal to a process or process group

- 例如 Examples

- /bin/kill -9 24818 发送SIGKILL给进程 24818 Send SIGKILL to process 24818
- /bin/kill -9 -24817 发送SIGKILL给进程组的每个进程 Send SIGKILL to every process in process group 24817

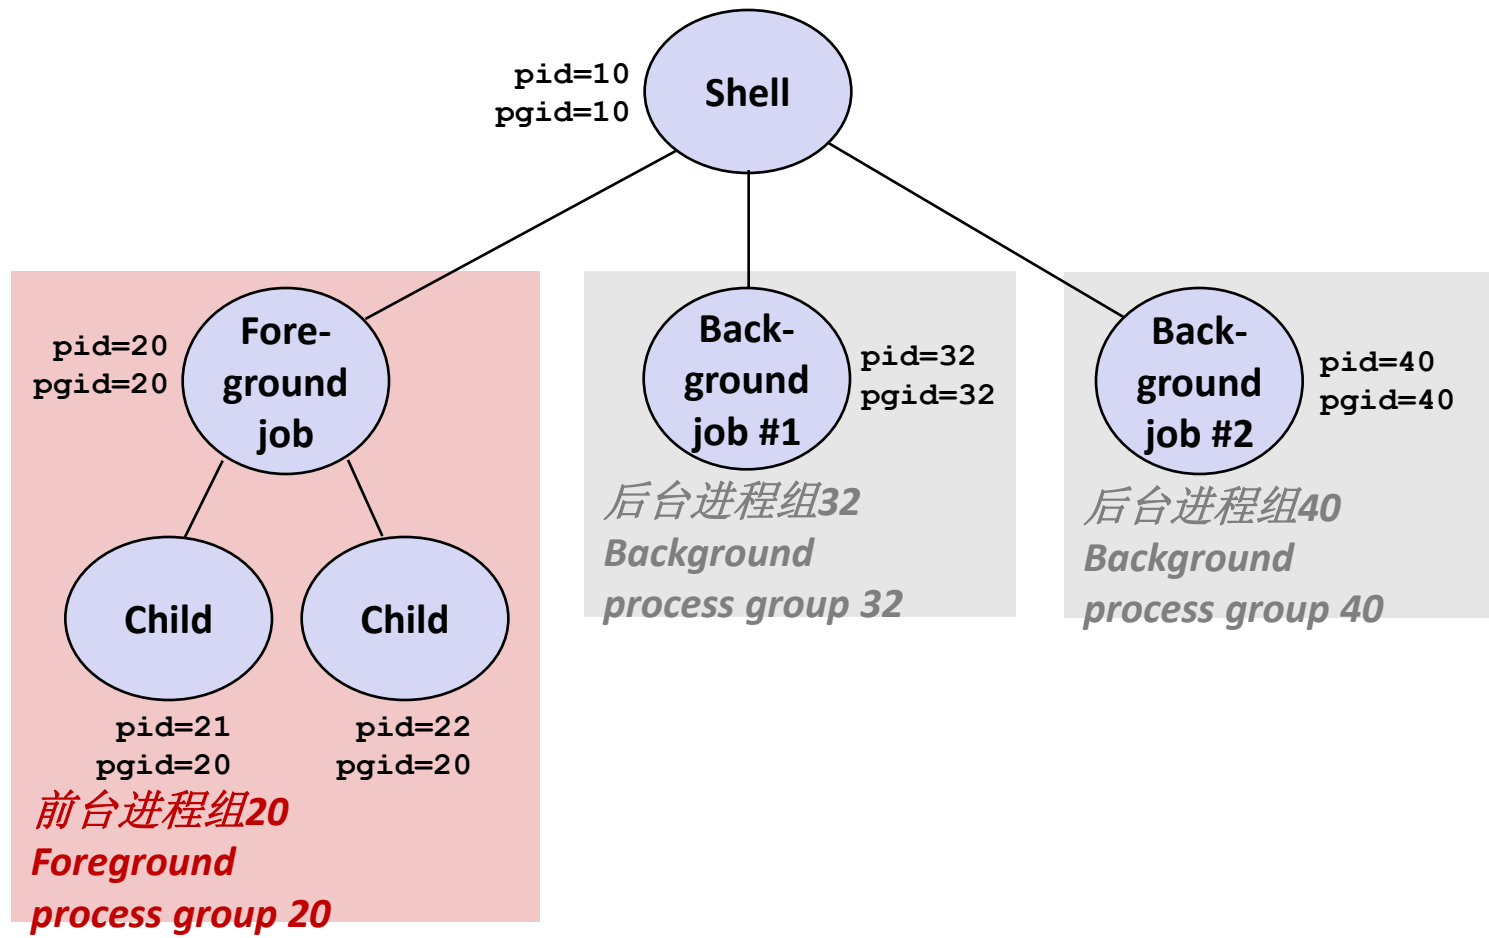
```
linux> ./forks 16
Child1: pid=24818 pgrp=24817
Child2: pid=24819 pgrp=24817

linux> ps
  PID TTY          TIME CMD
 24788 pts/2        00:00:00 tcsh
 24818 pts/2        00:00:02 forks
 24819 pts/2        00:00:02 forks
 24820 pts/2        00:00:00 ps
linux> /bin/kill -9 -24817
linux> ps
  PID TTY          TIME CMD
 24788 pts/2        00:00:00 tcsh
 24823 pts/2        00:00:00 ps
linux>
```

# 通过键盘发送信号 Sending Signals from the Keyboard



- 输入ctrl-c(ctrl-z)会导致系统内核发送一个SIGINT (SIGTSTP) 信号给前台进程组的每个作业 Typing ctrl-c (ctrl-z) causes the kernel to send a SIGINT (SIGTSTP) to every job in the foreground process group.
  - SIGINT – default action is to terminate each process 默认终止每个进程
  - SIGTSTP – default action is to stop (suspend) each process 默认停止（挂起）每个进程





# ctrl-c和ctrl-z示例

## Example of ctrl-c and ctrl-z

```
bluefish> ./forks 17
Child: pid=28108 pgrp=28107
Parent: pid=28107 pgrp=28107
<types ctrl-z>
Suspended
bluefish> ps w
```

PID	TTY	STAT	TIME	COMMAND
27699	pts/8	Ss	0:00	-tcsh
28107	pts/8	T	0:01	./forks 17
28108	pts/8	T	0:01	./forks 17
28109	pts/8	R+	0:00	ps w

```
bluefish> fg
./forks 17
<types ctrl-c>
bluefish> ps w
```

PID	TTY	STAT	TIME	COMMAND
27699	pts/8	Ss	0:00	-tcsh
28110	pts/8	R+	0:00	ps w

进程状态STAT标记 STAT  
(process state) Legend:

**First letter 第一个字母:**

S: sleeping 睡眠

T: stopped 停止

R: running 运行

**Second letter 第二个字母:**

s: session leader 会话首领

+: foreground proc group 前台  
进程组

参见“man ps”了解更多细节  
See “man ps” for more  
details



# 通过kill函数发送信号

## Sending Signals with kill Function

```
void fork12()
{
    pid_t pid[N];
    int i;
    int child_status;

    for (i = 0; i < N; i++)
        if ((pid[i] = fork()) == 0) {
            /* Child: Infinite Loop */
            while(1)
                ;
        }

    for (i = 0; i < N; i++) {
        printf("Killing process %d\n", pid[i]);
        kill(pid[i], SIGINT);
    }

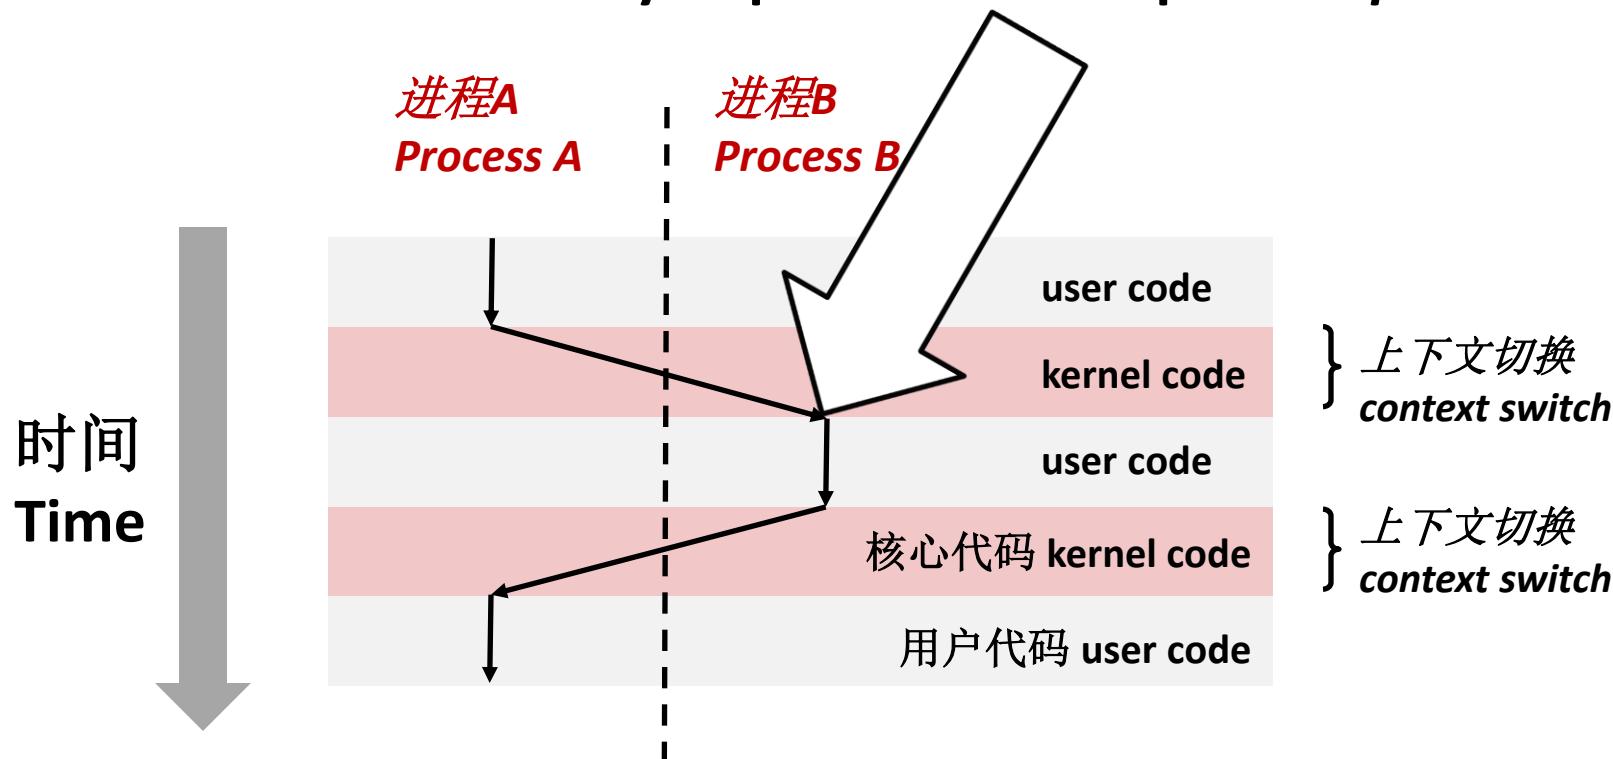
    for (i = 0; i < N; i++) {
        pid_t wpid = wait(&child_status);
        if (WIFEXITED(child_status))
            printf("Child %d terminated with exit status %d\n",
                wpid, WEXITSTATUS(child_status));
        else
            printf("Child %d terminated abnormally\n", wpid);
    }
}
```

*forks.c*



# 接收信号 Receiving Signals

- 假设内核正从异常处理函数返回，并准备把控制权传递给进程p  
Suppose kernel is returning from an exception handler and is ready to pass control to process  $p$





# 接收信号 Receiving Signals



- 假设内核正从异常处理函数返回，并准备把控制权传递给进程p  
Suppose kernel is returning from an exception handler and is ready to pass control to process  $p$
- 内核计算 Kernel computes  $pnb = pending \ \& \ \sim blocked$ 
  - 进程p挂起但非阻塞信号的集合 The set of pending nonblocked signals for process  $p$
- 如果集合为空 If ( $pnb == 0$ )
  - 将控制权交给进程p逻辑流的下一条指令 Pass control to next instruction in the logical flow for  $p$
- 否则 Else
  - 选择pnb中最低非0位k并强制进程p接收信号k Choose least nonzero bit  $k$  in  $pnb$  and force process  $p$  to **receive** signal  $k$
  - 信号的接收触发了p的某些动作 The receipt of the signal triggers some **action** by  $p$
  - 对pnb中每个非0位k重复上述过程 Repeat for all nonzero  $k$  in  $pnb$
  - 将控制权交给进程p逻辑流的下一条指令 Pass control to next instruction in logical flow for  $p$



# 默认动作 Default Actions

- 每种类型的信号有一个预定义的**默认动作**，可能是如下中的一个 Each signal type has a predefined **default action**, which is one of:
  - 终止进程 The process terminates
  - 停止进程，直到接收到SIGCONT时重启 The process stops until restarted by a SIGCONT signal
  - 进程忽略掉该信号 The process ignores the signal

# 安装信号处理程序 Installing Signal Handlers



- 函数signal修改接收信号signum对应的默认行为 The signal function modifies the default action associated with the receipt of signal signum:
  - `handler_t *signal(int signum, handler_t *handler)`
- 信号处理程序handler的不同值 Different values for handler:
  - SIG\_IGN: ignore signals of type signum 忽略signum类型的信号
  - SIG\_DFL: revert to the default action on receipt of signals of type signum 接收到signum类型的信号时按照默认动作处理
  - 否则handler是用户级信号处理程序的地址 Otherwise, handler is the address of a user-level **signal handler**
    - 当进程接收到类型为signum的信号时调用 Called when process receives signal of type signum
    - 称为安装信号处理程序 Referred to as **“installing”** the handler
    - 执行信号处理程序称为捕获或处理该信号 Executing handler is called **“catching”** or **“handling”** the signal
    - 当信号处理程序执行返回语句时，控制权交给进程接收到信号时被打断控制流中指令 When the handler executes its return statement, control passes back to instruction in the control flow of the process that was interrupted by receipt of the signal

# 信号处理例子 Signal Handling Example



```
void sigint_handler(int sig) /* SIGINT handler */
{
    printf("So you think you can stop the bomb with ctrl-c, do you?\n");
    sleep(2);
    printf("Well...");
    fflush(stdout);
    sleep(1);
    printf("OK. :-)\n");
    exit(0);
}

int main()
{
    /* Install the SIGINT handler */
    if (signal(SIGINT, sigint_handler) == SIG_ERR)
        unix_error("signal error");

    /* Wait for the receipt of a signal */
    pause();

    return 0;
}
```

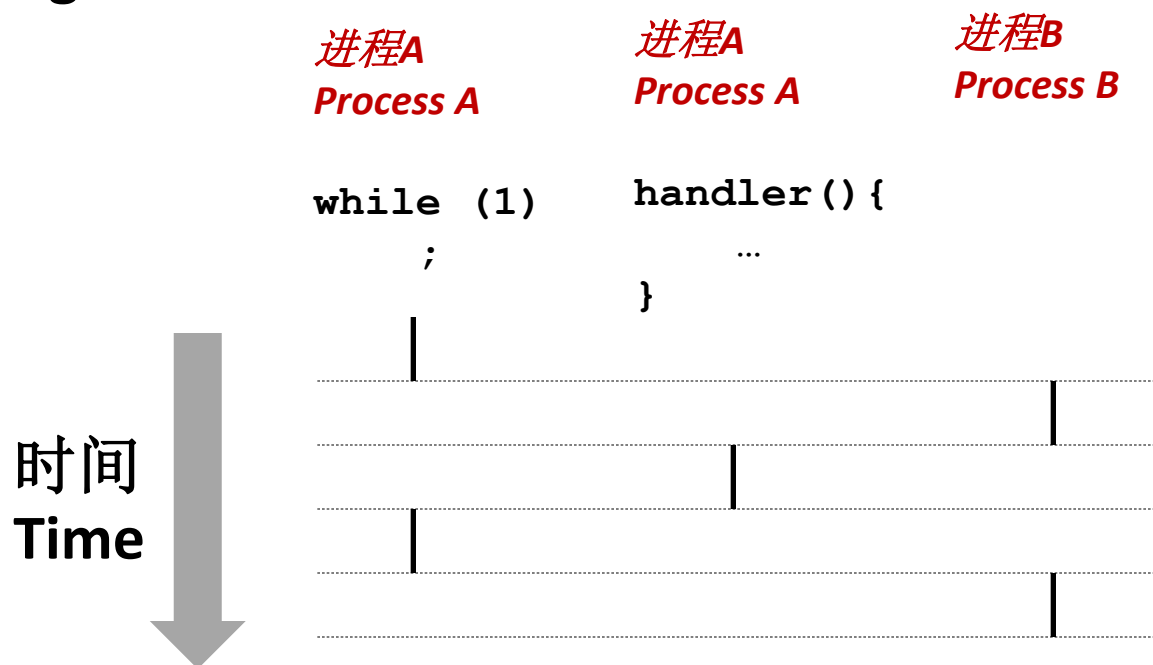
sigint.c

# 信号处理程序作为并发控制流

## Signals Handlers as Concurrent Flows



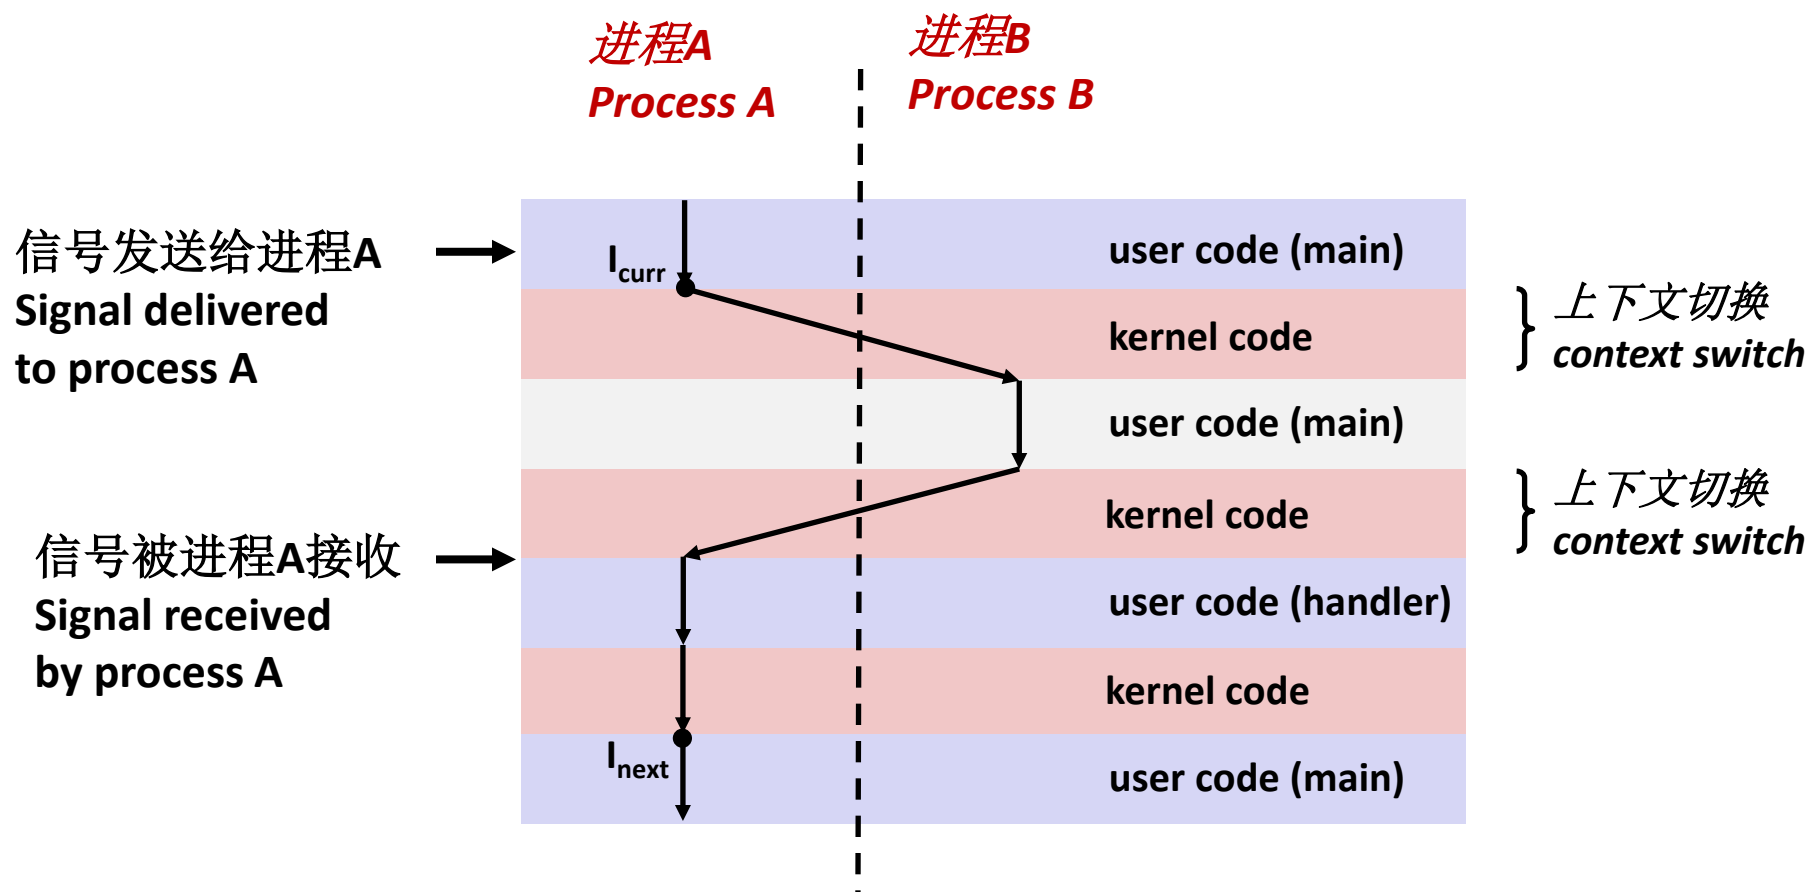
- 每个信号处理程序都是一个独立的逻辑控制流（非进程），与主程序并发执行 A signal handler is a separate logical flow (not process) that runs concurrently with the main program





# 信号处理程序作为并发控制流的另一个视图

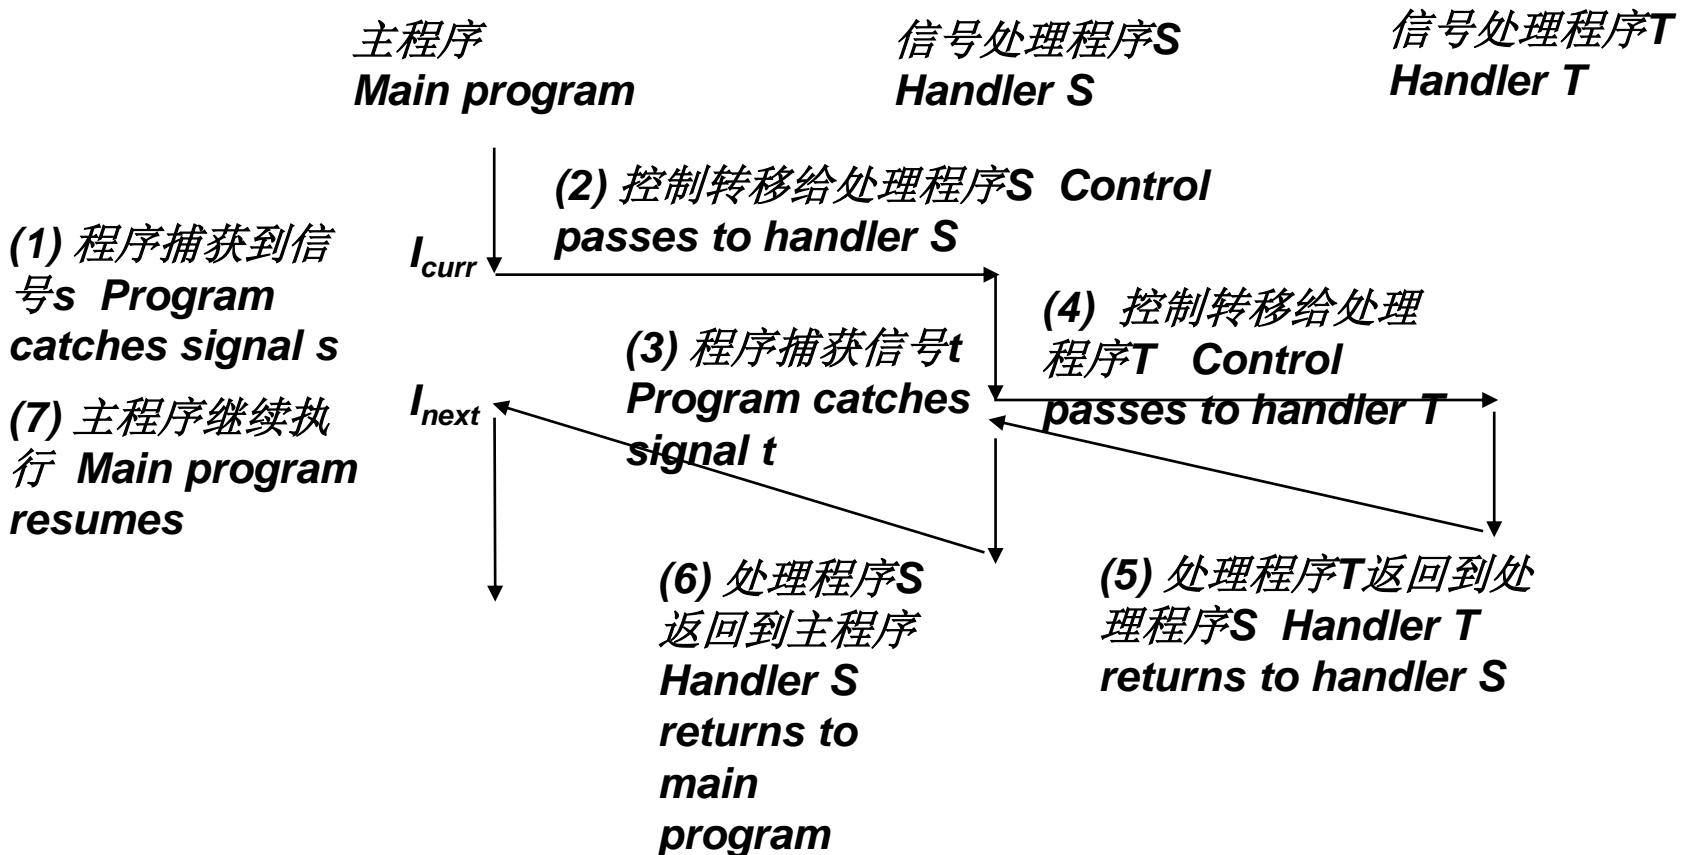
## Another View of Signal Handlers as Concurrent Flows





# 嵌套信号处理 Nested Signal Handlers

- 信号处理程序可能被另一个信号处理程序打断  
Handlers can be interrupted by other handlers





# 阻塞和解除信号阻塞

## Blocking and Unblocking Signals

- **隐式阻塞机制 Implicit blocking mechanism**
  - 内核会阻塞当前正在被处理的任何挂起信号类型 Kernel blocks any pending signals of type currently being handled.
  - 例如SIGINT信号处理程序不能被另一个SIGINT打断 E.g., A SIGINT handler can't be interrupted by another SIGINT
- **显式阻塞和解除阻塞机制 Explicit blocking and unblocking mechanism**
  - `sigprocmask`函数 `sigprocmask` function
- **支持函数 Supporting functions**
  - `sigemptyset` – Create empty set 创建一个空的集合
  - `sigfillset` – Add every signal number to set 对集合设置每个信号编号
  - `sigaddset` – Add signal number to set 对集合设置某个信号编号
  - `sigdelset` – Delete signal number from set 将信号编号从集合删除





# 临时阻塞信号

## Temporarily Blocking Signals

```
sigset_t mask, prev_mask;
```

```
Sigemptyset(&mask);
```

```
Sigaddset(&mask, SIGINT);
```

```
/* Block SIGINT and save previous blocked set */
```

```
Sigprocmask(SIG_BLOCK, &mask, &prev_mask);
```

```
⋮  
/* Code region that will not be interrupted by SIGINT */
```

```
/* Restore previous blocked set, unblocking SIGINT */
```

```
Sigprocmask(SIG_SETMASK, &prev_mask, NULL);
```

# 安全的信号处理

## Safe Signal Handling



- 信号处理程序比较复杂，是因为他们是和主程序并发运行的，并且共享同样的全局数据结构 **Handlers are tricky because they are concurrent with main program and share the same global data structures.**
  - 共享数据结构更容易被破坏 Shared data structures can become corrupted.
- 我们在这学期后面讨论并发的问题 **We'll explore concurrency issues later in the term.**
- 现在只给一些有助避免麻烦的提示 **For now here are some guidelines to help you avoid trouble.**

# 编写安全处理程序的提示

## Guidelines for Writing Safe Handlers



- **G0: 信号处理程序越简单越好 Keep your handlers as simple as possible**
  - 例如, 设置全局标记后返回 e.g., Set a global flag and return
- **G1: 在信号处理程序中只调用异步信号安全的函数 Call only async-signal-safe functions in your handlers**
  - `printf`, `sprintf`, `malloc`, and `exit` are not safe! 这些都不安全
- **G2: 进入和退出时保存和恢复 `errno` Save and restore `errno` on entry and exit**
  - 以便其它的信号处理程序不会覆盖你的 `errno` 值 So that other handlers don't overwrite your value of `errno`
- **G3: 临时阻塞所有的信号后再访问共享数据结构 Protect accesses to shared data structures by temporarily blocking all signals.**
  - 避免可能的破坏 To prevent possible corruption
- **G4: 将全局变量声明为 `volatile` Declare global variables as `volatile`**
  - 避免编译器将其存储在寄存器中 To prevent compiler from storing them in a register
- **G5: 将全局标记声明为 `volatile sig_atomic_t` Declare global flags as `volatile sig_atomic_t`**
  - *flag* 只读或只写的变量 (例如 `flag=1`, 不是 `flag++`) *flag*: variable that is only read or written (e.g. `flag = 1`, not `flag++`)
  - 按照这种方式声明的 `flag` 变量不需要像其他全局变量那样保护 Flag declared this way does not need to be protected like other globals

# 异步信号安全 Async-Signal-Safety



- 如果一个函数是可重入的（例如所有变量存储在栈帧，CS:APP3e 12.7.2）或者不可以被信号打断的则将其称为**异步信号安全***async-signal-safe* Function is *async-signal-safe* if either reentrant (e.g., all variables stored on stack frame, CS:APP3e 12.7.2) or non-interruptible by signals.
- Posix中有117个函数是异步信号安全*async-signal-safe* Posix guarantees 117 functions to be *async-signal-safe*
  - 来源：man命令 Source: “man 7 signal”
  - 在其中的常见函数包括： Popular functions on the list:
    - `_exit, write, wait, waitpid, sleep, kill`
  - 并不在其中得常见函数 Popular functions that are **not** on the list:
    - `printf, sprintf, malloc, exit`
    - 不幸的事实： `write`是唯一异步信号安全*async-signal-safe*输出函数  
Unfortunate fact: `write` is the only *async-signal-safe* output function

# 安全格式化输出：选项#1



## Safe Formatted Output: Option #1

- 在信号处理程序中使用csapp.c的可重入的SIO（安全I/O库）  
Use the reentrant SIO (Safe I/O library) from csapp.c in

your handlers

- `ssize_t sio_puts(char s[]) /* Put string */`
- `ssize_t sio_putl(long v) /* Put long */`
- `void sio_error(char s[]) /* Put msg & exit */`

```
void sigint_handler(int sig) /* Safe SIGINT handler */
{
    sio_puts("So you think you can stop the bomb"
            " with ctrl-c, do you?\n");
    sleep(2);
    sio_puts("Well...");
    sleep(1);
    sio_puts("OK. :-)\n");
    _exit(0);
}
```

sigintsafe.c

# 安全格式化输出：选项#2



## Safe Formatted Output: Option #2

- 使用新的且改进的可重入 `sio_printf`! Use the new & improved reentrant `sio_printf`!
  - 处理 `printf` 受限类的格式串 Handles restricted class of `printf` format strings
    - 识别: Recognizes: `%c %s %d %u %x %%`
    - 大小指定符: Size designators `'l'` and `'z'`

```
void sigint_handler(int sig) /* Safe SIGINT handler */
{
    sio_printf("So you think you can stop the bomb"
               " (process %d) with ctrl-%c, do you?\n",
               (int) getpid(), 'c');

    sleep(2);
    sio_puts("Well...");
    sleep(1);
    sio_puts("OK. :-)\n");
    _exit(0);
}
```

sigintsafe.c

# 正确的信号处理

## Correct Signal Handling

```
volatile int ccount = 0;
void child_handler(int sig) {
    int olderrno = errno;
    pid_t pid;
    if ((pid = wait(NULL)) < 0)
        Sio_error("wait error");
    ccount--;
    sio_puts("Handler reaped child ");
    sio_putl((long)pid);
    sio_puts(" \n");
    sleep(1);
    errno = olderrno;
}

void fork14() {
    pid_t pid[N];
    int i;
    ccount = N;
    signal(SIGCHLD, child_handler);

    for (i = 0; i < N; i++) {
        if ((pid[i] = fork()) == 0) {
            sleep(1);
            exit(0); /* Child exits */
        }
    }
    while (ccount > 0) /* Parent spins */
        ;
}
```

这段代码不正确！  
This code is incorrect!

N == 5

- 挂起的信号是不排队的  
Pending signals are not queued
  - 对每个信号类型，只用一个比特位来标识是否有信号被挂起 For each signal type, one bit indicates whether or not signal is pending...
  - 因此每种最多有一个挂起的信号 ...thus at most one pending signal of any particular type.
- 不可以使用信号对事件计数，例如子进程终止等 You can't use signals

```
whaleshark> ./forks 14
Handler reaped child 23240
Handler reaped child 23241
...(hangs)
```

as

# 正确信号处理 Correct Signal Handling



## ■ 必须等待所有终止的子进程 Must wait for all terminated child processes

- 将wait放入到循环中以回收所有终止的子进程 Put `wait` in a loop to reap all terminated children

```
void child_handler2(int sig)
{
    int olderrno = errno;
    pid_t pid;
    while ((pid = wait(NULL)) > 0) {
        ccount--;
        sio_puts("Handler reaped child ");
        sio_putl((long)pid);
        sio_puts(" \n");
    }
    if (errno != ECHILD)
        sio_error("wait error");
    errno = olderrno;
}
```

```
whaleshark> ./forks 15
Handler reaped child 23246
Handler reaped child 23247
Handler reaped child 23248
Handler reaped child 23249
Handler reaped child 23250
whaleshark>
```



# 可移植的信号处理

## Portable Signal Handling



- 不同的Unix版本有不同的信号处理语义      Ugh! Different versions of Unix can have different signal handling semantics
  - 一些早期的系统在捕获到信号后会恢复默认动作      Some older systems restore action to default after catching signal
  - 有些被中断的系统调用会返回 `errno == EINTR`      Some interrupted system calls can return with `errno == EINTR`
  - 有的系统并不阻塞正在被处理的信号类型      Some systems don't block signals of the type being handled
- 解决方案: `sigaction`      Solution: `sigaction`

```
handler_t *Signal(int signum, handler_t *handler)
{
    struct sigaction action, old_action;

    action.sa_handler = handler;
    sigemptyset(&action.sa_mask); /* Block sigs of type being handled */
    action.sa_flags = SA_RESTART; /* Restart syscalls if possible */

    if (sigaction(signum, &action, &old_action) < 0)
        unix_error("Signal error");
    return (old_action.sa_handler);
}
```

csapp.c



# 同步控制流避免竞争

## Synchronizing Flows to Avoid Races

- 简单shell的SIGCHLD处理程序 SIGCHLD handler for a simple shell
  - 当运行临界代码时阻塞所有信号 Blocks all signals while running critical code

```
void handler(int sig)
{
    int olderrno = errno;
    sigset_t mask_all, prev_all;
    pid_t pid;

    sigfillset(&mask_all);
    while ((pid = waitpid(-1, NULL, 0)) > 0) { /* Reap child */
        sigprocmask(SIG_BLOCK, &mask_all, &prev_all);
        deletejob(pid); /* Delete the child from the job list */
        sigprocmask(SIG_SETMASK, &prev_all, NULL);
    }
    if (pid != 0 && errno != ECHILD)
        sio_error("waitpid error");
    errno = olderrno;
}
```

procmask1.c

# 同步控制流避免竞争



## Synchronizing Flows to Avoid Races

- 简单的shell程序有个不易发现的同步问题，因为其假设父进程先于子进程  
Simple shell with a subtle synchronization error because it assumes parent runs before child

```
int main(int argc, char **argv)
{
    int pid;
    sigset_t mask_all, prev_all;
    int n = N; /* N = 5 */
    sigfillset(&mask_all);
    signal(SIGCHLD, handler);
    initjobs(); /* Initialize the job list */

    while (n-- > 0) {
        if ((pid = fork()) == 0) { /* Child */
            execve("/bin/date", argv, NULL);
        }
        sigprocmask(SIG_BLOCK, &mask_all, &prev_all); /* Parent */
        addjob(pid); /* Add the child to the job list */
        sigprocmask(SIG_SETMASK, &prev_all, NULL);
    }
    exit(0);
}
```

procmask1.c

# 没有竞争问题的修正shell程序

## Corrected Shell Program Without Race



```
int main(int argc, char **argv)
{
    int pid;
    sigset_t mask_all, mask_one, prev_one;
    int n = N; /* N = 5 */
    sigfillset(&mask_all);
    sigemptyset(&mask_one);
    sigaddset(&mask_one, SIGCHLD);
    signal(SIGCHLD, handler);
    initjobs(); /* Initialize the job list */

    while (n--) {
        sigprocmask(SIG_BLOCK, &mask_one, &prev_one); /* Block SIGCHLD */
        if ((pid = fork()) == 0) { /* Child process */
            sigprocmask(SIG_SETMASK, &prev_one, NULL); /* Unblock SIGCHLD */
            execve("/bin/date", argv, NULL);
        }
        sigprocmask(SIG_BLOCK, &mask_all, NULL); /* Parent process */
        addjob(pid); /* Add the child to the job list */
        sigprocmask(SIG_SETMASK, &prev_one, NULL); /* Unblock SIGCHLD */
    }
    exit(0);
}
```

# 显式等待信号

## Explicitly Waiting for Signals



- 信号处理程序显式等待SIGCHLD信号的到来 Handlers for program explicitly waiting for SIGCHLD to arrive

```
volatile sig_atomic_t pid;

void sigchld_handler(int s)
{
    int olderrno = errno;
    pid = waitpid(-1, NULL, 0); /* Main is waiting for nonzero pid */
    errno = olderrno;
}

void sigint_handler(int s)
{
}
```

waitforsignal.c

# 显式等待信号 Explicitly Waiting for Signals



类似于shell等待一个前台的作业终止  
Similar to a shell waiting  
for a foreground job to terminate.

```
int main(int argc, char **argv) {
    sigset_t mask, prev;
    int n = N; /* N = 10 */
    signal(SIGCHLD, sigchld_handler);
    signal(SIGINT, sigint_handler);
    sigemptyset(&mask);
    sigaddset(&mask, SIGCHLD);

    while (n--) {
        sigprocmask(SIG_BLOCK, &mask, &prev); /* Block SIGCHLD */
        if (fork() == 0) /* Child */
            exit(0);
        /* Parent */
        pid = 0;
        sigprocmask(SIG_SETMASK, &prev, NULL); /* Unblock SIGCHLD */

        /* Wait for SIGCHLD to be received (wasteful!) */
        while (!pid)
            ;
        /* Do some work after receiving SIGCHLD */
        printf(".");
    }
    printf("\n");
    exit(0);
}
```

waitforsignal.c

# 显式等待信号 Explicitly Waiting for Signals



```
while (!pid)
    ;
```

- 程序是对的，但是太浪费资源 Program is correct, but very wasteful

- 程序忙于等待循环 Program in busy-wait loop

```
while (!pid) /* Race! */
    pause();
```

- 可能存在竞争 Possible race condition

- 在检查pid和开始暂停之间，可能接收信号 Between checking pid and starting pause, might receive signal

```
while (!pid) /* Too slow! */
    sleep(1);
```

- 安全，但是很慢 Safe, but slow

- 会占用1秒钟才能响应 Will take up to one second to respond

# 使用sigsuspend等待信号

## Waiting for Signals with sigsuspend



- `int sigsuspend(const sigset_t *mask)`
- 等价于原子版本（无中断可能）的： **Equivalent to atomic (uninterruptable) version of:**

```
sigprocmask(SIG_SETMASK, &mask, &prev);  
pause();  
sigprocmask(SIG_SETMASK, &prev, NULL);
```



# 使用sigsuspend等待信号

## Waiting for Signals with sigsuspend



```
int main(int argc, char **argv) {
    sigset_t mask, prev;
    int n = N; /* N = 10 */
    signal(SIGCHLD, sigchld_handler);
    signal(SIGINT, sigint_handler);
    sigemptyset(&mask);
    sigaddset(&mask, SIGCHLD);
    while (n--) {
        sigprocmask(SIG_BLOCK, &mask, &prev); /* Block SIGCHLD */
        if (fork() == 0) /* Child */
            exit(0);

        /* Wait for SIGCHLD to be received */
        pid = 0;
        while (!pid)
            sigsuspend(&prev);
        /* Optionally unblock SIGCHLD */
        sigprocmask(SIG_SETMASK, &prev, NULL);
        /* Do some work after receiving SIGCHLD */
        printf(".");
    }
    printf("\n");
    exit(0);
}
```



# 议题

- 外壳 Shells
- 信号 Signals
- 非局部跳转 **Nonlocal jumps**
  - 参见教材和附加的幻灯片 Consult your textbook and additional slides



# 总结 Summary

- **信号提供进程级异常处理 Signals provide process-level exception handling**
  - 可以从用户程序产生 Can generate from user programs
  - 可以声明信号处理程序定义处理效果 Can define effect by declaring signal handler
  - 编写信号处理函数的时候要特别小心 Be very careful when writing signal handlers
  
- **非局部跳转给出了进程内部的异常控制流 Nonlocal jumps provide exceptional control flow within process**
  - 遵守栈相关的原则 Within constraints of stack discipline

# 附加的幻灯片 **Additional slides**





# 非局部跳转

## Nonlocal Jumps: `setjmp/longjmp`

- 将控制转移到任意位置的强大（但比较危险）用户级机制 **Powerful (but dangerous) user-level mechanism for transferring control to an arbitrary location**
  - 受控的打破call/return规则的方式 Controlled way to break the procedure call / return discipline
  - 通常用于错误恢复和信号处理 Useful for error recovery and signal handling
- `int setjmp(jmp_buf j)`
  - 必须在longjmp之前调用 Must be called before longjmp
  - 给出后续longjmp对应的返回位置 Identifies a return site for a subsequent longjmp
  - 一次调用，返回一次或者多次 Called **once**, returns **one or more** times
- 实现 **Implementation:**
  - 通过将当前寄存器上下文、栈指针和PC值存储在jmp\_buf中记住当前位置 Remember where you are by storing the current **register context**, **stack pointer**, and **PC value** in `jmp_buf`
  - 返回0 Return 0



# setjmp/longjmp (续 cont)

## ■ void longjmp(jmp\_buf j, int i)

### ■ 含义 Meaning:

- 从setjmp返回, 再次被跳转缓冲区j记住 return from the **setjmp** remembered by jump buffer **j** again ...
- 这次返回i而不是0 ... this time returning **i** instead of 0

### ■ setjmp之后调用 Called after **setjmp**

- 一次调用但是从不返回 Called **once**, but **never** returns

## ■ longjmp实现 longjmp Implementation:

- 从跳转缓冲区j中恢复寄存器上下文 (栈指针、基指针、PC值)  
Restore register context (stack pointer, base pointer, PC value) from jump buffer **j**
- 将返回值寄存器%eax设置为i Set **%eax** (the return value) to **i**
- 跳转到跳转缓冲j中PC指定的位置 Jump to the location indicated by the PC stored in jump buf **j**



# setjmp/longjmp Example 示例

- 目标：从深度嵌套的函数直接返回最开始的调用者
- Goal: return directly to original caller from a deeply-nested function

```
/* Deeply nested function foo */  
void foo(void)  
{  
    if (error1)  
        longjmp(buf, 1);  
    bar();  
}  
  
void bar(void)  
{  
    if (error2)  
        longjmp(buf, 2);  
}
```



```
jmp_buf buf;

int error1 = 0;
int error2 = 1;

void foo(void), bar(void);

int main()
{
    switch(setjmp(buf)) {
        case 0:
            foo();
            break;
        case 1:
            printf("Detected an error1 condition in foo\n");
            break;
        case 2:
            printf("Detected an error2 condition in foo\n");
            break;
        default:
            printf("Unknown error condition in foo\n");
    }
    exit(0);
}
```

## setjmp/longjmp 示例/Example (续/cont)



# 非局部跳转的限制

## Limitations of Nonlocal Jumps



### ■ 基于栈原理工作 Works within stack discipline

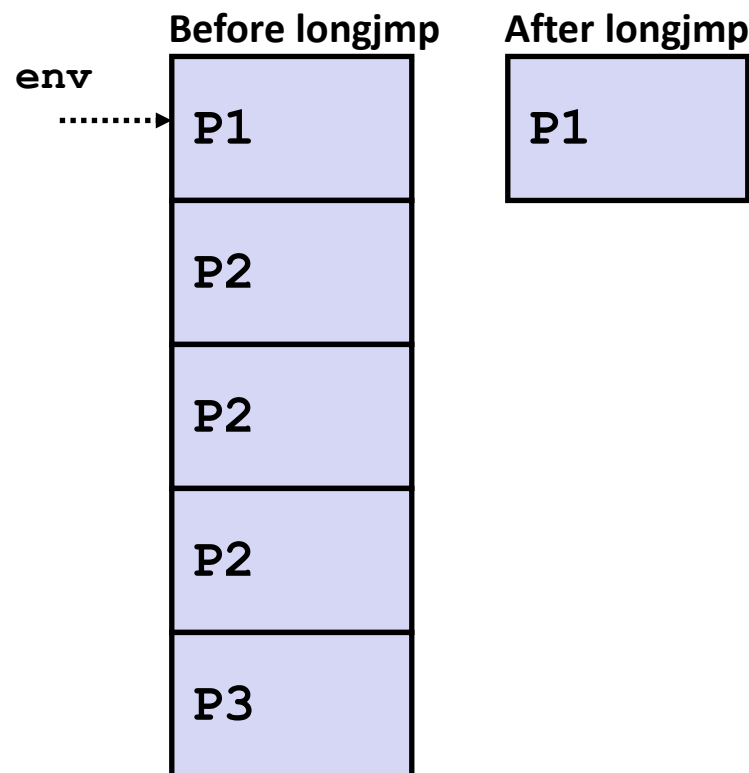
- 只能跳转到已经调用但是还没有完成的函数 Can only long jump to environment of function that has been called but not yet completed

```
jmp_buf env;

P1()
{
    if (setjmp(env)) {
        /* Long Jump to here */
    } else {
        P2();
    }
}

P2()
{
    . . . P2(); . . . P3();
}

P3()
{
    longjmp(env, 1);
}
```








# 非局部跳转的限制（续）

## Limitations of Long Jumps (cont.)

### ■ 基于栈原理工作 Works within stack discipline

- 只能跳转到已经调用但是还没有完成的函数/Can only long jump to environment of function that has been called but not yet completed

```
jmp_buf env;  
  
P1()  
{  
    P2(); P3();  
}  
  
P2()  
{  
    if (setjmp(env)) {  
        /* Long Jump to here */  
    }  
}  
  
P3()  
{  
    longjmp(env, 1);  
}
```



# 整合在一起：程序在按下ctrl-c或d时重启



## Putting It All Together: A Program That Restarts Itself When ctrl-c'd

```
#include "csapp.h"

sigjmp_buf buf;

void handler(int sig)
{
    siglongjmp(buf, 1);
}

int main()
{
    if (!sigsetjmp(buf, 1)) {
        Signal(SIGINT, handler);
        Sio_puts("starting\n");
    }
    else
        Sio_puts("restarting\n");

    while(1) {
        Sleep(1);
        Sio_puts("processing...\n");
    }
    exit(0); /* Control never reaches here */
}
```

```
greatwhite> ./restart
starting
processing...
processing...
processing...
restarting ← Ctrl-c
processing...
processing...
restarting ← Ctrl-c
processing...
processing...
processing...
```

restart.c