



# 第9章 虚拟内存

Dynamic Memory Allocation:

Basic Concepts

动态内存分配: 基本概念

100076202: 计算机系统导论

任课教师:

宿红毅    张艳    黎有琦    李秀星

原作者:

Randal E. Bryant and David R. O'Hallaron



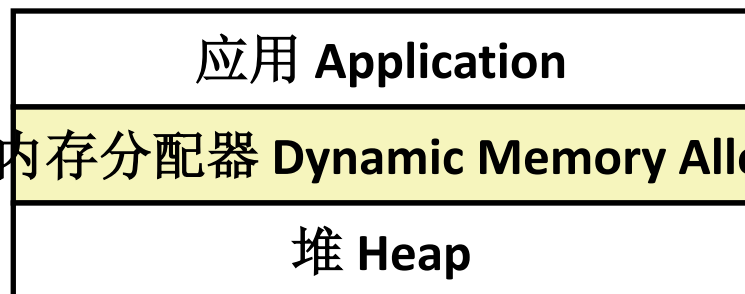
Carnegie  
Mellon  
University



# 议题 Today

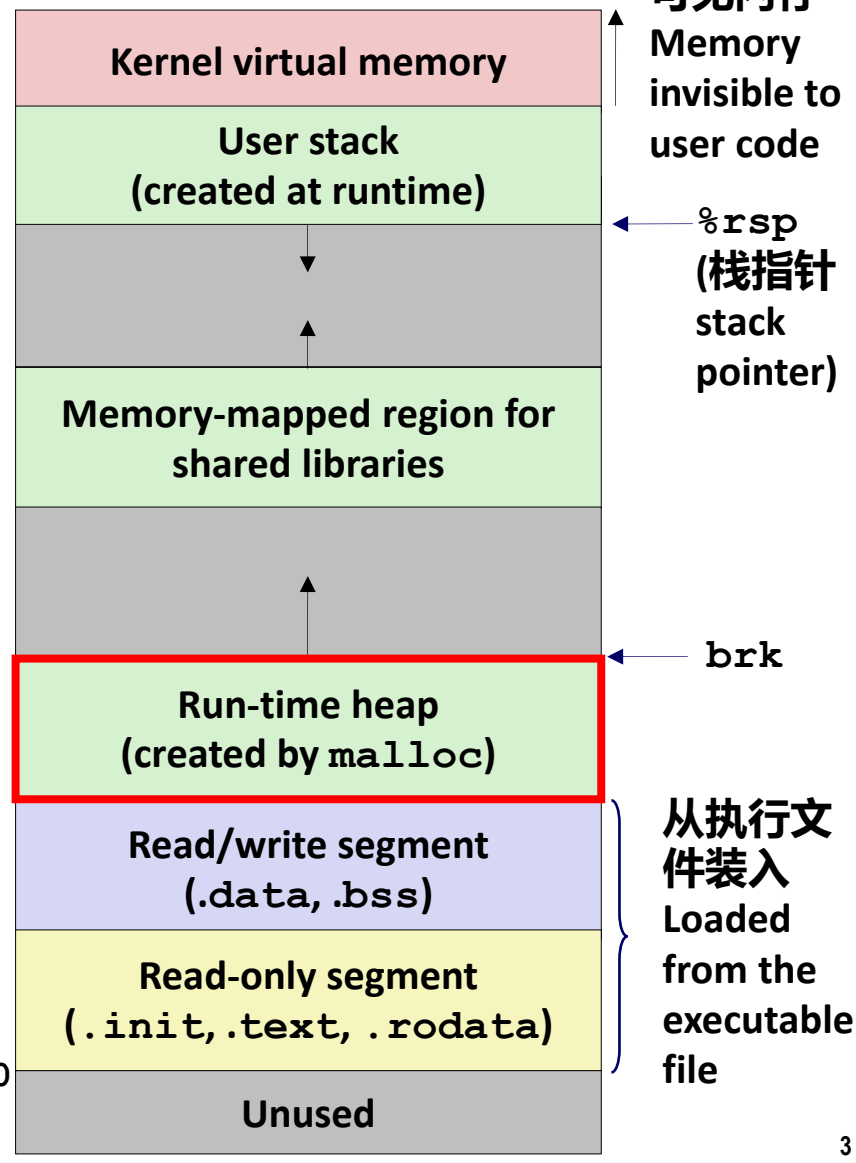
- **基本概念** Basic concepts
- 隐式空闲链表 Implicit free lists

# 动态内存分配 Dynamic Memory Allocation



动态内存分配器 Dynamic Memory Allocator

- 程序员使用**动态内存分配器** (malloc) 在运行时申请虚拟内存  
Programmers use **dynamic memory allocators** (such as malloc) to acquire virtual memory (VM) at runtime
  - 对于那些数据结构大小在运行时才能知道的数据结构 For data structures whose size is only known at runtime
- 动态内存分配器管理进程虚拟内存中一个称为**堆**的区域  
Dynamic memory allocators manage an area of process VM known as the **heap**



# 动态内存分配 Dynamic Memory Allocation



- 分配器将堆当做不同大小的**块**的集合进行管理，不是**已分配**就是**空闲** Allocator maintains heap as collection of variable sized **blocks**, which are either **allocated** or **free**
- 分配器类型 Types of allocators
  - **显式分配器**: 应用程序分配和释放空间 **Explicit allocator**: application allocates and frees space
    - 例如C中的malloc和free E.g., malloc and free in C
  - **隐式分配器**: 应用只负责分配但是不释放空间 **Implicit allocator**: application allocates, but does not free space
    - 例如Java、ML和Lisp中的垃圾收集 E.g. garbage collection in Java, ML, and Lisp
- 今天主要讨论简单的显式内存分配 Will discuss simple explicit memory allocation today



# malloc包 The malloc Package

```
#include <stdlib.h>
```

```
void *malloc(size_t size)
```

- 成功 Successful:
  - 返回大小至少是size的内存块指针，x86上是按8字节对齐，x86-64是按16字节对齐  
Returns a pointer to a memory block of at least **size** bytes aligned to an 8-byte (x86) or 16-byte (x86-64) boundary
  - 如果size为0，则返回NULL    If **size == 0**, returns NULL
- 不成功：返回NULL并设置errno    Unsuccessful: returns NULL (0) and sets **errno**

```
void free(void *p)
```

- 将p指向的内存块返回给可用内存池    Returns the block pointed at by **p** to pool of available memory
- p必须是之前调用malloc或者realloc获得的    **p** must come from a previous call to **malloc** or **realloc**

## 其他函数 Other functions

- **calloc**: malloc的另一个版本，会将分配的内存块初始化为0    Version of **malloc** that initializes allocated block to zero.
- **realloc**: 改变之前分配的块的大小    Changes the size of a previously allocated block.
- **sbrk**: 分配器内部用来增加或者减小堆的大小    Used internally by allocators to grow or shrink the heap



# malloc示例 malloc Example

```
#include <stdio.h>
#include <stdlib.h>

void foo(int n) {
    int i, *p;

    /* Allocate a block of n ints */
    p = (int *) malloc(n * sizeof(int));
    if (p == NULL) {
        perror("malloc");
        exit(0);
    }

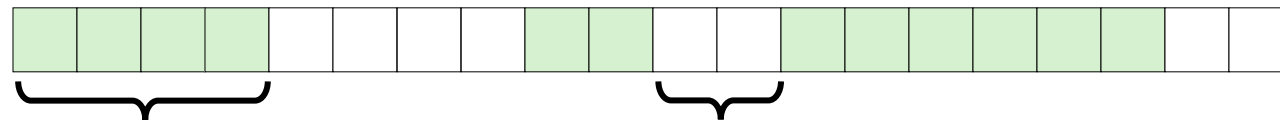
    /* Initialize allocated block */
    for (i=0; i<n; i++)
        p[i] = i;

    /* Return allocated block to the heap */
    free(p);
}
```



# 可视化展示规则 Visualization Conventions

- 显式8字节字为一个方块 Show 8-byte words as squares
- 分配采用双字对齐 Allocations are double-word aligned



Allocated block  
(4 words)

Free block  
(2 words)



空闲字 Free word



已分配字 Allocated word

# 分配示例 Allocation Example

## (概念上 Conceptual)

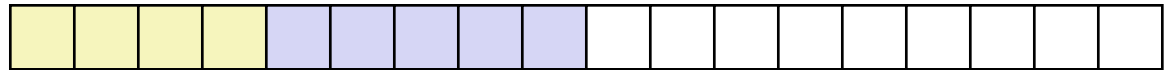


```
#define SIZ sizeof(size_t)
```

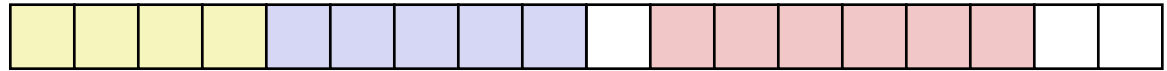
```
p1 = malloc(4*SIZ)
```



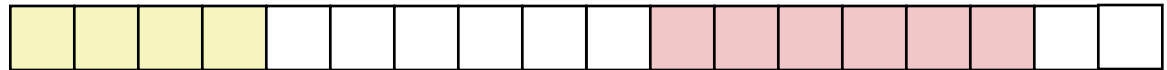
```
p2 = malloc(5*SIZ)
```



```
p3 = malloc(6*SIZ)
```



```
free(p2)
```



```
p4 = malloc(2*SIZ)
```







# 限制 Constraints

## ■ 应用 Applications

- 可以发出任意malloc和free请求序列 Can issue arbitrary sequence of **malloc** and **free** requests
- **free**请求必须针对一个malloc请求的块 **free** request must be to a **malloc**'d block

## ■ 显式分配器 Explicit Allocators

- 无法控制分配的块的数量和大小 Can't control number or size of allocated blocks
- 必须及时响应malloc请求 Must respond immediately to **malloc** requests
  - 例如, 不能对请求排序和缓冲 *i.e.*, can't reorder or buffer requests
- 必须从空闲空间分配内存块 Must allocate blocks from free memory
  - 例如, 分配的块必须在空闲内存中 *i.e.*, can only place allocated blocks in free memory
- 必须按照需求实现块对齐 Must align blocks so they satisfy all alignment requirements
  - Linux中x86是8字节对齐, x86-64是16字节对齐 8-byte (x86) or 16-byte (x86-64) alignment on Linux boxes
- 只能操作和修改空闲内存 Can manipulate and modify only free memory
- 一旦分配后不能移动内存块 Can't move the allocated blocks once they are **malloc**'d
  - 例如, 压缩是不允许的 *i.e.*, compaction is not allowed



# 性能目标：吞吐率 Performance Goal: Throughput

- 对于给定的malloc和free序列 Given some sequence of `malloc` and `free` requests:
  - $R_0, R_1, \dots, R_k, \dots, R_{n-1}$
- 目标：最大化吞吐率和峰值内存利用率 Goals: maximize throughput and peak memory utilization
  - 这些目标通常是互相冲突的 These goals are often conflicting
- 吞吐率 Throughput:
  - 单位时间内完成的请求数量 Number of completed requests per unit time
  - 例如： Example:
    - 10秒内完成5000次malloc和5000次free 5,000 `malloc` calls and 5,000 `free` calls in 10 seconds
    - 吞吐率就是1000次操作/秒 Throughput is 1,000 operations/second



# 性能目标：最小化开销

## Performance Goal: Minimize Overhead

- 对于给定的malloc和free某个请求序列 Given some sequence of malloc and free requests:
  - $R_0, R_1, \dots, R_k, \dots, R_{n-1}$
- *K次请求之后，我们得到： After k requests we have:*
- **定义：总有效载荷** **Def: Aggregate payload  $P_k$** 
  - `malloc(p)` 返回一个载荷为p字节的块 `malloc(p)` results in a block with a **payload** of p bytes
  - 请求 $R_k$ 完成后，总有效载荷 $P_k$ 是目前已分配的载荷的总大小 After request  $R_k$  has completed, the **aggregate payload**  $P_k$  is the sum of currently allocated payloads
- **定义：当前堆大小 $H_k$**  **Def: Current heap size  $H_k$** 
  - 假设 $H_k$ 单调不递减 Assume  $H_k$  is monotonically nondecreasing
    - 即当分配器使用sbrk时堆增加 i.e., heap only grows when allocator uses `sbrk`
- **定义： $k+1$ 次请求之后峰值内存利用率** **Def: Peak memory utilization after  $k+1$  requests**
  - $U_k = (\max_{i \leq k} P_i) / H_k$

# 性能目标：最小化开销

## Performance Goal: Minimize Overhead



- 对于给定的malloc和free一些请求序列 Given some sequence of malloc and free requests:
  - $R_0, R_1, \dots, R_k, \dots, R_{n-1}$
- $K$ 次请求之后，我们得到： After  $k$  requests we have:
- **定义：总有效载荷** **Def: Aggregate payload  $P_k$** 
  - `malloc(p)` 返回一个**载荷**为 $p$ 字节的块 / `malloc(p)` results in a block with a **payload** of  $p$  bytes
  - **总有效载荷** $P_k$ 是目前已分配的载荷的总和 The **aggregate payload**  $P_k$  is the sum of currently allocated payloads
  - **峰值总有效载荷**是请求序列中任何点处最大总有效载荷 The **peak aggregate payload**  $\max_{i \leq k} P_i$  is the maximum aggregate payload at any point in the sequence up to request

# 性能目标：最小化开销

## Performance Goal: Minimize Overhead



- 对于给定的malloc和free一些请求序列 Given some sequence of malloc and free requests:
  - $R_0, R_1, \dots, R_k, \dots, R_{n-1}$
- $K$ 次请求之后，我们得到： After  $k$  requests we have:
- **定义：** 当前堆大小  $H_k$  **Def: Current heap size  $H_k$** 
  - 假设当分配器使用sbrk时堆仅增加，从不收缩 Assume heap only grows when allocator uses **sbrk**, never shrinks
- **定义：** 开销，  $O_k$  **Def: Overhead,  $O_k$** 
  - 堆空间没有为程序数据使用的比例 Fraction of heap space NOT used for program data
  - $O_k = (H_k / \max_{i \leq k} P_i) - 1.0$

# 基准测试示例 Benchmark Example



## ■ 基准测试 Benchmark

### syn-array-short

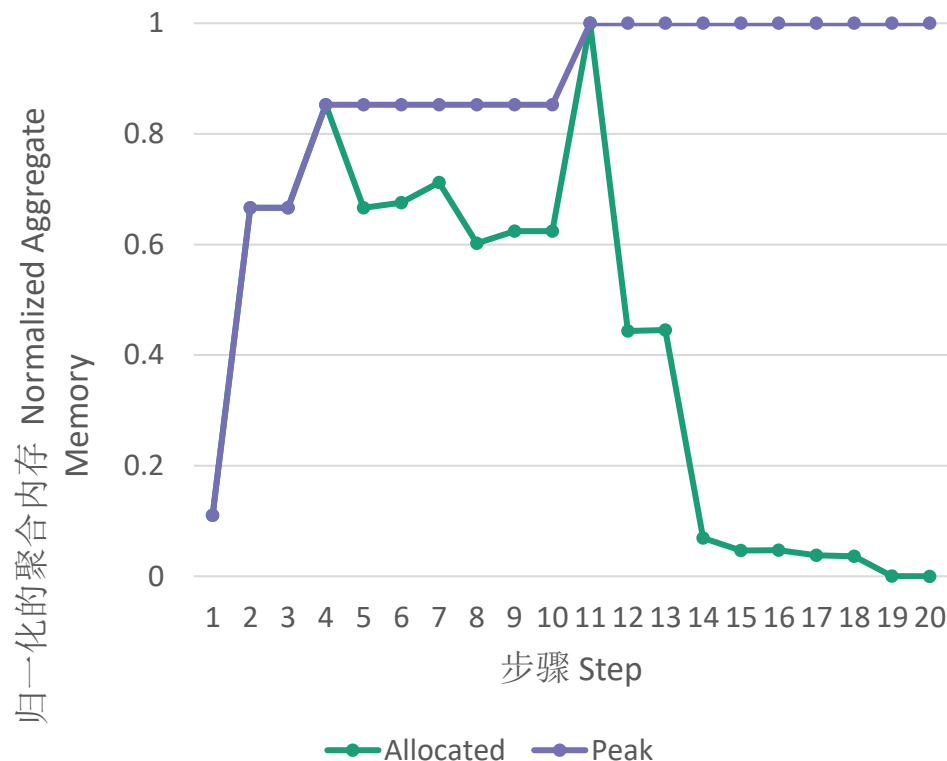
- malloc实验提供的跟踪 Trace provided with malloc lab
- 分配和释放各10个块 Allocate & free 10 blocks
- a代表分配 a = allocate
- f代表释放 f = free
- 偏置在开始时分配，在结束时释放 Bias toward allocate at beginning & free at end
- 块号1-10 Blocks number 1–10
- 已分配：所有分配量的和  
Allocated: Sum of all allocated amounts
- 峰值：曾经分配的最大值  
Peak: Max so far of Allocated

步骤 Step	命令 Command	偏置 Delta	已分配 Allocated	峰值 Peak
1	a 0 9904	9904	9904	9904
2	a 1 50084	50084	59988	59988
3	a 2 20	20	60008	60008
4	a 3 16784	16784	76792	76792
5	f 3	-16784	60008	76792
6	a 4 840	840	60848	76792
7	a 5 3244	3244	64092	76792
8	f 0	-9904	54188	76792
9	a 6 2012	2012	56200	76792
10	f 2	-20	56180	76792
11	a 7 33856	33856	90036	90036
12	f 1	-50084	39952	90036
13	a 8 136	136	40088	90036
14	f 7	-33856	6232	90036
15	f 6	-2012	4220	90036
16	a 9 20	20	4240	90036
17	f 4	-840	3400	90036
18	f 8	-136	3264	90036
19	f 5	-3244	20	90036
20	f 9	-20	0	90036

# 基准测试可视化 Benchmark Visualization



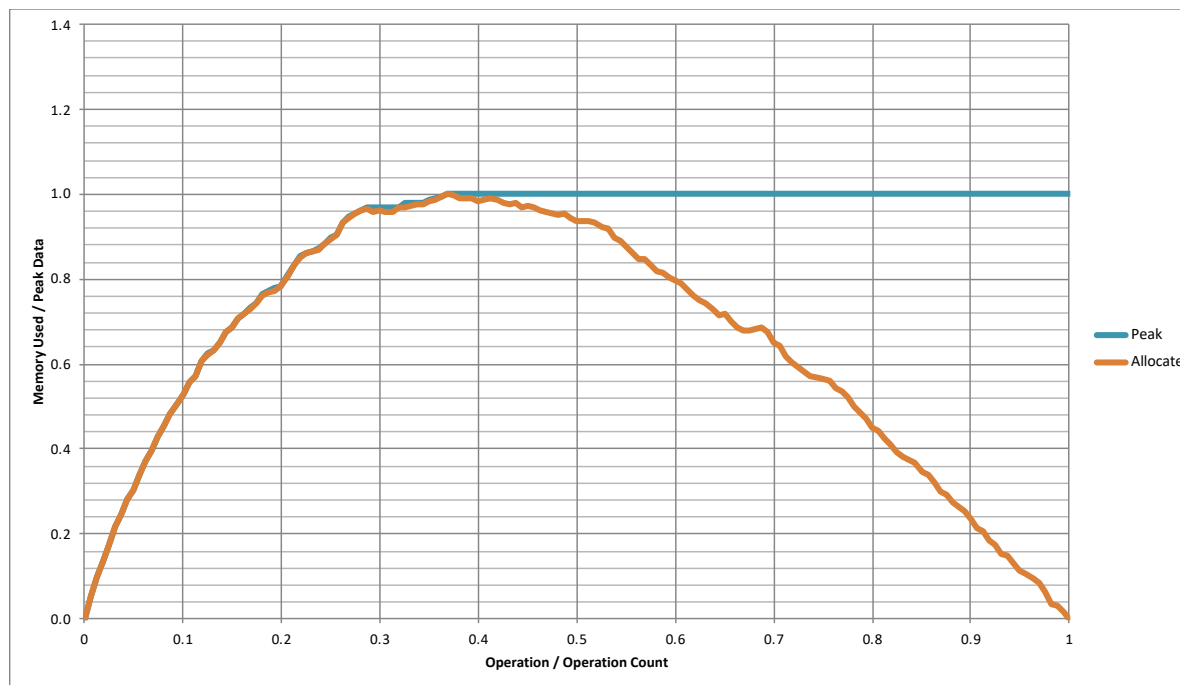
步骤 Step	命令 Command	偏置 Delta	已分配 Allocated	峰值 Peak
1	a 0 9904	9904	9904	9904
2	a 1 50084	50084	59988	59988
3	a 2 20	20	60008	60008
4	a 3 16784	16784	76792	76792
5	f 3	-16784	60008	76792
6	a 4 840	840	60848	76792
7	a 5 3244	3244	64092	76792
8	f 0	-9904	54188	76792
9	a 6 2012	2012	56200	76792
10	f 2	-20	56180	76792
11	a 7 33856	33856	90036	90036
12	f 1	-50084	39952	90036
13	a 8 136	136	40088	90036
14	f 7	-33856	6232	90036
15	f 6	-2012	4220	90036
16	a 9 20	20	4240	90036
17	f 4	-840	3400	90036
18	f 8	-136	3264	90036
19	f 5	-3244	20	90036
20	f 9	-20	0	90036



- 已分配内存和峰值内存是步骤 $k$ 的函数绘图 Plot  $P_k$  (allocated) and  $\max_{i \leq k} P_k$  (peak) as a function of  $k$  (step)
- Y轴归一化处理—占最大值的比例 Y-axis normalized — fraction of maximum

# 典型的基准测试行为

## Typical Benchmark Behavior



- 分配和释放内存的长序列（40000块） Longer sequence of mallocs & frees (40,000 blocks)
  - 开始都是分配内存，然后转向释放内存 Starts with all mallocs, and shifts toward all frees
- 分配器必须整个时间段内有效管理空间 Allocator must manage space efficiently the whole time
- 生产分配器可以收缩堆 Production allocators can shrink the heap





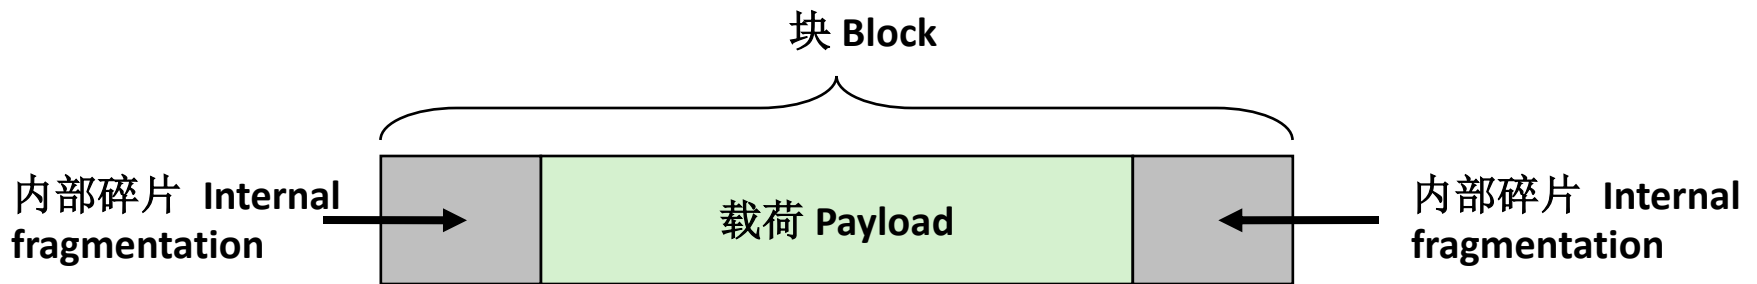
# 内存碎片 Fragmentation

- 由内存**碎片**导致的内存低利用率 **Poor memory utilization caused by *fragmentation***
  - **内部碎片** *internal* fragmentation
  - **外部碎片** *external* fragmentation



# 内部碎片 Internal Fragmentation

- 对于给定的块，如果载荷小于块大小就会导致**内部碎片** For a given block, **internal fragmentation** occurs if payload is smaller than block size

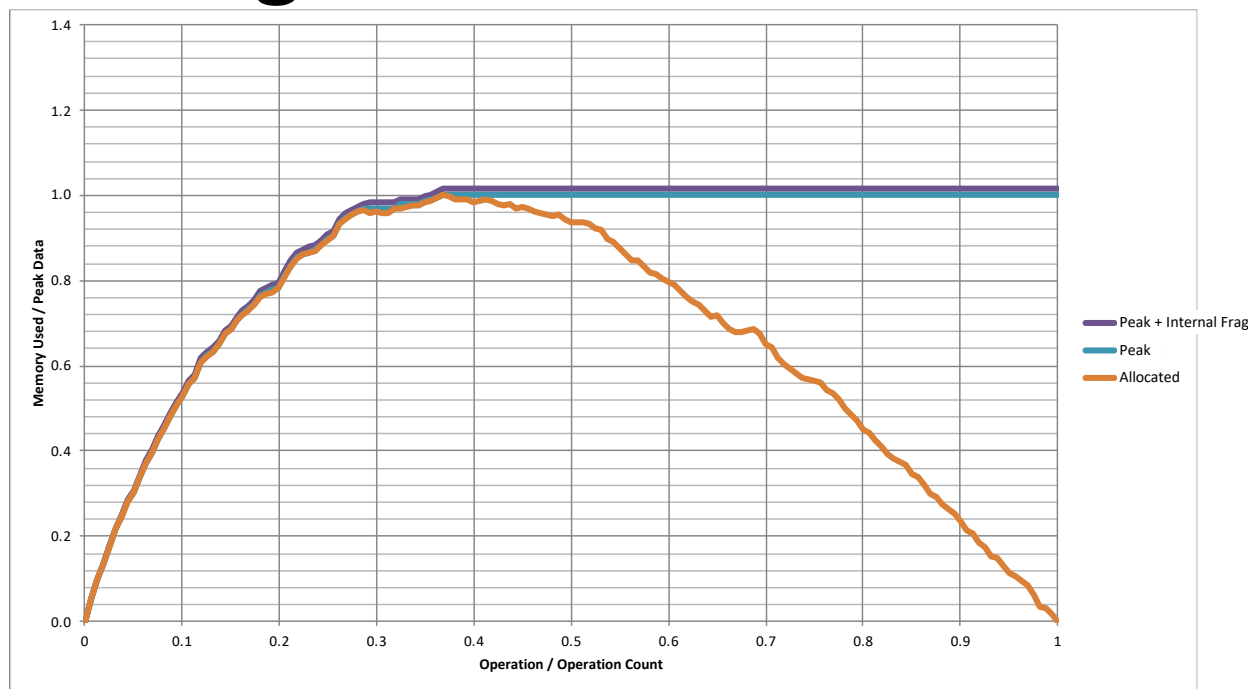


- 原因 **Caused by**
  - 维护堆数据结构开销 Overhead of maintaining heap data structures
  - 为了对齐填充的部分 Padding for alignment purposes
  - 显式策略导致 Explicit policy decisions  
(例如：为了满足一个小的请求返回一个大的块 e.g., to return a big block to satisfy a small request)
- 只是与**之前**的请求的模式相关 **Depends only on the pattern of *previous* requests**
  - 因此易于度量 Thus, easy to measure



# 内部碎片效应

## Internal Fragmentation Effect



- 紫色线条：由于分配器的数据+对齐填充，堆大小增加 **Purple line: additional heap size due to allocator's data + padding for alignment**

- 对于该基准，1.5%的开销 For this benchmark, 1.5% overhead
- 无法在实践中实现 Cannot achieve in practice
- 特别是因为无法移动已分配的块 Especially since cannot move allocated blocks

# 外部碎片 External Fragmentation



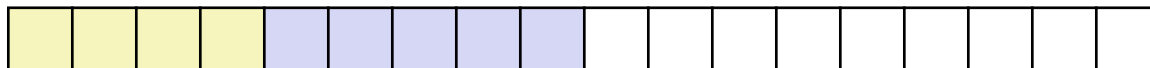
```
#define SIZ sizeof(size_t)
```

- 当有足够的聚合堆内存，但是没有单一的空闲块足够大时产生外部碎片 Occurs when there is enough aggregate heap memory, but no single free block is large enough

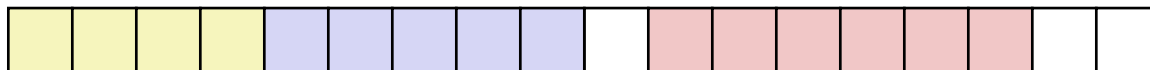
```
p1 = malloc(4*SIZ)
```



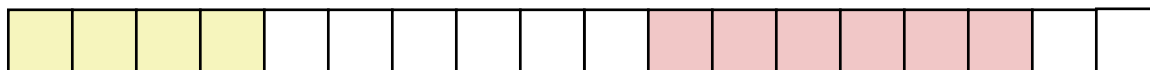
```
p2 = malloc(5*SIZ)
```



```
p3 = malloc(6*SIZ)
```



```
free(p2)
```



```
p4 = malloc(7*SIZ)
```

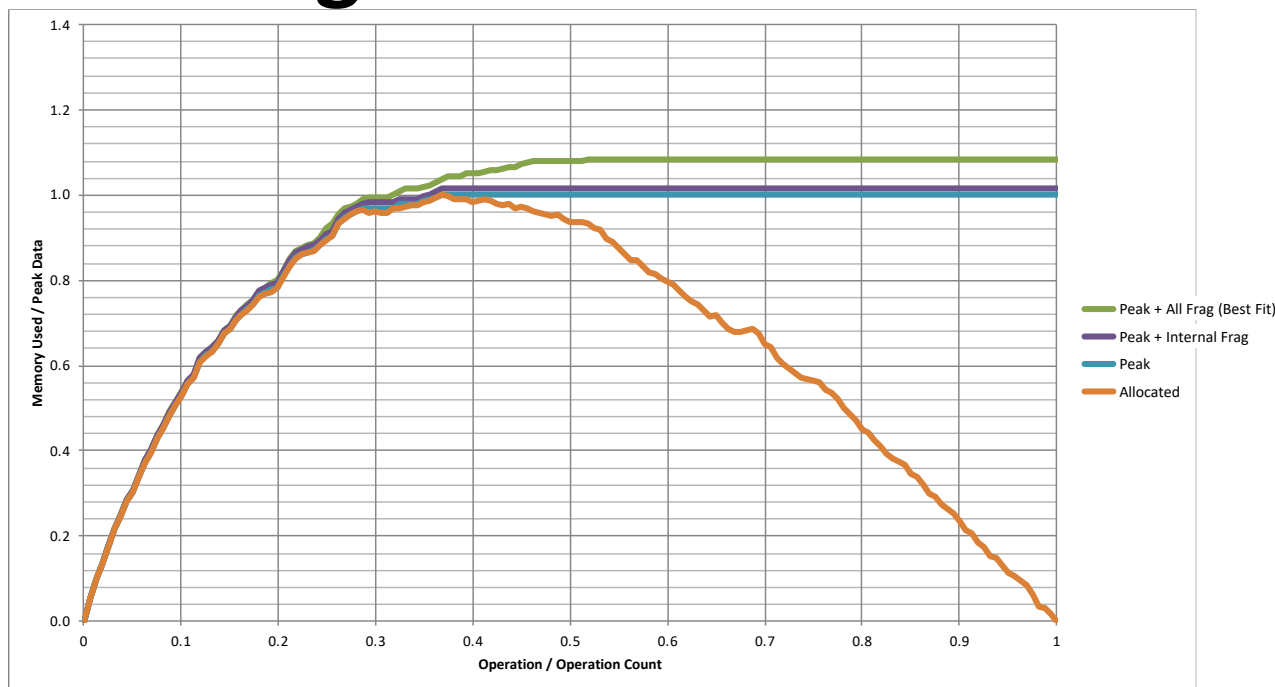
*诶呀（现在会发生什么？）*

*Yikes! (what would happen now?)*

- 取决于未来请求的模式 Depends on the pattern of future requests
  - 因此，难以测量 Thus, difficult to measure

# 外部碎片的效应

## External Fragmentation Effect



- **绿线：由于外部碎片导致的额外堆大小** **Green line: additional heap size due to external fragmentation**
- **最佳匹配：一种分配策略** **Best Fit: One allocation strategy**
  - (稍后讨论) (To be discussed later)
  - 总开销=本基准的8.3% Total overhead = 8.3% on this benchmark



# 实现问题 Implementation Issues

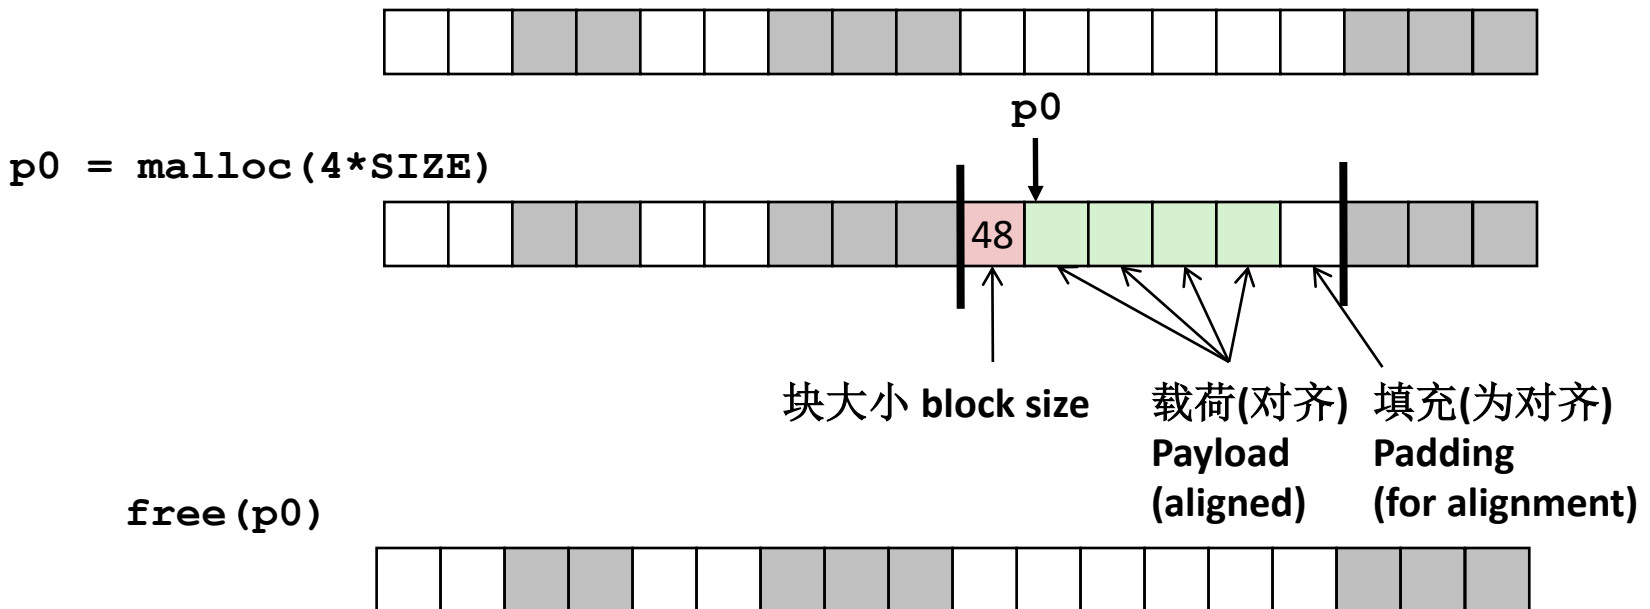
- 给定一个指针，我们怎么知道要释放多大的空间 **How do we know how much memory to free given just a pointer?**
- 我们怎么跟踪空闲块 **How do we keep track of the free blocks?**
- 当分配的结构大小小于选择的空闲块时怎么办？ **What do we do with the extra space when allocating a structure that is smaller than the free block it is placed in?**
- 当有多个块可用时我们应该怎么选？ **How do we pick a block to use for allocation -- many might fit?**
- 如何再次插入空闲块？ **How do we reinsert freed block?**

# 获取释放大小 Knowing How Much to Free



## ■ 标准方法 Standard method

- 在块之前的字中保存块长度 Keep the length (in bytes) of a block in the word *preceding* the block.
  - 包括头部 Including the header
  - 这个字称为头部域或者头部 This word is often called the **header field** or **header**
- 每个分配的块需要一个额外的字 Requires an extra word for every allocated block



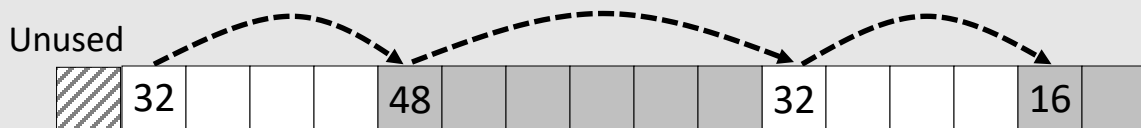
# 跟踪空闲块 Keeping Track of Free Blocks



- 方法1: **隐式链表**-使用长度链接所有块 Method 1: **Implicit list** using length—links all blocks

需要每个块标记为已分配/空闲

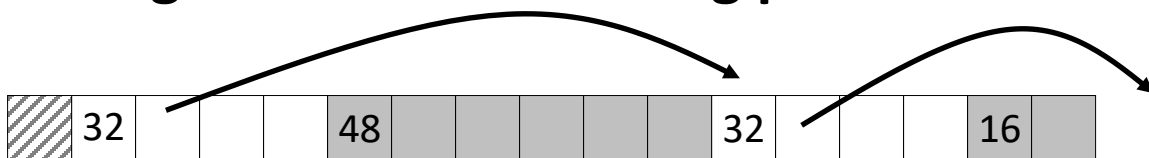
Need to tag each block as allocated/free



- 方法2: 空闲块之间使用指针的**显式链表** Method 2: **Explicit list** among the free blocks using pointers

指针需要占空间

Need space for pointers



- 方法3: **分离的空闲链表** Method 3: **Segregated free list**
  - 不同大小块使用不同的空闲链表 Different free lists for different size classes
- 方法4: **根据大小对块排序** Method 4: **Blocks sorted by size**
  - 可以使用一个平衡树（红黑树），每个空闲块内有指针和做为键值的长度 Can use a balanced tree (e.g., Red-Black tree) with pointers within each free block, and the length used as a key





# 议题 Today

- 基本概念 Basic concepts
- **隐式空闲链表** Implicit free lists

# 方法1：隐式空闲链表 Method 1: Implicit Free List



- 对每个块都需要大小和分配的状态 For each block we need both size and allocation status
  - 可以放在两个字中:浪费 Could store this information in two words: wasteful!
- 标准技巧 Standard trick
  - 如果块是对齐的, 则地址低位部分总是0 If blocks are aligned, some low-order address bits are always 0
  - 与其存储0, 还不如将其作为已分配/空闲的标志位 Instead of storing an always-0 bit, use it as a allocated/free flag
  - 读块大小那个字时需要将这些位屏蔽掉 When reading size word, must mask out this bit

1个字 1 word



a = 1: Allocated block 分配的块

a = 0: Free block 空闲块

Size: block size 块大小

Payload: application data 载荷: 应用数据  
(仅已分配的块)  
(allocated blocks only)

已分配和空闲块格式  
Format of  
allocated and  
free blocks



# 隐式链表：数据结构

## Implicit List: Data Structures



### ■ 块声明 Block declaration

头部 header	有效载荷 payload
-----------	--------------

```
typedef uint64_t word_t;

typedef struct block
{
    word_t header;
    unsigned char payload[0];           // Zero length array
} block_t;
```

### ■ 从块指针获得有效载荷 Getting payload from block pointer

```
return (void *) (block->payload);
```

### ■ 从有效载荷获得头部 Getting header from payload

```
return (block_t *) ((unsigned char *) bp
                    - offsetof(block_t, payload));
```

C语言函数`offsetof(struct, member)`返回`member`在`struct`中的偏移

C function `offsetof(struct, member)` returns offset of member within struct

# 隐式链表：访问头部

## Implicit List: Header access



大小 Size	a
---------	---

- 从头部获得分配位 Getting allocated bit from header

```
return header & 0x1;
```

- 从头部获得块大小 Getting size from header

```
return header & ~0xfL;
```

- 初始化头部 Initializing header // block\_t \*block

```
block->header = size | alloc;
```



# 隐式链表：遍历链表

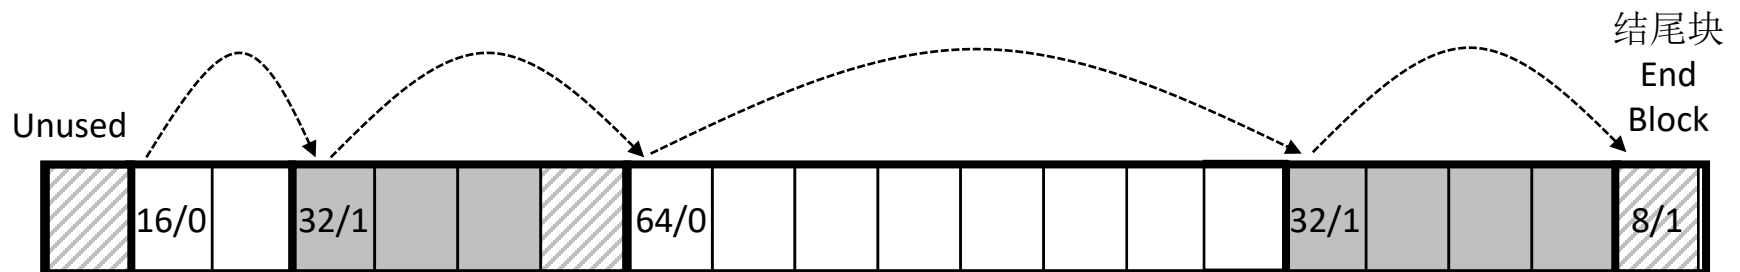
## Implicit List: Traversing list



————— 块大小 block size —————>

### ■ 查找下一个块 Find next block

```
static block_t *find_next(block_t *block)
{
    return (block_t *) ((unsigned char *) block
        + get_size(block));
}
```



# 隐式链表：查找空闲块

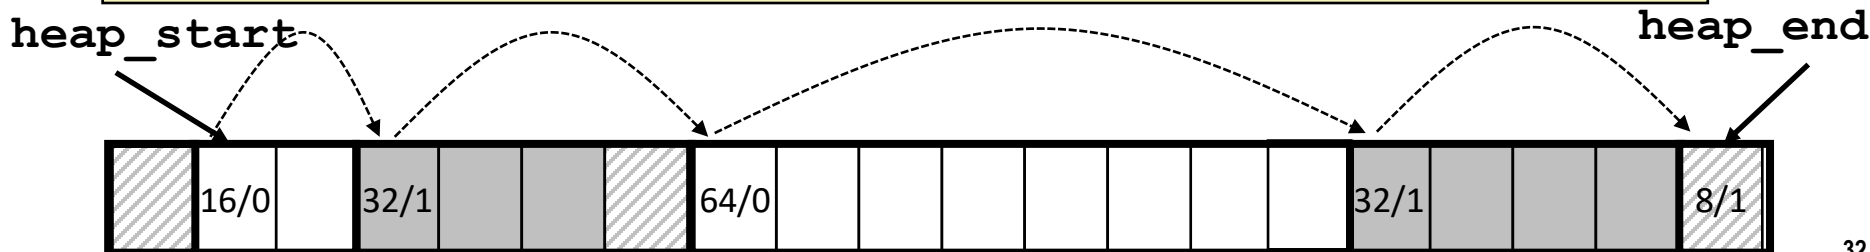


## Implicit List: Finding a Free Block

### ■ 首次匹配 *First fit*:

- 从链表开始搜索，选择第一个满足条件的空闲块 Search list from beginning, choose *first* free block that fits:
- 查找 *asize* 字节的空间（包括头部） Finding space for **asize** bytes (including header):

```
static block_t *find_fit(size_t asize)
{
    block_t *block;
    for (block = heap_start; block != heap_end;
         block = find_next(block)) {
        {
            if (!(get_alloc(block))
                && (asize <= get_size(block)))
                return block;
        }
    }
    return NULL; // No fit found
}
```



# 隐式链表：查找空闲块 Implicit List: Finding a Free Block



## ■ 首次匹配：First fit:

- 从链表开始搜索，选择**第一个**满足条件的空闲块 Search list from beginning, choose **first** free block that fits:
- 与总块数（分配和释放）成线性时间关系 Can take linear time in total number of blocks (allocated and free)
- 实际上会在链表开始时造成碎片 In practice it can cause “splinters” at beginning of list

## ■ 下一次匹配：Next fit:

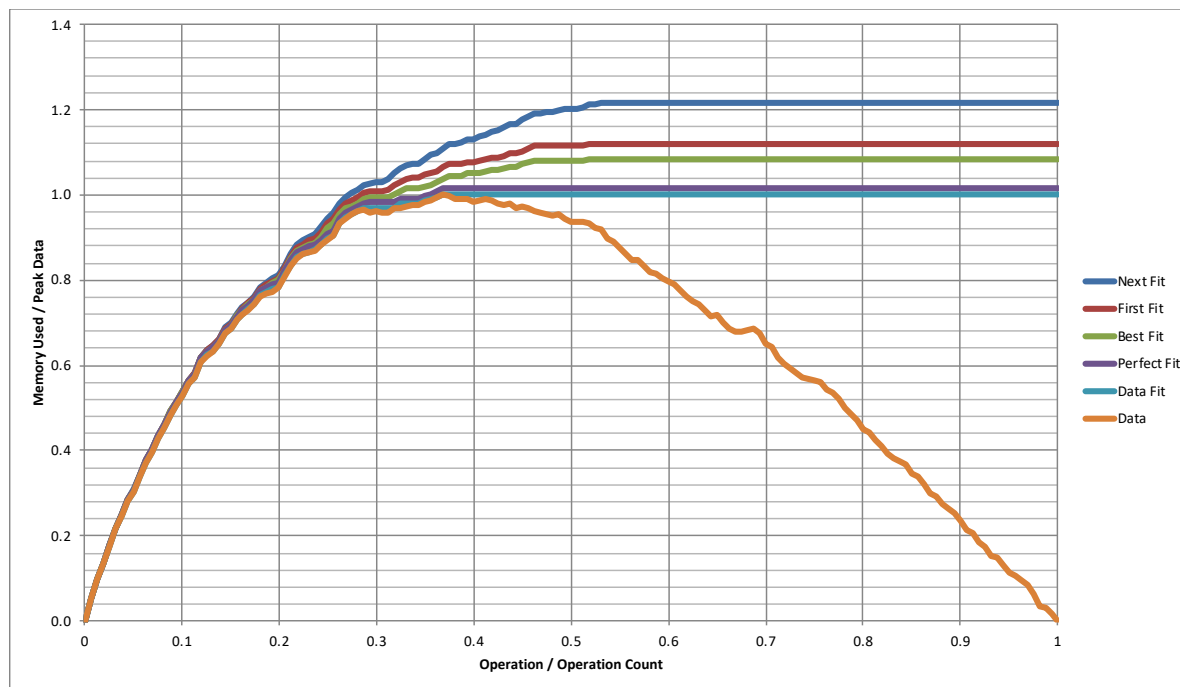
- 与first fit类似，但是从上一次搜索结束的位置开始查找 Like first fit, but search list starting where previous search finished
- 一般会比first fit快：避免了重扫描无用的块 Should often be faster than first fit: avoids re-scanning unhelpful blocks
- 部分研究表明更容易造成内存碎片 Some research suggests that fragmentation is worse

## ■ 最佳匹配：Best fit:

- 从链表中选择**最佳**的空闲块：最小满足需求的块 Search the list, choose the **best** free block: fits, with fewest bytes left over
- 保持内存碎片最小化-通常能改进内存利用率 Keeps fragments small—usually improves memory utilization
- 一般会比first fit慢 Will typically run slower than first fit



# 策略比较 Comparing Strategies



## ■ 总开销（对本基准） Total Overheads (for this benchmark)

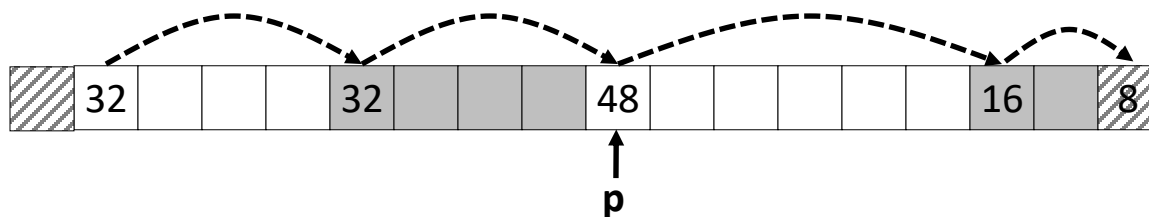
- 完美匹配 Perfect Fit: 1.6%
- 最佳匹配 Best Fit: 8.3%
- 首次匹配 First Fit: 11.9%
- 下次匹配 Next Fit: 21.6%

# 隐式链表：从空闲块中分配

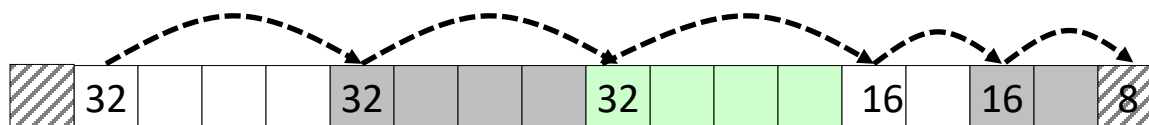
## Implicit List: Allocating in Free Block



- 从一个空闲块分配： **拆分** Allocating in a free block: ***splitting***
  - 由于分配的空间可能会比空闲空间小，因此可能会拆分空闲块  
Since allocated space might be smaller than free space, we might want to split the block



`split_block(p, 32)`

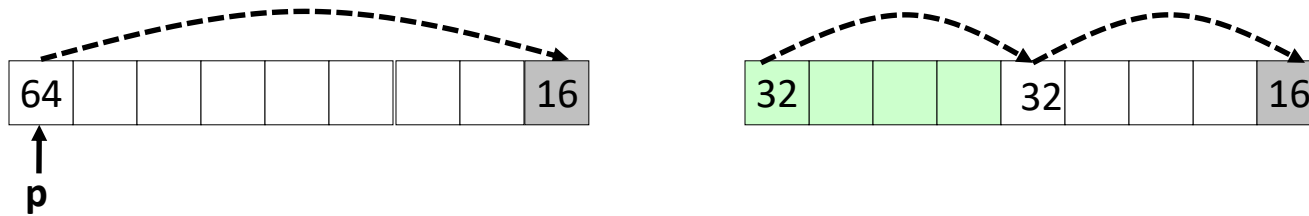




# 隐式链表：拆分空闲块

## Implicit List: Splitting Free Block

`split_block(p, 32)`



// Warning: This code is incomplete

```
static void split_block(block_t *block, size_t asize) {  
    size_t block_size = get_size(block);  
  
    if ((block_size - asize) >= min_block_size) {  
        write_header(block, asize, true);  
        block_t *block_next = find_next(block);  
        write_header(block_next, block_size - asize, false);  
    }  
}
```

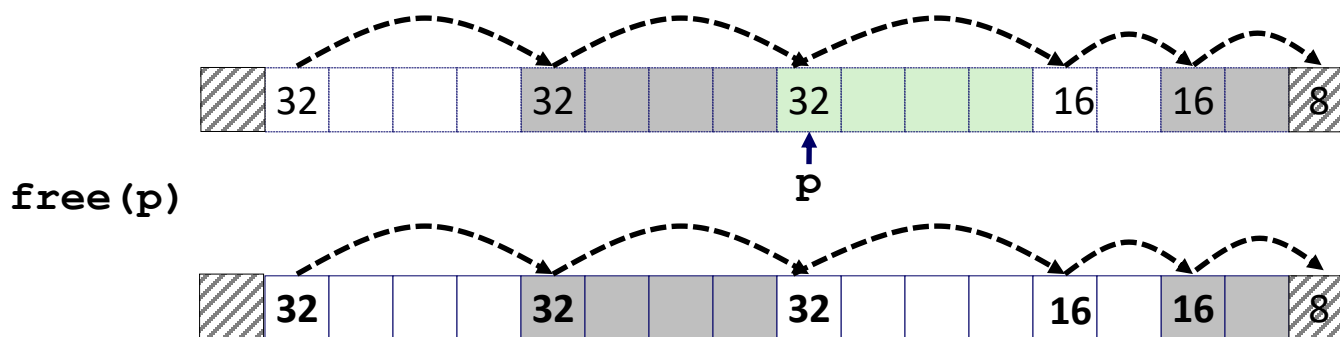


# 隐式链表：释放一个块

## Implicit List: Freeing a Block

### ■ 最简单的实现 Simplest implementation:

- 只需要清除“已分配”标记位 Need only clear the “allocated” flag
- 但是可能会导致“伪碎片” But can lead to “false fragmentation”



`malloc(5*SIZ)`

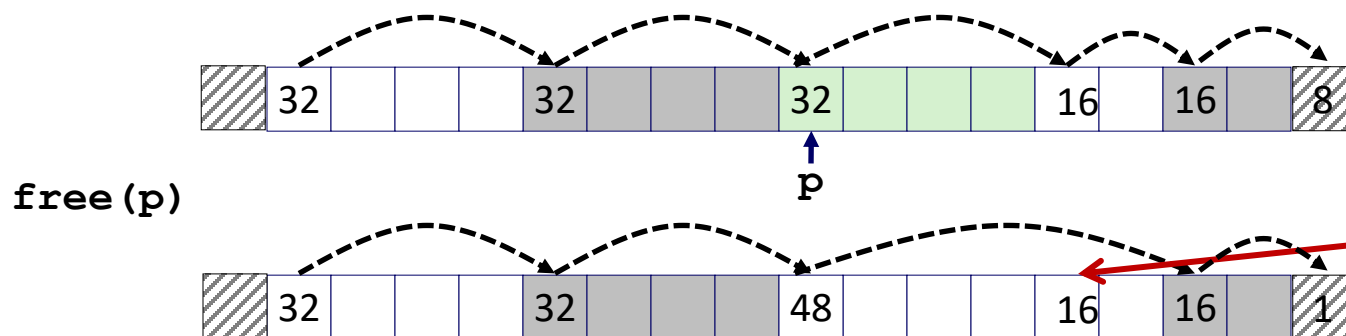
**诶呀! Yikes!**

**有足够的连续空闲空间，但是分配器找不到**  
**There is enough contiguous free space, but the allocator won't be able to find it**

# 隐式链表：合并 Implicit List: Coalescing



- 与下一个/前一个空闲块 **合并**, 如果有空闲块 Join (*coalesce*) with next/previous blocks, if they are free
  - 与下一个块合并 Coalescing with next block

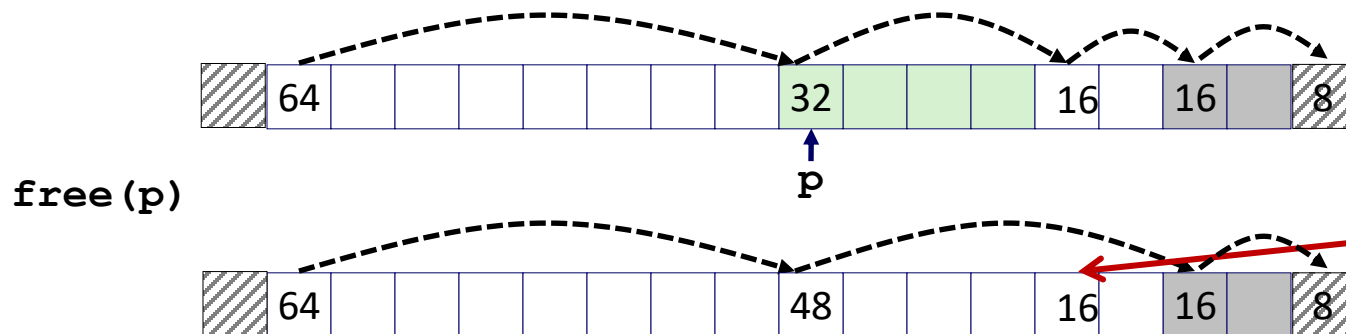


逻辑上  
不存在了  
*logically  
gone*

# 隐式链表：合并 Implicit List: Coalescing



- 与下一个/前一个空闲块 **合并**, 如果有空闲块 Join (*coalesce*) with next block, if it is free
  - 与下一个块合并 Coalescing with next block



逻辑上不存在了  
logically gone

- 但是怎么和前一个块合并? How do we coalesce with *previous* block?
  - 怎么知道从哪开始? How do we know where it starts?
  - 怎么能确定是否已经分配出去了? How can we determine whether its allocated?

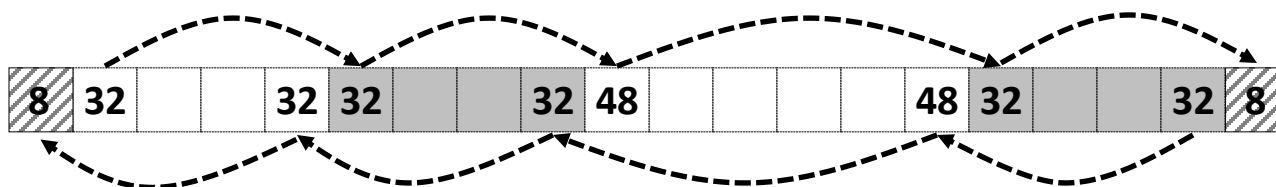


# 隐式链表：双向合并

## Implicit List: Bidirectional Coalescing

### ■ 边界标记 *Boundary tags* [Knuth73]

- 在空闲块“底部”（结束）的位置复制块大小/已分配字 Replicate size/allocated word at “bottom” (end) of free blocks
- 以额外的空间换取反向遍历链表功能 Allows us to traverse the “list” backwards, but requires extra space
- 重要和通用的技术 Important and general technique!



已分配和空闲块格式  
*Format of  
allocated and  
free blocks*

边界标记 Boundary tag  
(脚部 footer)

头部 Header

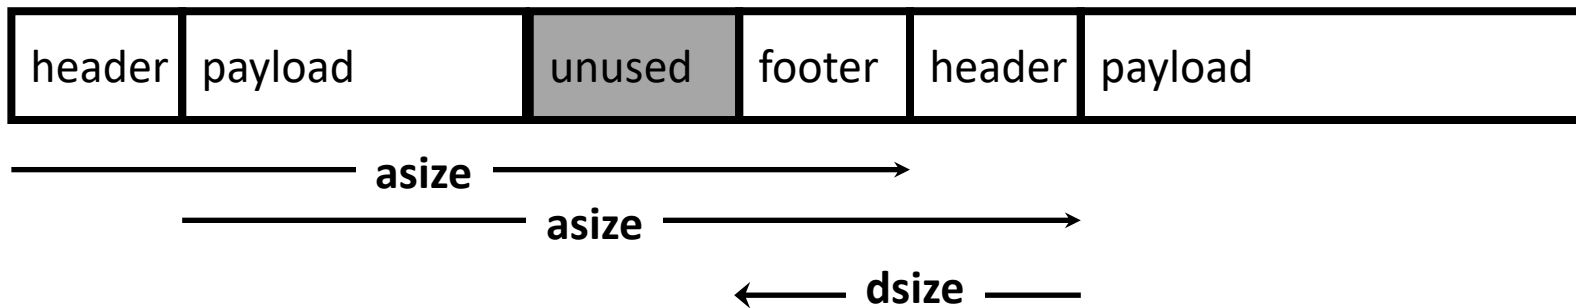


a = 1: Allocated block 已分配块  
a = 0: Free block 空闲块

Size: Total block size 大小：总的块大小

Payload: Application data 有效载荷：应用数据  
(allocated blocks only) (仅已分配块)

# 脚部的实现 Implementation with Footers



## ■ 定位当前块的脚部 Locating footer of current block

```
const size_t dsize = 2*sizeof(word_t);

static word_t *header_to_footer(block_t *block)
{
    size_t asize = get_size(block);
    return (word_t *) (block->payload + asize - dsize);
}
```



# 脚部的实现 Implementation with Footers



←  
1个字 1 word

## ■ 定位上一个块脚部 Locating footer of previous block

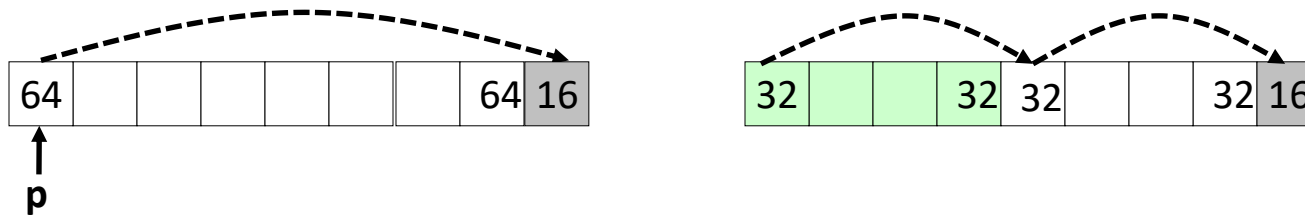
```
static word_t *find_prev_footer(block_t *block)
{
    return &(block->header) - 1;
}
```

# 拆分空闲块：完整版本

## Splitting Free Block: Full Version



`split_block(p, 32)`

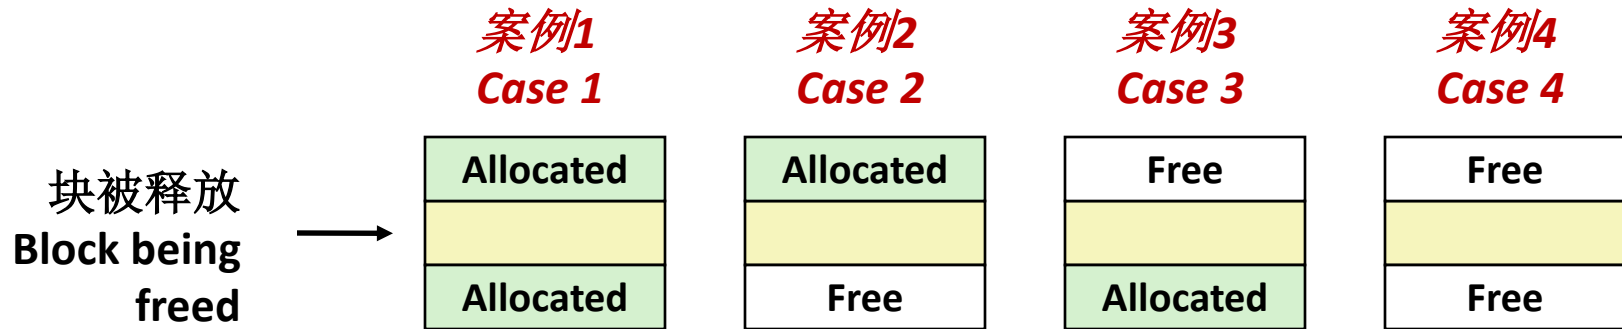


```
static void split_block(block_t *block, size_t asize) {
    size_t block_size = get_size(block);

    if ((block_size - asize) >= min_block_size) {
        write_header(block, asize, true);
        write_footer(block, asize, true);
        block_t *block_next = find_next(block);
        write_header(block_next, block_size - asize, false);
        write_footer(block_next, block_size - asize, false);
    }
}
```



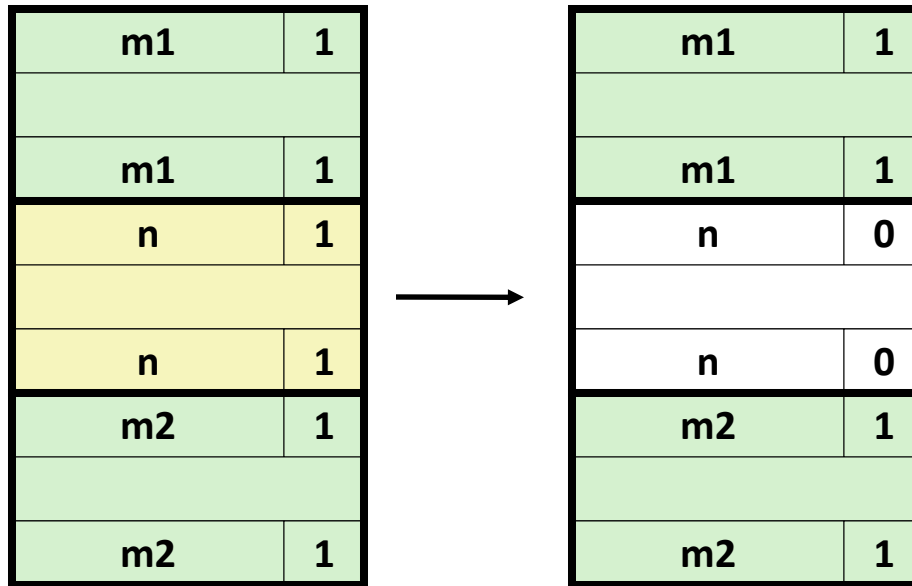
# 常量时间合并 Constant Time Coalescing





# 常量时间合并（案例1）

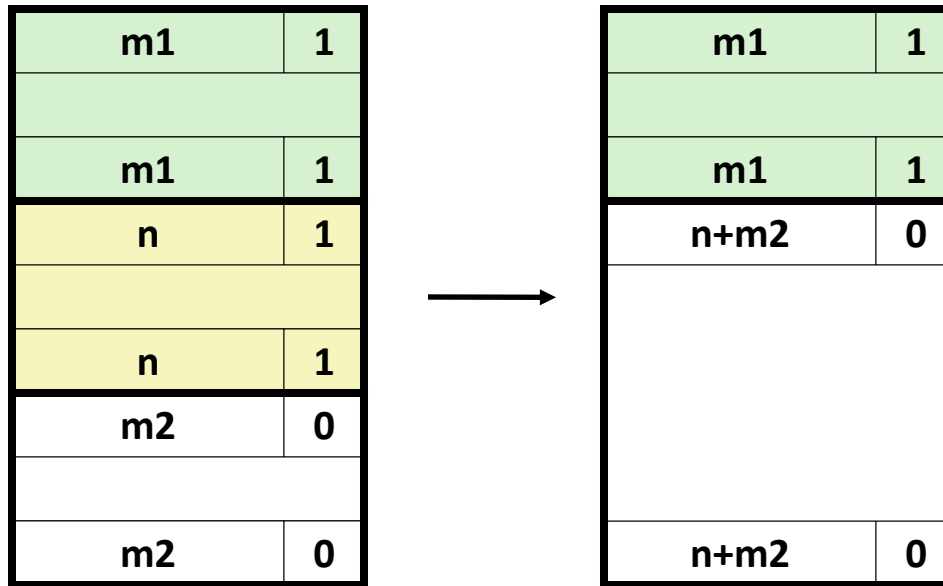
## Constant Time Coalescing (Case 1)





## 常量时间合并（案例2）

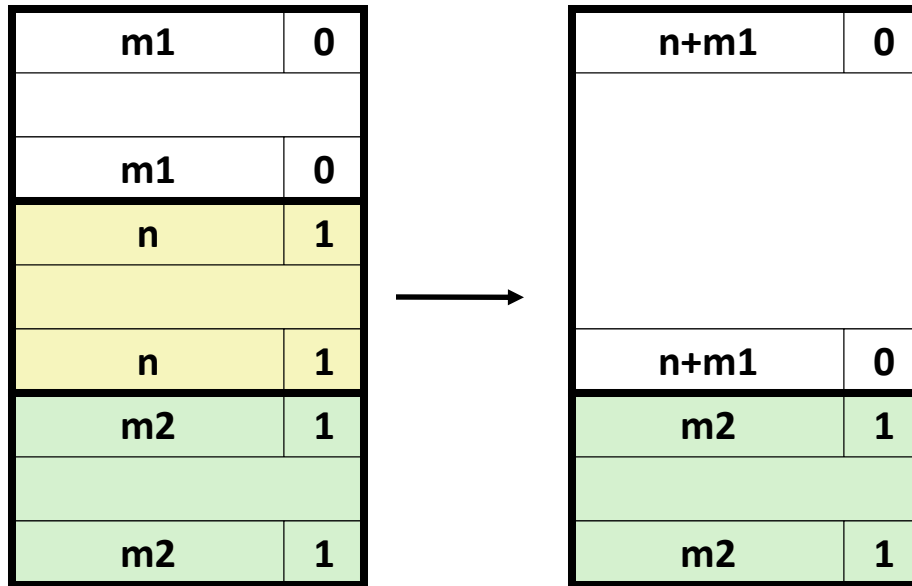
### Constant Time Coalescing (Case 2)





# 常量时间合并（案例3）

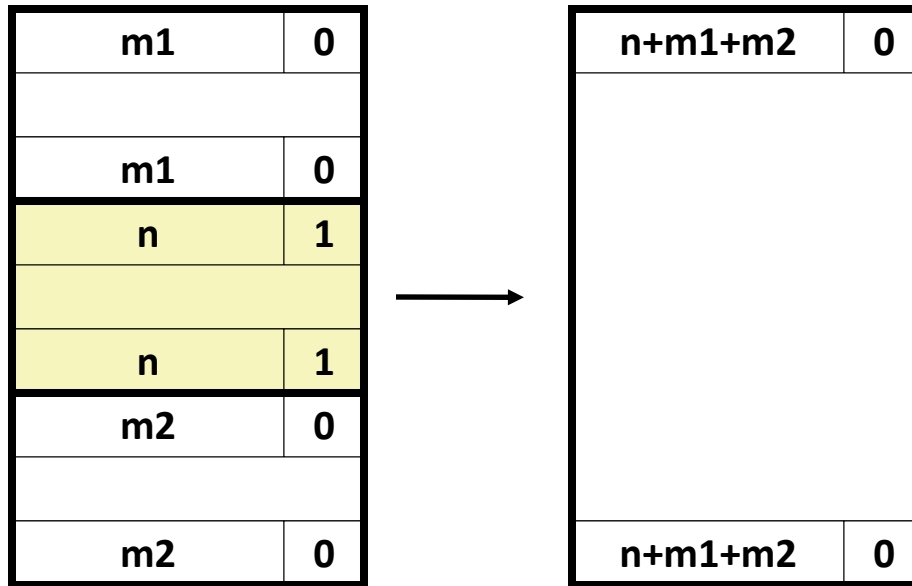
## Constant Time Coalescing (Case 3)



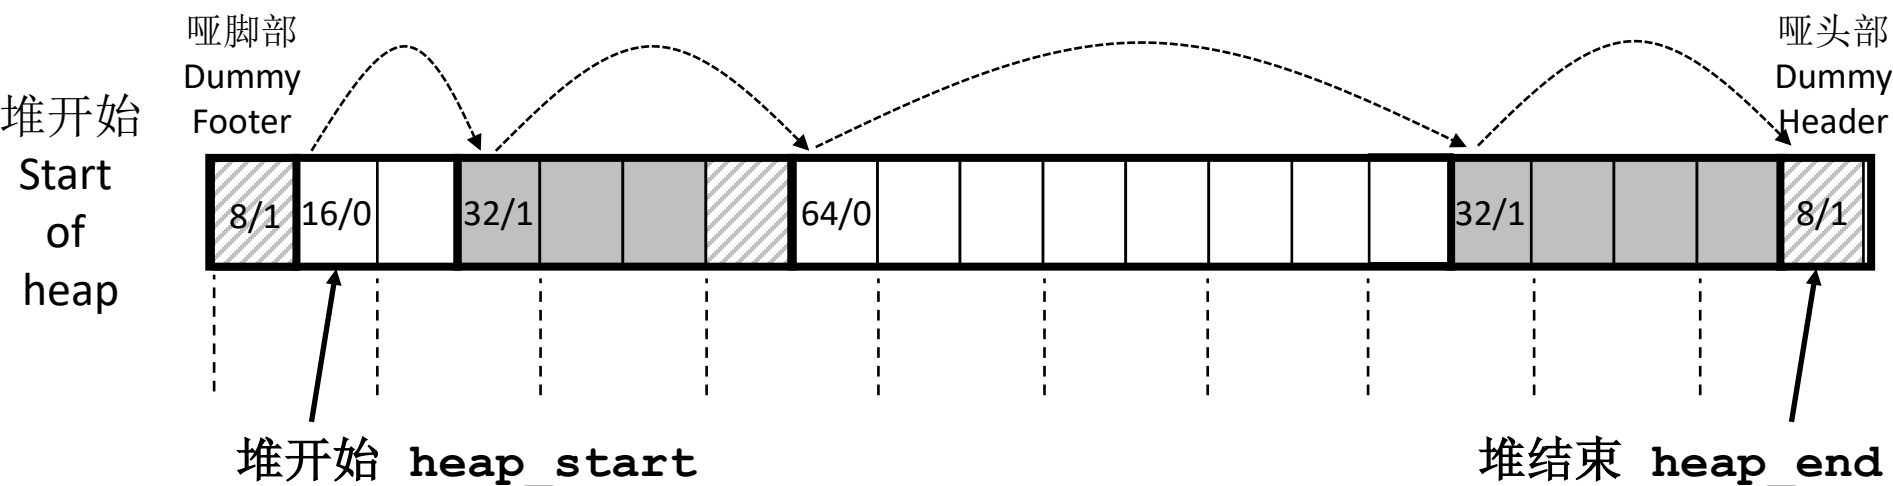


# 常量时间合并（案例4）

## Constant Time Coalescing (Case 4)



# 堆结构 Heap Structure



- **第一个头部之前的哑脚部 Dummy footer before first header**
  - 标记为已分配 Marked as allocated
  - 当释放第一个块时，防止意外合并 Prevents accidental coalescing when freeing first block
- **最后脚部之后的哑头部 Dummy header after last footer**
  - 在释放最后一块时，防止意外合并 Prevents accidental coalescing when freeing final block



# 顶层Malloc代码 Top-Level Malloc Code



```
const size_t dsize = 2*sizeof(word_t);

void *mm_malloc(size_t size)
{
    size_t asize = round_up(size + dsize, dsize);

    block_t *block = find_fit(asize);

    if (block == NULL)
        return NULL;

    size_t block_size = get_size(block);
    write_header(block, block_size, true);
    write_footer(block, block_size, true);

    split_block(block, asize);

    return header_to_payload(block);
}
```

$$\begin{aligned} \text{round\_up}(n, m) \\ &= \\ m * ((n+m-1) / m) \end{aligned}$$

# 顶层Free代码 Top-Level Free Code



```
void mm_free(void *bp)
{
    block_t *block = payload_to_header(bp);
    size_t size = get_size(block);

    write_header(block, size, false);
    write_footer(block, size, false);

    coalesce_block(block);
}
```



# 边界标记的缺点

## Disadvantages of Boundary Tags

- 内部碎片 Internal fragmentation
- 可以进一步优化吗？ Can it be optimized?
  - 哪些块需要脚部标记？ Which blocks need the footer tag?
  - 这意味着什么？ What does that mean?

大小 Size	a
载荷和填充 Payload and padding	
大小 Size	a



# 已分配块没有边界标记

## No Boundary Tag for Allocated Blocks

- 仅空闲块需要边界标记 Boundary tag needed only for free blocks
- 当块大小是16的整倍数，存在4个空闲位 When sizes are multiples of 16, have 4 spare bits

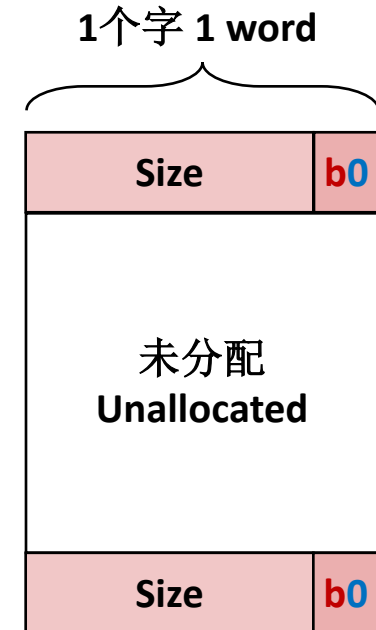


已分配块  
Allocated  
Block

a = 1: Allocated block 已分配  
a = 0: Free block 空闲块  
上一个块已分配  
b = 1: Previous block is allocated  
上一个块是空闲的  
b = 0: Previous block is free

Size: block size 大小: 块大小

Payload: application data 有效载荷:  
应用数据

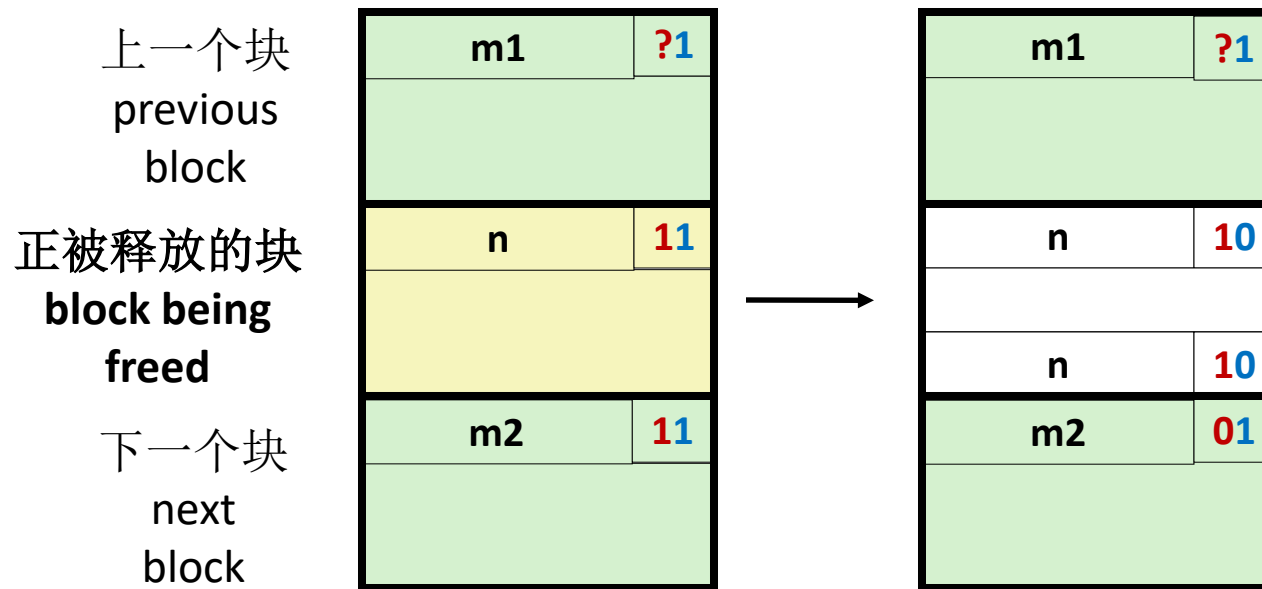


空闲块  
Free  
Block



# 已分配块没有边界标记（案例1）

## No Boundary Tag for Allocated Blocks (Case 1)



头部：使用2位（由于对齐的原因，这两个地址位始终为零）

Header: Use 2 bits (address bits always zero due to alignment):

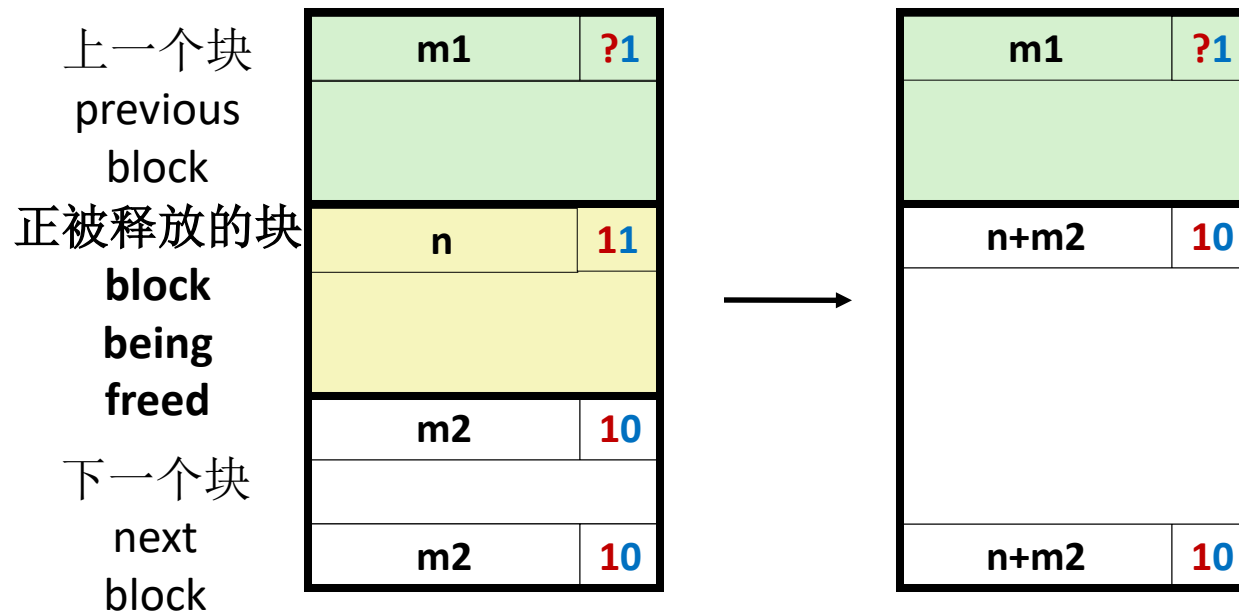
上一个分配的块<<1 | 当前分配的块

(previous block allocated)<<1 | (current block allocated)



# 已分配块没有边界标记（案例2）

## No Boundary Tag for Allocated Blocks (Case 2)



头部：使用2位（由于对齐的原因，这两个地址位始终为零）

Header: Use 2 bits (address bits always zero due to alignment):

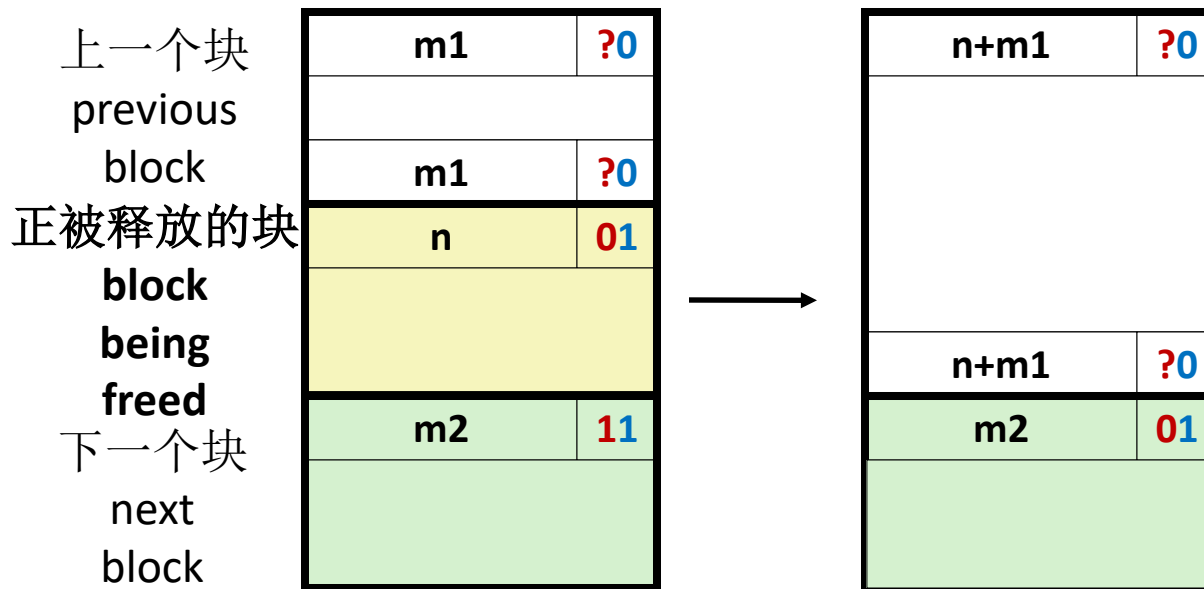
上一个分配的块<<1 | 当前分配的块

(previous block allocated)<<1 | (current block allocated)



# 已分配块没有边界标记（案例3）

## No Boundary Tag for Allocated Blocks (Case 3)



头部：使用2位（由于对齐的原因，这两个地址位始终为零）

Header: Use 2 bits (address bits always zero due to alignment):

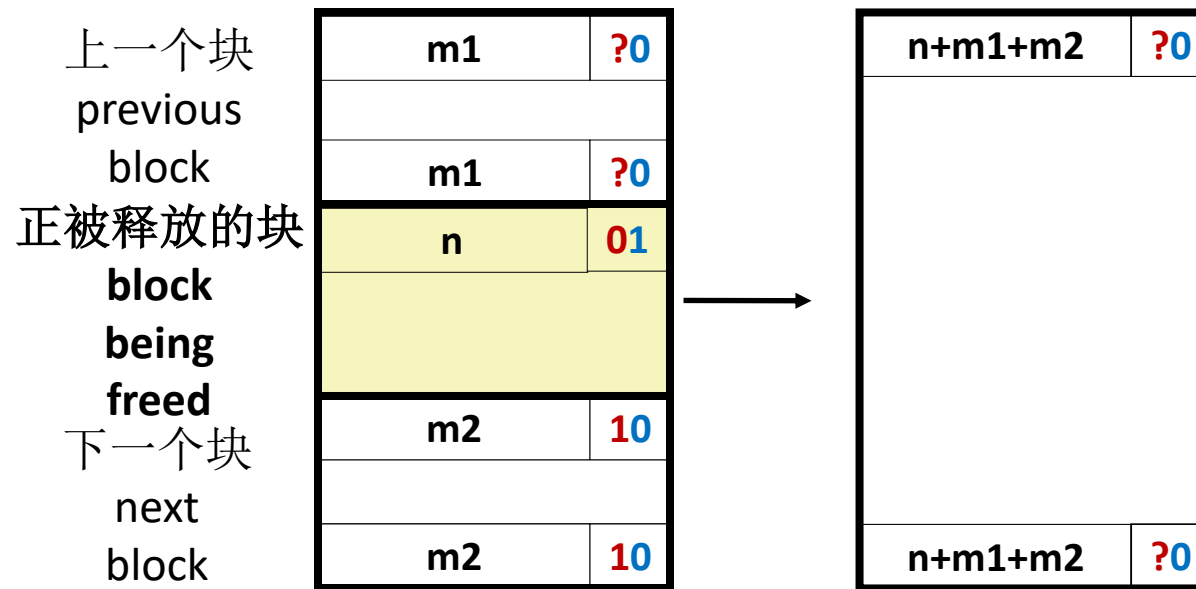
上一个分配的块<<1 | 当前分配的块

(previous block allocated)<<1 | (current block allocated)



# 已分配块没有边界标记（案例4）

## No Boundary Tag for Allocated Blocks (Case 4)



头部：使用2位（由于对齐的原因，这两个地址位始终为零）

Header: Use 2 bits (address bits always zero due to alignment):

上一个分配的块<<1 | 当前分配的块

(previous block allocated)<<1 | (current block allocated)



# 主要分配策略总结

## Summary of Key Allocator Policies



### ■ 选择策略 **Placement policy:**

- 首次匹配、下一次匹配、最佳匹配等 First-fit, next-fit, best-fit, etc.
- 在更低吞吐率和更少的碎片之间平衡 Trades off lower throughput for less fragmentation
- **有趣的观察:** 分离的空闲链表与最优选择策略接近, 且不用搜索整个链表  
*Interesting observation:* segregated free lists (next lecture) approximate a best fit placement policy without having to search entire free list

### ■ 拆分策略: **Splitting policy:**

- 什么时候需要拆分空闲块? When do we go ahead and split free blocks?
- 我们可能容忍多少内部碎片? How much internal fragmentation are we willing to tolerate?

### ■ 合并策略: **Coalescing policy:**

- **立即合并:** 每次free时合并 *Immediate coalescing:* coalesce each time **free** is called
- **延迟合并:** 为了提升free的性能, 当需要时再合并, 例如: *Deferred coalescing:* try to improve performance of **free** by deferring coalescing until needed. Examples:
  - 由于malloc扫描空闲列表时进行合并 Coalesce as you scan the free list for **malloc**
  - 当外部碎片超过某个阈值时进行合并 Coalesce when the amount of external fragmentation reaches some threshold



# 隐式链表：总结 Implicit Lists: Summary

- 实现：非常简单 **Implementation: very simple**
- 分配开销： **Allocate cost:**
  - 最差是线性时间 linear time worst case
- 释放开销： **Free cost:**
  - 最差常量时间 constant time worst case
  - 甚至包括合并 even with coalescing
- 内存使用 **Memory usage:**
  - 依赖于选择策略 will depend on placement policy
  - 首次匹配、下一次匹配或最佳匹配 First-fit, next-fit or best-fit
- 由于线性时间的分配开销，实际**malloc**和**free**并没有使用 **Not used in practice for malloc/free because of linear-time allocation**
  - 在很多特殊目的的应用中使用 used in many special purpose applications
- 然而拆分和基于边界标记的合并的概念对**所有**的分配器都是适用的  
**However, the concepts of splitting and boundary tag coalescing are general to *all* allocators**