



第12章 并发编程

同步：基础 Synchronization: Basics

100076202： 计算机系统导论

任课教师：

宿红毅 张艳 黎有琦 李秀星

原作者：

Randal E. Bryant and David R. O'Hallaron



**Carnegie
Mellon
University**



议题 Today

- **线程回顾** Threads review
- 共享 Sharing
- 互斥 Mutual exclusion
- 信号量 Semaphores

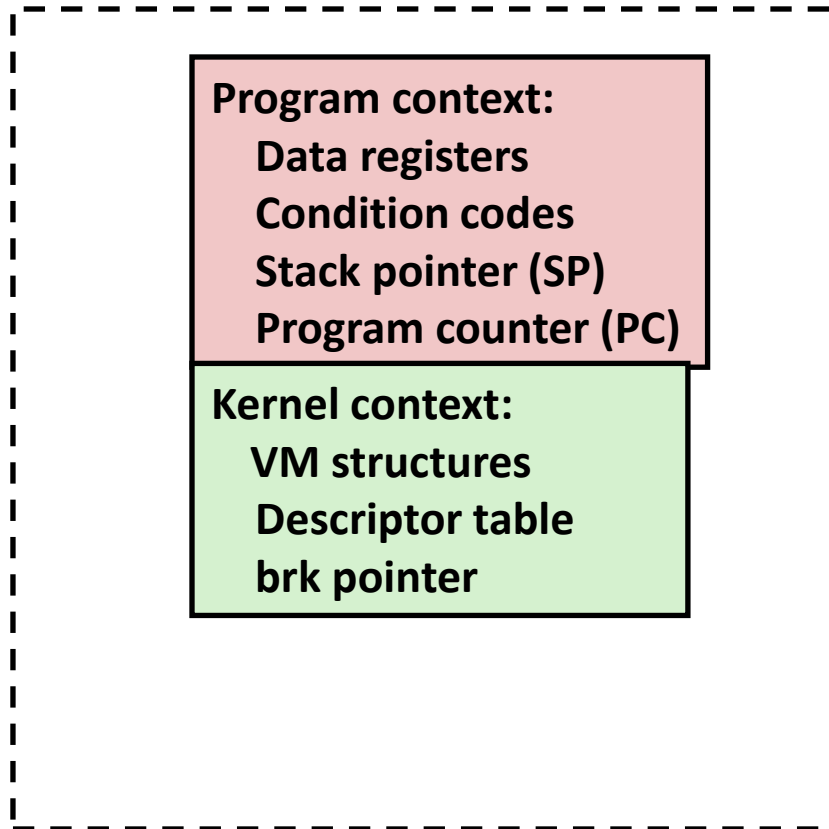


传统进程的视图 Traditional View of a Process

- 进程=进程上下文+代码、数据和栈 Process = process context + code, data, and stack


进程上下文

Process context



代码、数据和栈

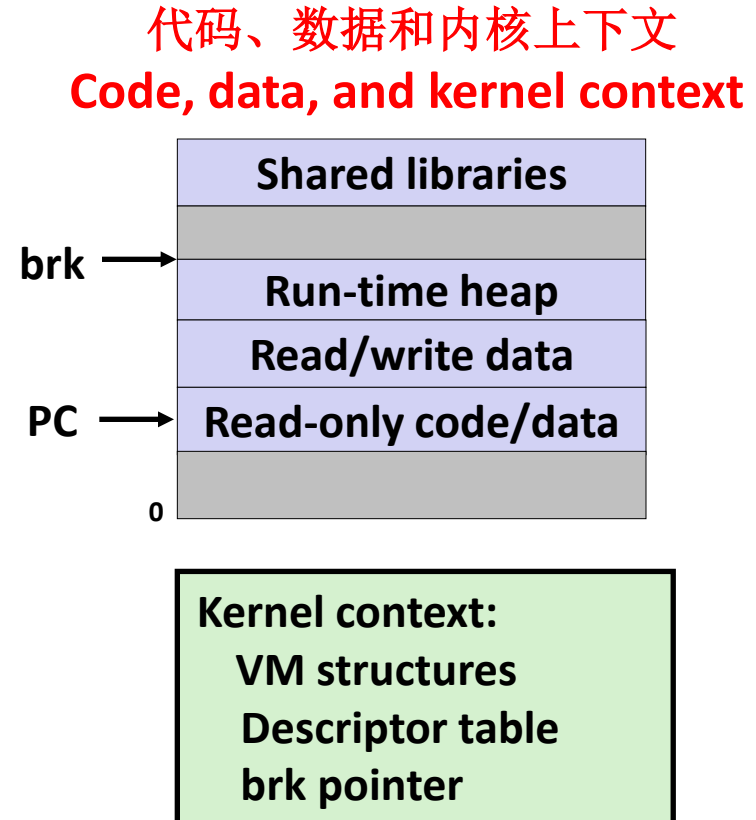
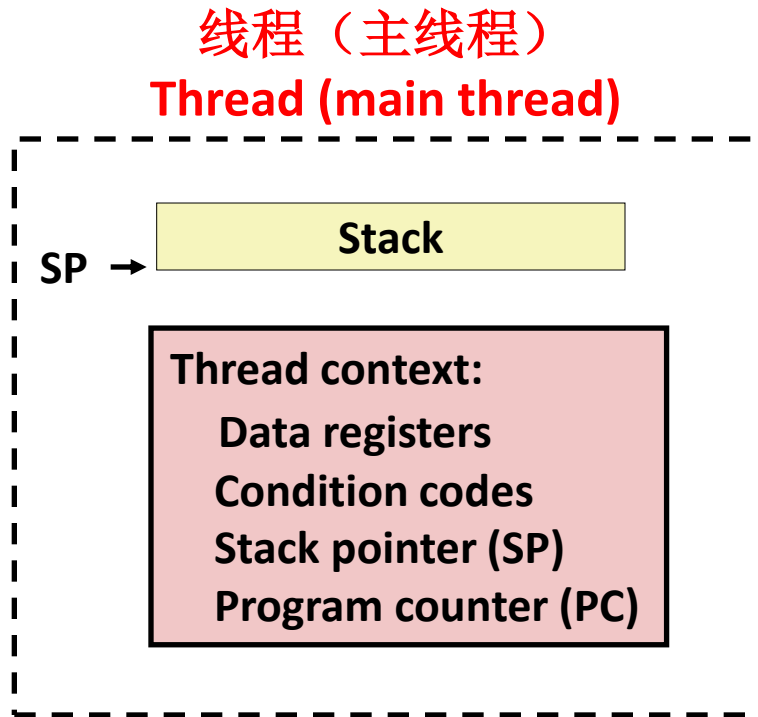
Code, data, and stack



进程的替代视图 Alternate View of a Process



- 进程=线程+（代码、数据和内核上下文） Process = thread + (code, data, and kernel context)





一个进程有多个线程-多线程进程

A Process With Multiple Threads

- 多个线程可以与一个进程关联 **Multiple threads can be associated with a process**
 - 每个线程都有自己的逻辑控制流 Each thread has its own logical control flow
 - 每个线程共享相同的代码、数据和内核上下文 Each thread shares the same code, data, and kernel context
 - 每个线程都有自己的局部变量栈 Each thread has its own stack for local variables
 - 但不受其他线程的保护 but not protected from other threads
 - 每个线程都有自己的线程id (TID) Each thread has its own thread id (TID)

线程1（主线程）
Thread 1 (main thread)

线程2（对等线程）
Thread 2 (peer thread)

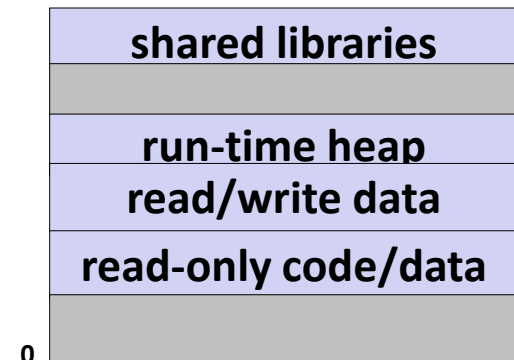
stack 1

stack 2

Thread 1 context:
Data registers
Condition codes
 SP_1
 PC_1

Thread 2 context:
Data registers
Condition codes
 SP_2
 PC_2

共享代码和数据
Shared code and data



Kernel context:
VM structures
Descriptor table
brk pointer



不要让图片迷惑你！

Don't let picture confuse you!

线程1（主线程） 线程2（对等线程）
Thread 1 (main thread) Thread 2 (peer thread)

stack 1

stack 2

Thread 1 context:
Data registers
Condition codes
 SP_1
 PC_1

Thread 2 context:
Data registers
Condition codes
 SP_2
 PC_2

共享代码和数据
Shared code and data

shared libraries

run-time heap

read/write data

read-only code/data

0

Kernel context:
VM structures
Descriptor table
brk pointer

内存在所有线程间共享

Memory is shared between all threads



议题 Today

- 线程回顾 Threads review
- **共享** Sharing
- 互斥 Mutual exclusion
- 信号量 Semaphores
- 生产者-消费者同步 Producer-Consumer Synchronization

在线程化的C语言程序中共享变量

Shared Variables in Threaded C Programs



- 问题：线程化C程序中的哪些变量是共享的？ Question: Which variables in a threaded C program are shared?
 - 答案并不像“全局变量是共享的”和“栈变量是私有的”那么简单 The answer is not as simple as “*global variables are shared*” and “*stack variables are private*”
- 定义：当且仅当多个线程引用x的某个实例时，变量x是共享的 Def: A variable **x** is *shared* if and only if multiple threads reference some instance of **x**.
- 需要以下问题的答案： Requires answers to the following questions:
 - 线程的内存模型是什么？ What is the memory model for threads?
 - 变量实例如何映射到内存？ How are instances of variables mapped to memory?
 - 有多少个线程可以引用每个实例？ How many threads might reference each of these instances?

线程内存模型：概念上



Threads Memory Model: Conceptual

- 多个线程在单个进程的上下文中运行 **Multiple threads run within the context of a single process**
- 每个线程都有自己独立的线程上下文 **Each thread has its own separate thread context**
 - 线程ID、栈、栈指针、PC、条件码和GP寄存器 Thread ID, stack, stack pointer, PC, condition codes, and GP registers
- 所有线程共享剩余的进程上下文 **All threads share the remaining process context**
 - 进程虚拟地址空间的代码、数据、堆和共享库段 Code, data, heap, and shared library segments of the process virtual address space
 - 打开文件和安装的信号处理程序 Open files and installed handlers

线程1 Thread 1 (私有 private)

stack 1

Thread 1 context:
Data registers
Condition codes
 SP_1
 PC_1

线程2 Thread 2 (私有 private)

stack 2

Thread 2 context:
Data registers
Condition codes
 SP_2
 PC_2

共享代码和数据 Shared code and data

shared libraries

run-time heap

read/write data

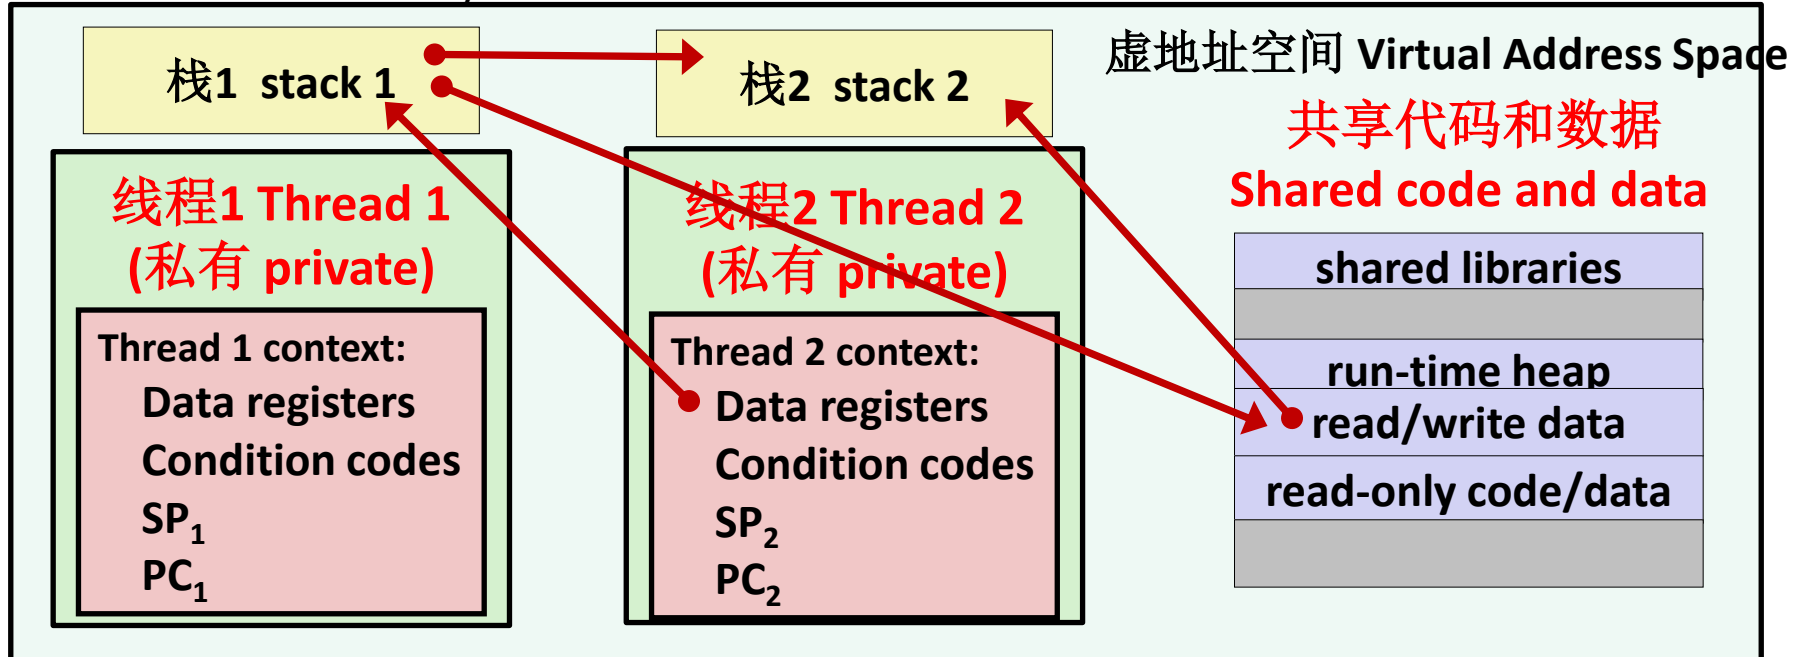
read-only code/data



线程内存模型：实际上

Threads Memory Model: Actual

- 未严格执行数据分离： Separation of data is not strictly enforced:
 - 寄存器值是真正独立和受保护的，但是... Register values are truly separate and protected, but...
 - 任何线程都可以读取和写入任何其他线程的栈 Any thread can read and write the stack of any other thread



概念模型和操作模型之间的不匹配是混淆和错误的根源

The mismatch between the conceptual and operation model is a source of confusion and errors

向线程传递参数 - 学究式方法



Passing an argument to a thread - Pedantic

```
int hist[N] = {0};

int main(int argc, char *argv[]) {
    long i;
    pthread_t tids[N];

    for (i = 0; i < N; i++) {
        long* p = Malloc(sizeof(long));
        *p = i;
        Pthread_create(&tids[i],
                       NULL,
                       thread,
                       (void *)p);
    }
    for (i = 0; i < N; i++)
        Pthread_join(tids[i], NULL);
    check();
}
```

```
void *thread(void *vargp)
{
    hist[(long *)vargp] += 1;
    Free(vargp);
    return NULL;
}
```

```
void check(void) {
    for (int i=0; i<N; i++) {
        if (hist[i] != 1) {
            printf("Failed at %d\n", i);
            exit(-1);
        }
    }
    printf("OK\n");
}
```

向线程传递参数 – 学究式方法

Passing an argument to a thread - Pedantic



```
int hist[N] = {0};

int main(int argc, char *argv[]) {
    long i;
    pthread_t tids[N];

    for (i = 0; i < N; i++) {
        long* p = Malloc(sizeof(long));
        *p = i;
        Pthread_create(&tids[i],
                       NULL,
                       thread,
                       (void *)p);
    }
    for (i = 0; i < N; i++)
        Pthread_join(tids[i], NULL);
    check();
}
```

```
void *thread(void *vargp)
{
    hist[(long *)vargp] += 1;
    Free(vargp);
    return NULL;
}
```

- 使用malloc为每个线程分配堆内存存放参数 Use malloc to create a per thread heap allocated place in memory for the argument
- 记得在线程中释放内存！ Remember to free in thread!
- 生产者-消费者模式 Producer-consumer pattern

向线程传递参数 – 另一种方法!

Passing an argument to a thread – Also OK!



```
int hist[N] = {0};

int main(int argc, char *argv[]) {
    long i;
    pthread_t tids[N];

    for (i = 0; i < N; i++)
        Pthread_create(&tids[i],
                       NULL,
                       thread,
                       (void *)i);

    for (i = 0; i < N; i++)
        Pthread_join(tids[i], NULL);
    check();
}
```

```
void *thread(void *vargp)
{
    hist[(long)vargp] += 1;
    return NULL;
}
```

- 使用强制转换也可以，因为长整数大小小于等于无类型指针的大小 **Ok to Use cast since sizeof(long) <= sizeof(void*)**
- 强制转换不会改变位模式 **Cast does NOT change bits**



向线程传递参数 – 警告！

Passing an argument to a thread – **WRONG!**

```
int hist[N] = {0};

int main(int argc, char *argv[]) {
    long i;
    pthread_t tids[N];

    for (i = 0; i < N; i++)
        Pthread_create(&tids[i],
                       NULL,
                       thread,
                       (void *) &i);

    for (i = 0; i < N; i++)
        Pthread_join(tids[i], NULL);
    check();
}
```

```
void *thread(void *vargp)
{
    hist[*(long*)vargp] += 1;
    return NULL;
}
```

- 取*i*的地址对所有的线程来说都指向同样的位置
&i points to same location for all threads!
- 产生数据竞争！ **Creates a data race!**

传递线程参数的三种方法

Three Ways to Pass Thread Arg



■ 申请/释放空间 **Malloc/free**

- 生产者申请空间, 传递指针给pthread_create Producer malloc's space, passes pointer to pthread_create
- 消费者释放指针空间 Consumer dereferences pointer

■ 指向栈槽位 **Ptr to stack slot**

- 生产者在pthread_create中传递生产者栈地址 Producer passes address to producer's stack in pthread_create
- 消费者释放指针 Consumer dereferences pointer

■ 强制转换成整数 **Cast of int**

- 在pthread_create中生产者强制转换整数/长整数为地址 Producer casts an int/long to address in pthread_create
- 消费者强制转换无类型指针参数回整数/长整数 Consumer casts void* argument back to int/long

示例程序说明共享

Example Program to Illustrate Sharing



```
char **ptr; /* global var */

int main(int argc, char *argv[])
{
    long i;
    pthread_t tid;
    char *msgs[2] = {
        "Hello from foo",
        "Hello from bar"
    };

    ptr = msgs;
    for (i = 0; i < 2; i++)
        Pthread_create(&tid,
            NULL,
            thread,
            (void *)i);
    Pthread_exit(NULL);
}
```

sharing.c

```
void *thread(void *vargp)
{
    long myid = (long)vargp;
    static int cnt = 0;

    printf("[%ld]: %s (cnt=%d)\n",
        myid, ptr[myid], ++cnt);
    return NULL;
}
```

对等线程间接通过全局ptr变量引用主线程的栈

Peer threads reference main thread's stack indirectly through global ptr variable

一种通用方法传递单个参数给一个线程例程
A common way to pass a single argument to a thread routine

在线程化的C语言程序中共享变量

Shared Variables in Threaded C Programs



- **问题：线程化C程序中的哪些变量是共享的？ Question: Which variables in a threaded C program are shared?**
 - 答案并不像“全局变量是共享的”和“栈变量是私有的”那么简单 The answer is not as simple as “*global variables are shared*” and “*stack variables are private*”
- **定义：当且仅当多个线程引用x的某个实例时，变量x是共享的 Def: A variable **x** is *shared* if and only if multiple threads reference some instance of **x**.**
- **需要以下问题的答案： Requires answers to the following questions:**
 - 线程的内存模型是什么？ What is the memory model for threads?
 - 变量实例如何映射到内存？ How are instances of variables mapped to memory?
 - 有多少个线程可以引用每个实例？ How many threads might reference each of these instances?

映射变量实例到内存

Mapping Variable Instances to Memory



■ 全局变量 **Global variables**

- 定义：在函数外部声明的变量 *Def: Variable declared outside of a function*
- 虚拟内存仅包含任何全局变量的一个实例 **Virtual memory contains exactly one instance of any global variable**

■ 局部变量 **Local variables**

- 定义：在函数内声明的没有静态属性的变量 *Def: Variable declared inside function without **static** attribute*
- 每个线程栈包含每个局部变量的一个实例 **Each thread stack contains one instance of each local variable**

■ 局部静态变量 **Local static variables**

- 定义：在函数内部声明的带有静态属性的变量 *Def: Variable declared inside function with the **static** attribute*
- 虚拟内存只包含任何本地静态变量的一个实例 **Virtual memory contains exactly one instance of any local static variable.**

映射变量实例到内存

Mapping Variable Instances to Memory



```
char **ptr;    /* global var */

int main(int main, char *argv[])
{
    long i;
    pthread_t tid;
    char *msgs[2] = {
        "Hello from foo",
        "Hello from bar"
    };

    ptr = msgs;
    for (i = 0; i < 2; i++)
        Pthread_create(&tid,
            NULL,
            thread,
            (void *)i);
    Pthread_exit(NULL);
}
```

sharing.c

```
void *thread(void *vargp)
{
    long myid = (long)vargp;
    static int cnt = 0;

    printf("[%ld]:  %s (cnt=%d)\n",
        myid, ptr[myid], ++cnt);
    return NULL;
}
```

映射变量实例到内存

Mapping Variable Instances to Memory



全局变量: 1个实例 **Global var:** 1 instance (ptr [data])

局部变量: 1个实例 **Local vars:** 1 instance (i.m, msgs.m, tid.m)

```
char **ptr; /* global var */

int main(int main, char *argv[])
{
    long i;
    pthread_t tid;
    char *msgs[2] = {
        "Hello from foo",
        "Hello from bar"
    };

    ptr = msgs;
    for (i = 0; i < 2; i++)
        Pthread_create(&tid,
            NULL,
            thread,
            (void *)i);
    Pthread_exit(NULL);
}
```

sharing.c

局部变量: 2个实例 **Local var:** 2 instances (myid.p0 [peer thread 0's stack], myid.p1 [peer thread 1's stack])

```
void *thread(void *vargp)
{
    long myid = (long)vargp;
    static int cnt = 0;

    printf("[%ld]: %s (cnt=%d)\n",
        myid, ptr[myid], ++cnt);
    return NULL;
}
```

局部静态变量: 1个实例
Local static var: 1 instance (cnt [data])

共享变量分析 Shared Variable Analysis



■ 哪些变量是共享的? Which variables are shared?

<i>Variable instance</i>	<i>Referenced by main thread?</i>	<i>Referenced by peer thread 0?</i>	<i>Referenced by peer thread 1?</i>
ptr	yes	yes	yes
cnt	no	yes	yes
i.m	yes	no	no
msgs.m	yes	yes	yes
myid.p0	no	yes	no
myid.p1	no	no	yes

```
char **ptr; /* global var */
int main(int main, char *argv[]) {
    long i; pthread_t tid;
    char *msgs[2] = {"Hello from foo",
                    "Hello from bar" };

    ptr = msgs;
    for (i = 0; i < 2; i++)
        Pthread_create(&tid,
                       NULL, thread, (void *)i);
    Pthread_exit(NULL); }
```

```
void *thread(void *vargp)
{
    long myid = (long)vargp;
    static int cnt = 0;

    printf("[%ld]: %s (cnt=%d)\n",
           myid, ptr[myid], ++cnt);
    return NULL;
}
```

共享变量分析 Shared Variable Analysis



■ 哪些变量是共享的？ Which variables are shared?

<i>Variable instance</i>	<i>Referenced by main thread?</i>	<i>Referenced by peer thread 0?</i>	<i>Referenced by peer thread 1?</i>
<code>ptr</code>	yes	yes	yes
<code>cnt</code>	no	yes	yes
<code>i.m</code>	yes	no	no
<code>msgs.m</code>	yes	yes	yes
<code>myid.p0</code>	no	yes	no
<code>myid.p1</code>	no	no	yes

- 答案：变量 x 是共享的，当且仅当多个线程引用最少一个 x 的实例，因此： Answer: A variable x is shared iff multiple threads reference at least one instance of x .
Thus:

- `ptr`、`cnt`和`msgs`是共享的 `ptr`, `cnt`, and `msgs` are shared
- `i`和`myid`不是共享的 `i` and `myid` are *not* shared



同步线程 Synchronizing Threads

- 共享变量很方便。。。 Shared variables are handy...
-但会引入严重同步错误的可能性 ...but introduce the possibility of nasty *synchronization* errors.

badcnt.c:不正确的同步

badcnt.c: Improper Synchronization



```
/* Global shared variable */
volatile long cnt = 0; /* Counter */

int main(int argc, char **argv)
{
    long niters;
    pthread_t tid1, tid2;

    niters = atoi(argv[1]);
    Pthread_create(&tid1, NULL,
                  thread, &niters);
    Pthread_create(&tid2, NULL,
                  thread, &niters);
    Pthread_join(tid1, NULL);
    Pthread_join(tid2, NULL);

    /* Check result */
    if (cnt != (2 * niters))
        printf("BOOM! cnt=%ld\n", cnt);
    else
        printf("OK cnt=%ld\n", cnt);
    exit(0);
}
```

badcnt.c

```
/* Thread routine */
void *thread(void *vargp)
{
    long i, niters =
        *((long *)vargp);

    for (i = 0; i < niters; i++)
        cnt++;

    return NULL;
}
```

```
linux> ./badcnt 10000
OK cnt=20000
linux> ./badcnt 10000
BOOM! cnt=13051
linux>
```

cnt应该等于20,000 cnt
should equal 20,000.
发生了什么错? What went
wrong?

计数循环的汇编代码

Assembly Code for Counter Loop



线程*i*中循环计数的C代码 C code for counter loop in thread *i*

```
for (i = 0; i < niters; i++)  
    cnt++;
```

线程*i*的汇编代码 *Asm code for thread *i**

<pre>movq (%rdi), %rcx testq %rcx,%rcx jle .L2 movl \$0, %eax</pre>	} H_i : Head 循环头
<pre>----- .L3: movq cnt(%rip), %rdx addq \$1, %rdx movq %rdx, cnt(%rip)</pre>	} L_i : Load cnt 装载cnt U_i : Update cnt 更新cnt S_i : Store cnt 存储cnt
<pre>----- addq \$1, %rax cmpq %rcx, %rax jne .L3 .L2:</pre>	} T_i : Tail 循环尾



并发执行 Concurrent Execution

- **关键思想:** 一般来说, 任何顺序一致的*指令交错执行都是可能的, 但有些会产生意想不到的结果! **Key idea:** In general, any **sequentially consistent*** interleaving is possible, but some give an unexpected result!

- I_i 表示线程i执行指令l I_i denotes that thread i executes instruction l
- $\%rdx_i$ 是线程i上下文中 $\%rdx$ 的内容 $\%rdx_i$ is the content of $\%rdx$ in thread i's context

i (thread)	instr _i	$\%rdx_1$	$\%rdx_2$	cnt
1	H ₁	-	-	0
1	L ₁	0	-	0
1	U ₁	1	-	0
1	S ₁	1	-	1
2	H ₂	-	-	1
2	L ₂	-	1	1
2	U ₂	-	2	1
2	S ₂	-	2	2
2	T ₂	-	2	2
1	T ₁	1	-	2

OK

*现在。实际上, 在x86上, 甚至可以进行非顺序一致的指令交错执行

*For now. In reality, on x86 even non-sequentially consistent interleavings are possible

并发执行 Concurrent Execution



- **关键思想:** 一般来说, 任何顺序一致的*指令交错执行都是可能的, 但有些会产生意想不到的结果! **Key idea:** In general, any sequentially consistent interleaving is possible, but some give an unexpected result!

- I_i 表示线程*i*执行指令 I_i denotes that thread *i* executes instruction *I*
- $\%rdx_i$ 是线程*i*上下文中 $\%rdx$ 的内容 $\%rdx_i$ is the content of $\%rdx$ in thread *i*'s context

<i>i</i> (thread)	$instr_i$	$\%rdx_1$	$\%rdx_2$	cnt
1	H_1	-	-	0
1	L_1	0	-	0
1	U_1	1	-	0
1	S_1	1	-	1
2	H_2	-	-	1
2	L_2	-	1	1
2	U_2	-	2	1
2	S_2	-	2	2
2	T_2	-	2	2
1	T_1	1	-	2



线程1临界区
Thread 1
critical section



线程2临界区
Thread 2
critical section

OK

并发执行（续）



Concurrent Execution (cont)

- 不正确的顺序：两个线程递增计数器，但结果是1而不是2
Incorrect ordering: two threads increment the counter, but the result is 1 instead of 2

i (thread)	instr _i	%rdx ₁	%rdx ₂	cnt
1	H ₁	-	-	0
1	L ₁	0	-	0
1	U ₁	1	-	0
2	H ₂	-	-	0
2	L ₂	-	0	0
1	S ₁	1	-	1
1	T ₁	1	-	1
2	U ₂	-	1	1
2	S ₂	-	1	1
2	T ₂	-	1	1

哎呀! Oops!



并发执行（续）

Concurrent Execution (cont)

- 这个顺序会怎么样？ How about this ordering?

i (thread)	instr _i	%rdx ₁	%rdx ₂	cnt
1	H ₁			0
1	L ₁	0		
2	H ₂			
2	L ₂		0	
2	U ₂		1	
2	S ₂		1	1
1	U ₁	1		
1	S ₁	1		1
1	T ₁			1
2	T ₂			1

哎呀！ Oops!

- 我们可以使用进度图分析行为 We can analyze the behavior using a *progress graph*

进度图 Progress Graphs




进度图描述了并发线程的离散执行状态空间 A **progress graph** depicts the discrete **execution state space** of concurrent threads.

每个轴对应于线程中的指令顺序 Each axis corresponds to the sequential order of instructions in a thread.

每个点对应于可能的执行状态 Each point corresponds to a possible **execution state** (Inst₁, Inst₂).

例如 (L₁, S₂) 表示状态, 其中线程1已完成L₁和线程2已完成S₂ E.g., (L₁, S₂) denotes state where thread 1 has completed L₁ and thread 2 has completed S₂.

线程2 Thread 2




进度图中的轨迹

Trajectories in Progress Graphs



线程2 Thread 2



轨迹是一系列合法状态转换，描述了线程的一种可能并发执行。

A **trajectory** is a sequence of legal state transitions that describes one possible concurrent execution of the threads.

例如: Example:


H1, L1, U1, H2, L2, S1, T1, U2, S2, T2

进度图中的轨迹

Trajectories in Progress Graphs



线程2 Thread 2



轨迹是一系列合法状态转换，描述了线程的一种可能并发执行。

A **trajectory** is a sequence of legal state transitions that describes one possible concurrent execution of the threads.

例如: Example:

H1, L1, U1, H2, L2, S1, T1, U2, S2, T2


线程1 Thread 1

临界区和不安全区域

Critical Sections and Unsafe Regions



线程2 Thread 2



L、U和S形成关于共享变量cnt的
临界区 L, U, and S form a **critical section** with respect to the shared variable cnt

临界区中的指令（写入一些共享变量）不应交错 Instructions in critical sections (wrt some shared variable) should not be interleaved

发生这种交错的状态集形成不
安全区域 Sets of states where such interleaving occurs form **unsafe regions**


线程1 Thread 1

临界区和不安全区域

Critical Sections and Unsafe Regions



线程2 Thread 2



定义: 当且仅当轨迹不进入任何不安全区域, 轨迹是安全的

Def: A trajectory is **safe** iff it does not enter any unsafe region

声明: 当且仅当轨迹是安全的, 轨迹是正确的 (写入cnt)

Claim: A trajectory is correct (wrt cnt) iff it is safe

badcnt.c: 不正确的同步

badcnt.c: Improper Synchronization



```
/* Global shared variable */
volatile long cnt = 0; /* Counter */

int main(int argc, char **argv)
{
    long niters;
    pthread_t tid1, tid2;

    niters = atoi(argv[1]);
    Pthread_create(&tid1, NULL,
                  thread, &niters);
    Pthread_create(&tid2, NULL,
                  thread, &niters);
    Pthread_join(tid1, NULL);
    Pthread_join(tid2, NULL);

    /* Check result */
    if (cnt != (2 * niters))
        printf("BOOM! cnt=%ld\n", cnt);
    else
        printf("OK cnt=%ld\n", cnt);
    exit(0);
}
```

badcnt.c

```
/* Thread routine */
void *thread(void *vargp)
{
    long i, niters =
        *((long *)vargp);

    for (i = 0; i < niters; i++)
        cnt++;

    return NULL;
}
```

Variable	main	thread1	thread2
cnt			
niters.m			
tid1.m			
i.1			
i.2			
niters.1			
niters.2			

badcnt.c: 不正确的同步

badcnt.c: Improper Synchronization



```
/* Global shared variable */
volatile long cnt = 0; /* Counter */

int main(int argc, char **argv)
{
    long niters;
    pthread_t tid1, tid2;

    niters = atoi(argv[1]);
    Pthread_create(&tid1, NULL,
                  thread, &niters);
    Pthread_create(&tid2, NULL,
                  thread, &niters);
    Pthread_join(tid1, NULL);
    Pthread_join(tid2, NULL);

    /* Check result */
    if (cnt != (2 * niters))
        printf("BOOM! cnt=%ld\n", cnt);
    else
        printf("OK cnt=%ld\n", cnt);
    exit(0);
}
```

badcnt.c

```
/* Thread routine */
void *thread(void *vargp)
{
    long i, niters =
        *((long *)vargp);

    for (i = 0; i < niters; i++)
        cnt++;

    return NULL;
}
```

Variable	main	thread1	thread2
cnt	yes*	yes	yes
niters.m	yes	no	no
tid1.m	yes	no	no
i.1	no	yes	no
i.2	no	no	yes
niters.1	no	yes	no
niters.2	no	no	yes



议题 Today

- 线程回顾 Threads review
- 共享 Sharing
- **互斥** Mutual exclusion
- 信号量 Semaphores
- 生产者-消费者同步 Producer-Consumer Synchronization

执行互斥 Enforcing Mutual Exclusion



- **问题：** 我们如何保证安全的轨迹？ **Question:** How can we guarantee a safe trajectory?
- **答：** 我们必须**同步**线程的执行，以便它们永远不会有不安全的轨迹 **Answer:** We must **synchronize** the execution of the threads so that they can never have an unsafe trajectory.
 - 即需要保证每个临界区的**互斥访问** i.e., need to guarantee **mutually exclusive access** for each critical section.
- **经典解决方案：** **Classic solution:**
 - 互斥锁 (pthreads) Mutex (pthreads)
 - 信号量 (Edsger Dijkstra) Semaphores (Edsger Dijkstra)
- **其他方法（超出我们的讨论范围）** **Other approaches (out of our scope)**
 - 条件变量 (pthreads) Condition variables (pthreads)
 - 监视器 (Java) Monitors (Java)

互斥锁 (mutex)

MUTual EXclusion (mutex)



- **互斥锁**: 布尔型同步变量 **Mutex**: boolean synchronization variable
- `enum {locked = 0, unlocked = 1}`
- **lock(m)**
 - 如果互斥锁当前未锁定, 请锁定它并返回 If the mutex is currently not locked, lock it and return
 - 否则, 等待 (挂起、休眠等) 并重试 Otherwise, wait (spinning, yielding, etc) and retry
- **unlock(m)**
 - 将互斥锁状态更新为解锁 Update the mutex state to unlocked

互斥锁 (mutex)

MUTual EXclusion (mutex)



- **互斥锁**: 布尔型同步变量* **Mutex**: boolean synchronization variable *
- **Swap(*a, b)**
[t = *a; *a = b; return t;]
// [] -通过硬件/OS的魔力实现原子操作 atomic by the magic of hardware / OS
- **Lock(m):**
while (swap(&m->state, locked) == locked) ;
- **Unlock(m):**
m->state = unlocked;

**现在。实际上，许多其他实现和设计选择（参见15-410、418等）。*

** For now. In reality, many other implementations and design choices (c.f., 15-410, 418, etc).*



badcnt.c: 不正确的同步

badcnt.c: Improper Synchronization

```
/* Global shared variable */
volatile long cnt = 0; /* Counter */

int main(int argc, char **argv)
{
    long niters;
    pthread_t tid1, tid2;

    niters = atoi(argv[1]);
    Pthread_create(&tid1, NULL,
                  thread, &niters);
    Pthread_create(&tid2, NULL,
                  thread, &niters);
    Pthread_join(tid1, NULL);
    Pthread_join(tid2, NULL);

    /* Check result */
    if (cnt != (2 * niters))
        printf("BOOM! cnt=%ld\n", cnt);
    else
        printf("OK cnt=%ld\n", cnt);
    exit(0);
}
```

badcnt.c

```
/* Thread routine */
void *thread(void *vargp)
{
    long i, niters =
        *((long *)vargp);

    for (i = 0; i < niters; i++)
        cnt++;

    return NULL;
}
```

如何使用同步解决此问题？
How can we fix this using
synchronization?



goodmcent.c: 互斥锁同步

goodmcent.c: Mutex Synchronization

- 为共享变量cnt定义并初始化互斥锁: Define and initialize a mutex for the shared variable cnt:

```
volatile long cnt = 0; /* Counter */  
pthread_mutex_t mutex;  
pthread_mutex_init(&mutex, NULL); // No special attributes
```

- 用加锁和解锁包围临界区: Surround critical section with *lock* and *unlock*:

```
for (i = 0; i < niters; i++) {  
    pthread_mutex_lock(&mutex);  
    cnt++;  
    pthread_mutex_unlock(&mutex);  
}
```


```
linux> ./goodmcent 10000  
OK cnt=20000  
linux> ./goodmcent 10000  
OK cnt=20000  
linux>
```

Function	badcnt	goodmcent
Time (ms) niters = 10 ⁶	12.0	214.0
减速 Slowdown	1.0	17.8

为什么互斥锁有效 Why Mutexes Work



线程2 Thread 2



通过加锁和解锁操作围绕临界区，提供对共享变量的互斥访问 Provide mutually exclusive access to shared variable by surrounding critical section with *lock* and *unlock* operations


线程1 Thread 1

初始 Initially
 $m = 1$

为什么互斥锁有效 Why Mutexes Work



线程2 Thread 2



通过加锁和解锁操作围绕临界区，提供对共享变量的互斥访问 Provide mutually exclusive access to shared variable by surrounding critical section with *lock* and *unlock* operations

互斥锁恒定大于等于零的特性创建了一个封闭不安全区域的禁区，任何轨迹都无法进入 Mutex invariant creates a *forbidden region* that encloses unsafe region and that cannot be entered by any trajectory.


线程1 Thread 1

初始 Initially
 $m = 1$

为什么互斥锁有效 Why Mutexes Work



线程2 Thread 2



通过加锁和解锁操作围绕临界区，提供对共享变量的互斥访问 Provide mutually exclusive access to shared variable by surrounding critical section with *lock* and *unlock* operations

互斥锁恒定大于等于零的特性创建了一个封闭不安全区域的禁区，任何轨迹都无法进入 Mutex invariant creates a *forbidden region* that encloses unsafe region and that cannot be entered by any trajectory.


线程1 Thread 1

初始 Initially
 $m = 1$

为什么互斥锁有效 Why Mutexes Work



线程2 Thread 2



通过加锁和解锁操作围绕临界区，提供对共享变量的互斥访问 Provide mutually exclusive access to shared variable by surrounding critical section with *lock* and *unlock* operations

互斥锁恒定大于等于零的特性创建了一个封闭不安全区域的**禁区**，任何轨迹都无法进入 Mutex invariant creates a **forbidden region** that encloses unsafe region and that cannot be entered by any trajectory.

线程1 Thread 1

初始 Initially
 $m = 1$



议题 Today

- 线程回顾/Threads review
- 共享 Sharing
- 互斥 Mutual exclusion
- **信号量 Semaphores**
- **生产者-消费者同步 Producer-Consumer Synchronization**

信号量 Semaphores



- **信号量:** 非负全局整数同步变量, 由P和V操作操纵 **Semaphore: non-negative global integer synchronization variable. Manipulated by P and V operations.**
- **P(s)**
 - 如果s为非零, 则将s减1并立即返回 If s is nonzero, then decrement s by 1 and return immediately.
 - 测试和减1操作以原子方式发生 (不可分割) Test and decrement operations occur atomically (indivisibly)
 - 如果s为零, 则挂起线程, 直到s变为非零, 并通过V操作重新启动线程 If s is zero, then suspend thread until s becomes nonzero and the thread is restarted by a V operation.
 - 重新启动后, P操作将s减1并将控制权返回给调用者 After restarting, the P operation decrements s and returns control to the caller.
- **V(s):**
 - 将s递增1 Increment s by 1.
 - 增量操作以原子方式发生 Increment operation occurs atomically
 - 如果在P操作中有任何线程被阻塞, 等待s变为非零, 那么只重新启动其中一个线程, 然后通过将s减1来完成P操作 If there are any threads blocked in a P operation waiting for s to become non-zero, then restart exactly one of those threads, which then completes its P operation by decrementing s.
- **信号量恒定大于等于零: Semaphore invariant: ($s \geq 0$)**

信号量 Semaphores



- **信号量:** 非负全局整数同步变量 **Semaphore:** non-negative global integer synchronization variable
- **由P和V操作操纵 Manipulated by P and V operations:**
 - $P(s): [\text{while } (s == 0) \text{ wait}(); s--;]$
 - 荷兰语单词“Proberen” (测试) Dutch for "Proberen" (test)
 - $V(s): [s++;]$
 - 荷兰语单词“Verhogen” (增加) Dutch for "Verhogen" (increment)
- **OS内核保证括号[]之间的操作不可分割地执行 OS kernel guarantees that operations between brackets [] are executed indivisibly**
 - 一次只能一个P或V操作修改s Only one P or V operation at a time can modify s .
 - 当P中的while循环终止时, 只有该P操作可以减少s When **while** loop in P terminates, only that P can decrement s
- **信号量恒定大于等于零: Semaphore invariant: ($s \geq 0$)**

C语言信号量操作

C Semaphore Operations



Pthread函数 Pthreads functions:

```
#include <semaphore.h>

int sem_init(sem_t *s, 0, unsigned int val); /* s = val */

int sem_wait(sem_t *s); /* P(s) */
int sem_post(sem_t *s); /* V(s) */
```

CS: APP包装器函数 CS:APP wrapper functions:

```
#include "csapp.h"

void P(sem_t *s); /* Wrapper function for sem_wait */
void V(sem_t *s); /* Wrapper function for sem_post */
```




使用信号量协调共享资源的访问

Using Semaphores to Coordinate Access to Shared Resources

- **基本思想：线程使用信号量操作通知另一个线程某些条件已变为真** **Basic idea: Thread uses a semaphore operation to notify another thread that some condition has become true**
 - 使用计数信号量来跟踪资源状态 Use counting semaphores to keep track of resource state.
 - 使用二元信号量通知其他线程 Use binary semaphores to notify other threads.
- **生产者-消费者问题 The Producer-Consumer Problem**
 - 对进程之间的交互操作进行协调，一个进程产生信息，另一个进程使用这些消息 Mediating interactions between processes that generate information and that then make use of that information

生产者-消费者问题


Producer-Consumer Problem



- 通用同步模式: **Common synchronization pattern:**
 - 生产者等待空槽, 将项目插入缓冲区, 并通知消费者 Producer waits for empty **slot**, inserts item in buffer, and notifies consumer
 - 消费者等待项目, 将其从缓冲区中删除, 并通知生产者 Consumer waits for **item**, removes it from buffer, and notifies producer

生产者-消费者问题

Producer-Consumer Problem



■ 示例 Examples

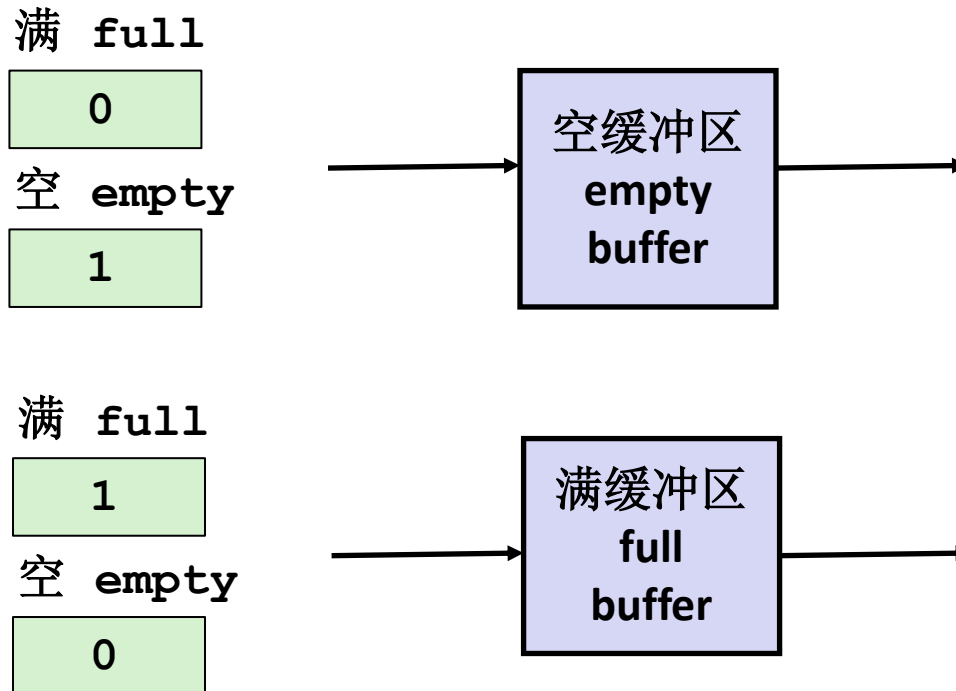
- 多媒体处理：Multimedia processing:
 - 生产者创建视频帧，消费者对其进行渲染 Producer creates video frames, consumer renders them
- 事件驱动的图形用户界面 Event-driven graphical user interfaces
 - 生产者检测鼠标点击、鼠标移动和键盘点击，并在缓冲区中插入相应的事件 Producer detects mouse clicks, mouse movements, and keyboard hits and inserts corresponding events in buffer
 - 消费者从缓冲区检索事件并绘制显示 Consumer retrieves events from buffer and paints the display

生产者和消费者之间有1个元素的缓冲区

Producer-Consumer on 1-element Buffer




- 维护两个信号量：缓冲区满+缓冲区空 Maintain two semaphores: `full` + `empty`



生产者和消费者之间有1个元素的缓冲区

Producer-Consumer on 1-element Buffer



```
#include "csapp.h"

#define NITERS 5

void *producer(void *arg);
void *consumer(void *arg);

struct {
    int buf; /* shared var */
    sem_t full; /* sems */
    sem_t empty;
} shared;
```

```
int main(int argc, char** argv) {
    pthread_t tid_producer;
    pthread_t tid_consumer;

    /* Initialize the semaphores */
    Sem_init(&shared.empty, 0, 1);
    Sem_init(&shared.full, 0, 0);

    /* Create threads and wait */
    Pthread_create(&tid_producer, NULL,
                  producer, NULL);
    Pthread_create(&tid_consumer, NULL,
                  consumer, NULL);

    Pthread_join(tid_producer, NULL);
    Pthread_join(tid_consumer, NULL);

    return 0;
}
```

生产者和消费者之间有1个元素的缓冲区

Producer-Consumer on 1-element Buffer



初始: Initially: $empty == 1$, $full == 0$

生产者线程 Producer Thread

```
void *producer(void *arg) {
    int i, item;

    for (i=0; i<NITERS; i++) {
        /* Produce item */
        item = i;
        printf("produced %d\n",
            item);

        /* Write item to buf */
        P(&shared.empty);
        shared.buf = item;
        V(&shared.full);
    }
    return NULL;
}
```

消费者线程 Consumer Thread

```
void *consumer(void *arg) {
    int i, item;

    for (i=0; i<NITERS; i++) {
        /* Read item from buf */
        P(&shared.full);
        item = shared.buf;
        V(&shared.empty);


        /* Consume item */
        printf("consumed %d\n", item);
    }
    return NULL;
}
```


为何对一个条目的缓冲区使用2个信号量?



Why 2 Semaphores for 1-Entry Buffer?

- 考虑多个生产者和多个消费者 Consider multiple producers & multiple consumers



- 生产者将与每个人竞争以获得空缓冲区 Producers will contend with each to get empty
- 消费者将相互竞争以获得满缓冲区 Consumers will contend with each other to get full


生产者 Producers

```
P(&shared.empty);  
shared.buf = item;  
V(&shared.full);
```

空 empty



满 full




消费者 Consumers

```
P(&shared.full);  
item = shared.buf;  
V(&shared.empty);
```

生产者和消费者之间有n个元素的缓冲区

Producer-Consumer on an n -element Buffer



- 使用名为sbuf的共享缓冲区包实现 Implemented using a shared buffer package called sbuf.

环形缓冲区 (n=10)

Circular Buffer (n = 10)



- 将元素存储在大小为n的数组中 Store elements in array of size n
- 项目: 缓冲区中的元素数 items: number of elements in buffer
- 空缓冲区: Empty buffer:
 - $\text{front} = \text{rear}$
- 非空缓冲区 Nonempty buffer
 - rear: 最近插入的元素的索引 | rear: index of most recently inserted element
 - front: (要删除的下一个元素的索引-1) mod n | front: (index of next element to remove - 1) mod n

- 初始 Initially:

front	0
rear	0
项目 items	0


0	1	2	3	4	5	6	7	8	9




环形缓冲区操作 (n=10)

Circular Buffer Operation (n = 10)


■ 插入7个元素 Insert 7 elements




■ 删除5个元素 Remove 5 elements



■ 插入6个元素 Insert 6 elements



■ 删除8个元素 Remove 8 elements



顺序环形缓冲区代码

Sequential Circular Buffer Code




```
init(int v)
{
    items = front = rear = 0;
}
```

```
insert(int v)
{
    if (items >= n)
        error();
    if (++rear >= n) rear = 0;
    buf[rear] = v;
    items++;
}
```

```
int remove()
{
    if (items == 0)
        error();
    if (++front >= n) front = 0;
    int v = buf[front];
    items--;
    return v;
}
```

生产者和消费者之间有 n 个元素的缓冲区

Producer-Consumer on an n -element Buffer



- 需要一个互斥锁和两个计数信号量: **Requires a mutex and two counting semaphores:**
 - 互斥锁: 执行对缓冲区和计数器进行互斥访问 `mutex`: enforces mutually exclusive access to the buffer and counters
 - 槽位数: 统计缓冲区中的可用槽位 `slots`: counts the available slots in the buffer
 - 项目: 统计缓冲区中的可用项目 `items`: counts the available items in the buffer
- 使用通用信号量 **Makes use of general semaphores**
 - 值范围从0到 n Will range in value from 0 to n



sbuf包-声明 sbuf Package - Declarations

```
#include "csapp.h"

typedef struct {
    int *buf;          /* Buffer array */
    int n;             /* Maximum number of slots */
    int front;         /* buf[front+1 (mod n)] is first item */
    int rear;          /* buf[rear] is last item */
    pthread_mutex_t mutex; /* Protects accesses to buf */
    sem_t slots;        /* Counts available slots */
    sem_t items;        /* Counts available items */
} sbuf_t;

void sbuf_init(sbuf_t *sp, int n);
void sbuf_deinit(sbuf_t *sp);
void sbuf_insert(sbuf_t *sp, int item);
int sbuf_remove(sbuf_t *sp);
```

sbuf.h



sbuf包-实现

sbuf Package - Implementation

初始化和释放共享缓冲区 Initializing and deinitializing a shared buffer:

```
/* Create an empty, bounded, shared FIFO buffer with n slots */
void sbuf_init(sbuf_t *sp, int n)
{
    sp->buf = Calloc(n, sizeof(int));
    sp->n = n; /* Buffer holds max of n items */
    sp->front = sp->rear = 0; /* Empty buffer iff front == rear */
    pthread_mutex_init(&sp->mutex, NULL); /* lock */
    Sem_init(&sp->slots, 0, n); /* Initially, buf has n empty slots */
    Sem_init(&sp->items, 0, 0); /* Initially, buf has zero items */
}

/* Clean up buffer sp */
void sbuf_deinit(sbuf_t *sp)
{
    Free(sp->buf);
}
```

sbuf.c



sbuf包-实现

sbuf Package - Implementation

插入一个项目到共享缓冲区 Inserting an item into a shared buffer:

```
/* Insert item onto the rear of shared buffer sp */
void sbuf_insert(sbuf_t *sp, int item)
{
    P(&sp->slots); /* Wait for available slot */
    pthread_mutex_lock(&sp->mutex); /* Lock the buffer */
    if (++sp->rear >= sp->n) /* Increment index (mod n) */
        sp->rear = 0;
    sp->buf[sp->rear] = item; /* Insert the item */
    pthread_mutex_unlock(&sp->mutex); /* Unlock the buffer */
    V(&sp->items); /* Announce available item */
}
```

sbuf.c



sbuf包-实现

sbuf Package - Implementation

从共享缓冲区删除一个项目 Removing an item from a shared buffer:

```
/* Remove and return the first item from buffer sp */
int sbuf_remove(sbuf_t *sp)
{
    int item;
    P(&sp->items);          /* Wait for available item */
    pthread_mutex_lock(&sp->mutex); /* Lock the buffer */
    if (++sp->front >= sp->n)      /* Increment index (mod n) */
        sp->front = 0;
    item = sp->buf[sp->front];    /* Remove the item */
    pthread_mutex_unlock(&sp->mutex); /* Unlock the buffer */
    V(&sp->slots);              /* Announce available slot */
    return item;
}
```

sbuf.c



演示 Demonstration

- 参见code目录中的程序produce-consume.c See program produce-consume.c in code directory
- 10个条目的共享环形缓冲区 10-entry shared circular buffer
- 5个生产者 5 producers
 - 代理 i 生成从 $20*i$ 到 $20*i-1$ 的数字 Agent i generates numbers from $20*i$ to $20*i - 1$.
 - 将它们放入缓冲区 Puts them in buffer
- 5个消费者 5 consumers
 - 每个从缓冲区中检索20个元素 Each retrieves 20 elements from buffer
- 主程序 Main program
 - 确保0到99之间的每个值检索一次 Makes sure each value between 0 and 99 retrieved once



小结 Summary

- 程序员需要一个线程如何共享变量的清晰模型。
Programmers need a clear model of how variables are shared by threads.
- 必须保护多个线程共享的变量，以确保互斥访问
Variables shared by multiple threads must be protected to ensure mutually exclusive access.
- 信号量是执行互斥的基本机制 **Semaphores are a fundamental mechanism for enforcing mutual exclusion.**